




Pure Distillation



Pure RL



Hybrid KD + RL



Unified Objective

$$\max_{\theta} \mathbb{E}_{y \sim p_{\theta}}[R(y)] - \beta \cdot D_{\text{KL}}(p_{\theta} \| p_T)$$

RL objective

KD regularizer

G-OPD proved: standard OPD is a special case ($\beta = 1$, $R = \log p_T$)