

A Survey of On-Policy Distillation for Large Language Models

Mingyang Song & Mao Zheng

Large Language Model Department

Tencent, China

nickmysong@tencent.com

 [nick7nlp/Awesome-LLM-On-Policy-Distillation](https://github.com/nick7nlp/Awesome-LLM-On-Policy-Distillation)

Abstract

As Large Language Models (LLMs) continue to grow in both capability and cost, transferring frontier capabilities into smaller, deployable students has become a central engineering problem, and knowledge distillation remains the dominant technique for this transfer. The prevailing recipe in industrial pipelines, static imitation of teacher-generated text, carries a structural weakness that grows more severe as tasks become longer and more reasoning-intensive. Because the student is trained on flawless teacher prefixes but must generate its own at inference, small errors tend to accumulate into trajectories it has rarely been trained to recover from, and the resulting exposure bias has been shown to scale roughly with the square of sequence length. On-Policy Distillation (OPD) reorganizes the training loop around this observation by having the teacher provide feedback on what the student *actually produces*, with the goal of reducing the compounding term toward linear and reframing distillation as an iterative correction process rather than single-pass imitation. The resulting literature has expanded along divergence design, reward-guided optimization, and self-play, yet contributions remain scattered across the knowledge distillation, RLHF, and imitation learning communities without a unified treatment. This survey provides such a treatment. We formalize OPD as f -divergence minimization over student-sampled trajectories, organize the field along three design axes (what to optimize, where the signal comes from, and how to stabilize training in practice), and consolidate success conditions, recurring failure modes, and the connection between OPD and KL-constrained reinforcement learning. We close with open problems that emerge from this synthesis, including distillation scaling laws, uncertainty-aware feedback, agent-level distillation, and the growing overlap between knowledge distillation and RL.

1 Introduction

The capability of Large Language Models (LLMs) has advanced rapidly over the past two years, with recent systems reaching state-of-the-art performance on reasoning, code generation, and multilingual benchmarks (Yang et al., 2025a; DeepSeek-AI et al., 2025; Gemma Team et al., 2024). This progress has come at a training cost that only a few organizations can sustain. Transferring the resulting capabilities from large frontier models into smaller deployable students has therefore become a core step in the modern LLM pipeline. Knowledge distillation, first formalized by Hinton et al. (2015) as a way for a student to inherit the soft output structure of a teacher, has broadened in scope over this period. What began as a model compression technique for transferring knowledge within the same architecture now serves as a general mechanism for moving capabilities across model scales and architectural families. The release of DeepSeek-R1 (DeepSeek-AI et al., 2025) made this shift concrete, distilling a 671B mixture-of-experts teacher into dense students spanning 1.5B to 70B parameters while preserving long chain-of-thought reasoning largely intact.

The dominant recipe behind these successes, however, has a known structural weakness. In almost all large-scale pipelines, distillation proceeds *off-policy*. The student is trained to match the teacher’s next-token distribution over a fixed corpus, typically the original pre-training data or traces generated in advance by the teacher, and every gradient step conditions on flawless teacher prefixes. At inference, this assumption no longer holds. The student must generate autoregressively from its own partial outputs, and any divergence between the training states and the states it visits at deployment translates into compounding error along the trajectory. This corresponds to the *compounding error* identified in interactive imitation learning (Ross et al., 2011), where autoregressive generation amplifies the distribution mismatch quadratically in sequence length, producing an expected discrepancy of $O(\epsilon T^2)$ over a horizon T . The consequences are most pronounced on reasoning tasks, where a single early misstep can derail an entire proof or program. Gudibande et al. (2023) report that off-policy imitation of proprietary LLMs often yields students that reproduce surface style without acquiring the underlying reasoning competence.

Addressing this mismatch requires changing *where* the training data comes from rather than *what* is being matched. **On-Policy Distillation (OPD)** addresses this by shifting the sampling distribution from a static corpus to the student’s own evolving policy. The student proposes trajectories, populates the states it will visit at deployment, and receives teacher feedback on those states rather than on idealized expert prefixes. Depending on the available teacher access, this feedback ranges from full token-level distributions in white-box settings (Agarwal et al., 2024; Gu et al., 2024), through scalar rewards or pairwise preferences when only black-box APIs are available (Ye et al., 2025; Jiang et al., 2023), to contrastive signals the model derives from its own outputs in teacher-free settings (Chen et al., 2024). The theoretical payoff mirrors the classical DAgger result (Ross et al., 2011), where querying an expert on the learner’s own visited states reduces the $O(\epsilon T^2)$ compounding of off-policy imitation to $O(\epsilon T)$. Translated into autoregressive language modeling, this turns distillation from a single-pass matching problem into an iterative optimization loop in which the training distribution co-evolves with the model, allowing distributional gaps between training and deployment to be corrected during training rather than deferred to inference.

The body of work that has emerged from this reformulation can be organized around a small number of recurring design choices. Early methods focused on *how* to sample and *which* divergence to apply on student trajectories (Agarwal et al., 2024; Gu et al., 2024; Ko et al., 2024), establishing that an appropriate divergence depends on where in the sequence one is training and what kind of supervision the teacher can supply. A parallel line of work established that the resulting objective is formally equivalent to a KL-constrained form of reinforcement learning (Yang et al., 2026c), bridging two previously separate research communities. A third branch established that the “teacher” need not be external at all, with methods such as SPIN (Chen et al., 2024) and OPSD (Zhao et al., 2026b) showing that a model can bootstrap improvement from its own generations by exploiting asymmetries in its output distribution across contexts. These ideas have since moved from papers into production pipelines, with Qwen3 (Yang et al., 2025a), DeepSeek-V4 (DeepSeek-AI, 2026), Gemma 2 (Gemma Team et al., 2024), and MiMo-V2 (Xiaomi LLM-Core Team et al., 2026) all adopting OPD as a core training ingredient. The adoption cuts across architectural lines, covering both mixture-of-experts and dense systems and spanning proprietary and open-weight release strategies, because all of these systems face the same structural problem once reasoning chains grow long enough for compounding error to dominate. DeepSeek-V4 went the furthest, replacing its mixed RL stage with pure multi-teacher OPD for model consolidation.

GKD (Agarwal et al., 2024), alongside the concurrent MiniLLM (Gu et al., 2024), brought on-policy distillation into the LLM era in mid-2023, and in less than two years the literature has expanded to over one hundred papers spanning divergence design, reward-guided optimization, self-play, multi-teacher debate, agentic trajectory distillation, and cross-modal transfer. The research focus has shifted accordingly. Where early work asked whether on-policy training helps at all, the current frontier asks how to make it efficient enough for trillion-parameter teacher inference, how to extend it beyond single-turn generation into multi-step agent trajectories, and how to unify it with the reinforcement learning pipelines that increasingly share its computational infrastructure. This trajectory parallels

a broader shift in LLM development. As models become more capable and reasoning chains grow longer, the gap between off-policy training states and deployment-time states widens, which makes on-policy correction increasingly attractive for the next generation of reasoning-capable systems (Fu et al., 2026; DeepSeek-AI, 2026).

Despite this rapid adoption, the research literature on OPD has yet to catch up with the practice. Methods approach the same problem through the separate lenses of knowledge distillation, RLHF, and interactive imitation learning, inheriting different notations, benchmarks, and failure taxonomies from their parent communities. Existing surveys of LLM distillation (Xu et al., 2024) generally retain the classical compression framing, treating off-policy and on-policy methods as interchangeable variants instead of as regimes with materially different theoretical guarantees and failure modes. As a result, no current treatment we are aware of offers a unified mathematical account of the on-policy dynamics that connect these threads, and a systematic comparison across the white-box, black-box, and teacher-free instantiations that researchers and engineers must now weigh against one another in practice remains missing.

This survey addresses that gap through four contributions that also structure the rest of the paper.

- **A unified theoretical framework.** We recast the transition from off-policy to on-policy distillation as a sequential decision-making problem and show that the core OPD algorithms are instances of f -divergence minimization over student-sampled trajectories (Section 2). This framework supplies a common analytical vocabulary for methods previously studied in isolation and makes their relationships explicit in terms of divergence choice, argument ordering, and sampling mixture.
- **A design-centric taxonomy.** Rather than grouping methods by surface similarity, we organize them along three core design axes (Section 3), namely objective function design (Section 4), signal source architecture (Section 5), and training dynamics optimization (Section 6). The resulting one-method-one-category classification makes the relationships among methods explicit in terms of the specific design choices each paper optimizes.
- **A theoretical account of failure.** We consolidate the empirical record into a coherent view of when and why OPD breaks down (Section 7), covering the flawed prefix trap, self-play saturation, diversity collapse, the calibration-capability gap, and emerging multi-turn agentic failure modes, and we state formal conditions under which OPD provably improves over off-policy alternatives.
- **Comprehensive coverage and comparative analysis.** We survey the white-box, black-box, and self-distillation regimes within this common framework, summarize them in three comparison tables (Tables 1, 2, and 3), and document method selection factors (Section 3.3) that map deployment constraints to applicable methods and compute budgets.

The structure of the paper follows the same design-axis logic. Section 2 develops the mathematical preliminaries, formalizes autoregressive generation as a sequential decision problem, and derives the exposure bias that motivates on-policy training. Section 3 presents the method landscape and selection considerations. Sections 4–6 then traverse the three design axes from “what to optimize” through “where the signal comes from” to “how to stabilize training in practice,” with each section building on the choices surfaced in the previous one. Section 7 consolidates the theoretical perspectives, showing that the patterns observed across Sections 4–6 follow from a small number of unifying principles. Section 8 discusses industrial deployments, system-level optimization, and emerging domains, and Section 9 lays out the open problems that emerge from this synthesis.

Our scope is centered on the on-policy regime. We cover methods in which the student generates its own training data during distillation, and we include work from adjacent fields such as imitation learning, online RL, and preference optimization whenever it operates on student-generated data under teacher or verifier supervision. We exclude generic off-policy KD, compression techniques orthogonal to the training regime such as pruning and quantization, and inference-time methods that leave model weights unchanged.

2 Background and Unified Math

To understand why on-policy training is often preferable, we formalize the compounding geometry of autoregressive generation. We begin with classical KD and its limitations, then establish the unified f -divergence framework that subsumes modern OPD objectives.

Defining “on-policy.” A distillation method is *on-policy* if the training data for the student is sampled from the student’s own current policy p_θ at training time, rather than from a fixed external corpus $\mathcal{D}$ or from the teacher’s generation distribution p_T . Formally, on-policy training optimizes:

$$\min_{\theta} \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y \sim p_\theta(\cdot|x)} [\mathcal{L}(y, x; \theta, T)] \quad (1)$$

where $\mathcal{L}$ is the distillation loss (divergence, reward, or hybrid). The key distinction is that the *outer* expectation is over the student’s own generations, not a static dataset. This creates a non-stationary optimization landscape because as θ updates, the data distribution p_θ shifts, requiring fresh rollouts at each training step. The computational cost of this fresh generation is a central systems-level challenge of OPD (Section 6.3).

Notation. Throughout this survey, we use lowercase p for token-level conditional distributions ($p_T(\cdot|x, y_{<t})$ for teacher, $p_\theta(\cdot|x, y_{<t})$ for student) and uppercase P for sequence-level distributions ($P_T(y|x)$, $P_\theta(y|x)$). $\mathcal{D}$ denotes the training dataset, p_{data} the data distribution, and $|V|$ the vocabulary size. D_{KL} and D_f denote KL divergence and general f -divergence respectively. In the RL formalization, π denotes policies (e.g., π_{mix} for a mixture policy, π_{teach} for a sequence-level teacher policy in DPO-style derivations, π_{ref} for a reference policy) and $R(x, y)$ the outcome reward. When discussing individual methods, we follow each paper’s conventions where clarity demands.

2.1 Classical Knowledge Distillation

The original KD formulation (Hinton et al., 2015) trains the student to match the teacher’s temperature-softened output distribution:

$$p_T^{(\tau)}(y|x) = \frac{\exp(z_y/\tau)}{\sum_j \exp(z_j/\tau)} \quad (2)$$

where z_y are the teacher’s logits and $\tau > 1$ is a temperature parameter. When $\tau \rightarrow \infty$, the softened distribution approaches a uniform distribution over the vocabulary, exposing the teacher’s “dark knowledge” about inter-class similarities that hard labels hide. The gradient of the distillation loss with respect to the student’s logits z_i^S shows why temperature matters:

$$\frac{\partial \mathcal{L}_{\text{KD}}}{\partial z_i^S} = \frac{1}{\tau} (p_i^S - p_i^T) \quad (3)$$

In the high-temperature limit ($\tau \rightarrow \infty$), a Taylor expansion $\exp(z_i/\tau) \approx 1 + z_i/\tau$ together with the zero-mean assumption $\sum_i z_i^S = \sum_i z_i^T = 0$ reduces the gradient to $\frac{1}{|V|\tau^2} (z_i^S - z_i^T)$ (Hinton et al., 2015), so classical KD in this regime is equivalent to minimizing the mean squared error between raw logits. This forces the student to replicate the teacher’s *entire* logit structure, transferring both primary predictive modes and nuanced sub-optimal probability masses (the “dark knowledge”). The $1/\tau^2$ factor explains why the soft-target term is commonly upweighted by τ^2 relative to the hard-label cross-entropy in the combined KD objective. In practice, LLM distillation operates at $\tau = 1$ or low-to-moderate temperatures (Zhong et al., 2024) because the large vocabulary ($|V| > 30,000$) already produces richly structured non-peak probabilities without explicit softening, and higher temperatures may amplify noise in the teacher’s poorly calibrated tail distribution.

For autoregressive LLMs, the token-level distillation loss factors as:

$$\mathcal{L}_{\text{Token-KD}} = \mathbb{E}_{x, y \sim \mathcal{D}} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_T(\cdot|x, y_{<t}) \parallel p_\theta(\cdot|x, y_{<t})) \right] \quad (4)$$

This is computationally tractable because it relies on static, pre-computed prefixes from the dataset, but it optimizes next-token accuracy while assuming an error-free history $y_{<t}$. Kim & Rush (2016) extended this to Sequence-Level Knowledge Distillation (Seq-KD), which instead minimizes the global KL divergence over the entire sequence space:

$$\mathcal{L}_{\text{Seq-KD}} = D_{\text{KL}}(P_T(y|x) \parallel P_\theta(y|x)) = \sum_{y \in \mathcal{Y}} P_T(y|x) \log \frac{P_T(y|x)}{P_\theta(y|x)} \quad (5)$$

Since $|\mathcal{Y}| = |V|^T$ grows exponentially with length, exact computation is intractable. Seq-KD approximates $P_T(y|x)$ with a Dirac delta at the beam-search output $\hat{y} = \arg \max_y P_T(y|x)$, reducing to standard negative log-likelihood on teacher-generated sequences:

$$\mathcal{L}_{\text{Seq-KD}} \approx -\log P_\theta(\hat{y}|x) \quad (6)$$

This approximation works well when the teacher is highly confident but discards all information about the teacher’s uncertainty and alternative modes. The practical consequence is twofold. Seq-KD collapses the teacher’s distributional richness into a single point estimate, and it still trains on a *static* dataset disconnected from the student’s own generation behavior. These two limitations (loss of distributional signal and off-policy training) motivate the richer objectives below.

Subsequent work (Zhong et al., 2024) revisited these foundations for modern LLMs and showed that classical assumptions (small teacher-student gap, shared vocabulary, similar architectural inductive biases) frequently fail when distilling across heterogeneous model families. The gap between classical assumptions and modern reality motivates the more general f -divergence framework below.

From classical KD to on-policy. The historical progression from Hinton’s classical KD to modern OPD amounts to a series of relaxations of four classical assumptions: (1) shared vocabulary, (2) i.i.d. data, (3) static teacher, and (4) off-policy training data. Seq-KD (Kim & Rush, 2016) relaxes assumption (2) by operating at the sequence level but retains the static data assumption. GKD (Agarwal et al., 2024) relaxes assumption (4) by training on student-generated sequences. DSKD (Zhang et al., 2025b) relaxes assumption (1) through dual-space alignment. G-OPD (Yang et al., 2026c) relaxes assumptions (2)–(4) simultaneously while adding RL-augmented reward extrapolation that pushes the student beyond the teacher’s frontier. Each relaxation targets a specific failure mode that intensifies with model scale: vocabulary mismatch is irrelevant for same-family distillation but catastrophic for cross-family transfer (relaxation 1), while exposure bias is negligible for short sequences but devastating for multi-step reasoning (relaxation 4). This progression is not merely historical but informative: the relevant assumptions in a given deployment scenario determine which level of relaxation is necessary, and unnecessary relaxation introduces avoidable complexity.

2.2 Off-Policy Exposure Bias

The token-level loss in Eq. (4) conditions on ground-truth prefixes $y_{<t}$ drawn from a static dataset. During inference, the student generates its own prefixes, creating a distribution mismatch between training and deployment that compounds over sequence length. We now formalize why this mismatch is catastrophic.

Standard knowledge distillation minimizes the KL divergence over states drawn from the dataset distribution $d_{\mathcal{D}}(s)$:

$$\mathcal{L}_{\text{Off-Policy}} = \mathbb{E}_{x, y \sim p_{\text{data}}} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_T(\cdot|x, y_{<t}) \parallel p_\theta(\cdot|x, y_{<t})) \right] \quad (7)$$

At inference, the student acts according to its own policy, inducing a different state visitation distribution $d_{\pi_\theta}(s)$. By the DAGger theorem (Ross et al., 2011), if a policy mimics an expert with per-step error ϵ under the training distribution, the expected total discrepancy over a trajectory of length T under the learner’s own state visitation scales as $O(\epsilon T^2)$.

This quadratic compounding has severe practical consequences for reasoning tasks. Even under the simplest independent-error assumption, a mathematical proof requiring 10 reasoning steps with per-step accuracy 95% yields only $(0.95)^{10} \approx 60\%$ trajectory-level correctness.

The DAgger bound predicts something strictly worse. Because off-policy training never exposes the student to its own error states, each mistake pushes the model into regions where its next-step accuracy degrades further (distributional shift), producing $O(\epsilon T^2)$ total error rather than the $O(\epsilon T)$ of independent compounding. On-policy training offers a way to mitigate this vicious cycle by exposing the student to its own error states during training, so that it can learn recovery strategies that keep distributional drift from amplifying local mistakes into trajectory-level collapse.

Remark: The DAgger bound in LLMs. While the theoretical reduction from $O(\epsilon T^2)$ to $O(\epsilon T)$ is appealing, applying the DAgger bound to LLMs demands careful qualification. The original theorem assumes an *interactive expert* that emits optimal actions in any state. In white-box OPD, the “expert” emits a next-token distribution $p_T(y_t | \hat{y}_{<t})$ conditioned on the student’s prefix $\hat{y}_{<t}$. If the student hallucinates a severely out-of-distribution prefix, the teacher’s conditional distribution may itself become poorly calibrated, because the teacher was never trained on such inputs. Forcing the student to match this noisy distribution violates the core assumption of interactive imitation learning, potentially destabilizing training rather than recovering the $O(\epsilon T)$ bound. Jeong (2026) offer direct empirical evidence that naive OPD without stable teacher dynamics suffers catastrophic instability. Unconditional token-level matching with periodic hard-copy teacher resets leads to abrupt output collapse (KL divergence dropping from 2.637 to 0.343 at a single teacher-reset event), demonstrating that unconditional token-level matching alone is insufficient for robust OPD. This observation motivates the adaptive trust mechanisms surveyed in Section 6.1 and the stable teacher dynamics discussed in Section 7.2.

2.3 The Unified f -Divergence Framework

OPD resolves exposure bias by altering the training expectation to sample from the student’s own rollouts, or a mixture policy π_{mix} . The generalized on-policy objective decouples the sampling trajectory from the local matching metric:

$$\mathcal{L}_{\text{OPD}}(\theta) = \mathbb{E}_{y \sim \pi_{\text{mix}}} \left[\sum_{t=1}^{|y|} \mathcal{D}_f(p_T(\cdot | x, y_{<t}), p_\theta(\cdot | x, y_{<t})) \right] \quad (8)$$

Here, $\mathcal{D}_f$ represents a divergence from the f -divergence family (Wen et al., 2023). Formally, given two distributions P and Q , an f -divergence is defined as:

$$D_f(P \parallel Q) = \mathbb{E}_{y \sim Q} \left[f \left(\frac{P(y)}{Q(y)} \right) \right] \quad (9)$$

where $f : (0, \infty) \rightarrow \mathbb{R}$ is a convex generator with $f(1) = 0$. The choice of f determines the implicit weighting of likelihood ratios $p_T(y) / p_\theta(y)$:

- **Forward KL** ($f(u) = u \log u$): Mode-covering (zero-avoiding). The student must allocate probability mass wherever the teacher does, often bridging distinct teacher modes and hallucinating in the inter-mode space.
- **Reverse KL** ($f(u) = -\log u$): Mode-seeking (zero-forcing). The student collapses its mass onto the teacher’s single highest peak, ignoring secondary acceptable outputs but guaranteeing high precision.
- **JSD** ($f(u) = u \log u - (u + 1) \log \frac{u+1}{2}$): Symmetric, bounded, and smoothly interpolates between mode-covering and mode-seeking behavior.
- **α -divergence**: Parameterized family that continuously interpolates between Forward KL ($\alpha \rightarrow 1$) and Reverse KL ($\alpha \rightarrow 0$), allowing fine-grained control over the mode-seeking/covering tradeoff.

In practice, these divergences each suit different task geometries. For tasks with a *unique* correct answer (mathematical proofs, code correctness), Reverse KL’s mode-seeking behavior concentrates the student on the single optimal solution path. For tasks with *many* acceptable outputs (creative writing, open-ended QA), Forward KL’s mode-covering behavior preserves output diversity. JSD offers a compromise but lacks theoretical optimality guarantees for either extreme. All these divergences are computable as expectations under the student’s

own policy (since $D_f(P_T \parallel P_\theta) = \mathbb{E}_{y \sim p_\theta}[f(p_T(y)/p_\theta(y))]$), making them directly amenable to on-policy optimization. The sampling distribution p_θ is the policy being optimized, so gradients flow through the reparameterized samples without importance weights or off-policy corrections that would add additional variance.

This framework subsumes all modern OPD objectives. The choice of f determines the geometric behavior, while the choice of π_{mix} controls the degree of on-policy exploration. We illustrate by mapping the three foundational methods into this space.

GKD (Agarwal et al., 2024) defines $\pi_{\text{mix}} = \lambda p_\theta + (1 - \lambda)p_{\text{data}}$ as an explicit interpolation, and is agnostic to $\mathcal{D}_f$, empirically testing Forward KL, Reverse KL, and JSD. Setting $\lambda \rightarrow 1$ makes GKD purely on-policy, while $\lambda = 0$ reduces to standard off-policy KD. The λ parameter offers a smooth dial between the extremes, and moderate values ($\lambda \approx 0.5$) tend to combine the benefits of on-policy exposure with the stability of grounded, dataset-derived prefixes. GKD’s experiments report that on-policy sampling ($\lambda = 1$) consistently outperforms off-policy across the divergence choices they tested. JSD performs best on translation tasks, where the output space has moderate diversity (neither a single correct answer nor fully open-ended generation) and the symmetric nature of JSD naturally balances mode-covering and mode-seeking pressures appropriate for this intermediate regime.

Empirically, all three divergences yield competitive results on summarization and instruction-following, validating the exposure bias hypothesis. This suggests that the sampling policy (on-policy vs. off-policy) matters more than the specific divergence choice when task geometry does not strongly favor one extreme.

MiniLLM (Gu et al., 2024) selects Reverse KL as $\mathcal{D}_f = D_{\text{KL}}(p_\theta \parallel p_T)$ and employs $\pi_{\text{mix}} = (1 - \alpha)p_\theta + \alpha p_T$ with $\alpha = 0.2$ for stability. Because Reverse KL places the student in both the expectation and the log-ratio, MiniLLM reformulates optimization via REINFORCE, treating $\log(p_T/p_\theta)$ as a per-step reward.

DistiLLM (Ko et al., 2024) computes its token-level loss on student-generated prefixes but uses an adaptive scheduler that gradually increases the fraction of student-generated outputs while caching earlier rollouts in a replay buffer for sample efficiency (an “adaptive off-policy” strategy in the authors’ terminology). Its core contribution is a skewed mixture target $\tilde{p} = \alpha p_T + (1 - \alpha)p_\theta$, computing $D_{\text{KL}}(p_T \parallel \tilde{p})$. This avoids zero-division instability when $p_\theta(y) \approx 0$, because $\tilde{p}(y) \geq \alpha p_T(y) > 0$ always, achieving tractability and stability without policy gradients.

Comparative synthesis. These three methods expose a core tradeoff triangle in OPD design. GKD prioritizes *generality* (divergence-agnostic) and *simplicity* (no REINFORCE variance) at the cost of limited theoretical guidance on divergence selection. MiniLLM commits to Reverse KL for its mode-seeking precision but inherits REINFORCE’s high variance, requiring reward baselines, length penalties, and careful α tuning. DistiLLM avoids both the divergence selection problem and the variance problem by engineering a mixture target that is numerically stable by construction, but at the cost of adding a replay buffer (mixing on-policy and off-policy data) and an adaptive scheduler that adds hyperparameter complexity. The progression from GKD through MiniLLM to DistiLLM thus traces a shift from algorithmic simplicity toward computational and engineering sophistication, with each method solving a specific failure mode of its predecessor.

Modern OPD, viewed through this decomposition, is not a series of disparate heuristics but a systematic exploration of a design space governed by three interacting choices, namely trajectory sampling (π_{mix}), divergence generator (f), and argument ordering within the chosen divergence.

2.4 Distillation Scaling Laws

The unified f -divergence framework characterizes *what* to optimize, but leaves open *how performance scales* with computational investment. Unlike pre-training, where Hoffmann et al. (2022) established compute-optimal scaling laws (the Chinchilla laws), distillation creates a qualitatively different resource allocation problem. The budget is typically split between

teacher inference, student rollouts, and gradient updates, with non-trivial interactions between these axes.

Busbridge et al. (2025) initiated the study of distillation-specific scaling and found that the teacher-student relationship obeys qualitatively different dynamics from pre-training alone. Their key finding is that a *capacity gap* emerges when the teacher becomes too powerful relative to the student. The student’s loss follows a power law in teacher cross-entropy that transitions between two regimes: initially, stronger teachers improve the student (denser gradient signal from better-calibrated next-token probabilities), but beyond a critical teacher capacity the student’s ability to absorb the signal saturates. This capacity-gap phenomenon connects directly to the divergence choice in Section 2.3: Forward KL’s mode-covering pressure exacerbates the gap (the student must spread mass across all teacher modes, including those beyond its capacity), while Reverse KL’s mode-seeking behavior partially mitigates it (the student selectively matches its highest-confidence subset).

No comprehensive framework yet predicts how distillation quality scales jointly with teacher size N_T , student size N_S , data volume D , and on-policy rollout budget R . A natural extension of the Busbridge law might take the form $\text{Quality} \propto N_T^\alpha \cdot N_S^\beta \cdot D^\gamma \cdot R^\delta$, but pinning down the exponents and interaction terms for on-policy methods remains open (Section 9). The rollout budget R is unique to on-policy distillation, adding a new scaling axis absent from both pre-training and off-policy distillation. Each rollout step requires a full forward pass through the student, making compute cost scale linearly with R while the learning benefit likely exhibits diminishing returns.

3 Landscape and Method Selection

Before examining individual methods in detail, this section provides a structured overview of the field, covering the organizational taxonomy, a comparative table of representative methods, and a discussion of method selection considerations.

3.1 Method Landscape

Existing OPD methods can be organized along three axes that correspond to sequential design decisions in the training pipeline: (1) the objective function to optimize, (2) the source of the supervisory signal, and (3) the mechanism for stabilizing training dynamics. Figure 1 illustrates this organization.

To illustrate the taxonomy in practice, GKD (Agarwal et al., 2024) selects forward KL (or JSD) as its divergence objective, relies on white-box teacher logits as its signal source, and interpolates between student-generated and ground-truth sequences (π_{mix}) to stabilize on-policy training dynamics. The DeepSeek-R1 distilled models (DeepSeek-AI et al., 2025) follow the same three-stage logic at industrial scale, adopting cross-entropy as the objective, 800K curated samples (predominantly reasoning trajectories) distilled from R1 as the signal source, and cosine learning-rate decay with difficulty-aware curricula for stability.

The three stages interact in both constraining and reinforcing ways. Some combinations are incompatible. Forward KL in its exact token-level form requires the teacher’s full output distribution, ruling out API-constrained settings where only generated text is available. Other combinations reinforce each other, as RL-augmented objectives naturally couple with external feedback sources (verifiers, reward models), while fixed divergences align with white-box logit access that permits exact gradient computation. An objective choice therefore constrains the viable signal sources and dynamics strategies.

Chronologically, the distribution of methods reflects shifting research emphasis. Early work (2023–2024) concentrated on the objective axis, debating forward KL, reverse KL, and JSD (Agarwal et al., 2024; Gu et al., 2024). By mid-2025, the focus shifted to signal architecture, particularly self-distillation methods that eliminate external teacher dependence. The most recent work (late 2025–2026) addresses training dynamics, specifically the instabilities that on-policy sampling introduces but static-dataset distillation avoids. Industrial systems such as DeepSeek-V4 and Qwen3 increasingly combine all three axes, pairing




Figure 1: Taxonomy of On-Policy Distillation methods organized along three design axes: (1) Objective function design (§4), (2) Signal source and teacher architecture (§5), and (3) Training dynamics and efficiency (§6). Each leaf lists representative methods with a badge showing the method count. Methods are placed under their primary contribution. Those contributing to secondary dimensions (e.g., PAINT and PRISM contributing implicit curricula) are discussed in the relevant section prose rather than duplicated in the tree.

adaptive or RL-augmented objectives with multi-teacher signals and curriculum-based stabilization (Yang et al., 2026c; 2025a).

3.2 Method Comparison Table

Classification methodology. Each method is assigned to exactly one primary category based on its *core contribution dimension*. A paper whose novelty lies in the loss function (e.g., G-OPD’s KL-constrained RL reformulation) belongs to “Objective.” A new signal architecture (e.g., DSKD’s dual-space projection) places a method under “Signal.” A training stability or efficiency contribution (e.g., TIP’s token weighting) maps to “Dynamics.” Methods that contribute to multiple dimensions are classified by their *most distinctive* contribution, reducing the categorization ambiguity that arises when a single method spans multiple dimensions (Xu et al., 2024).

Table organization. Table 1 provides a compact reference mapping 15 representative methods to their primary category, signal source, and loss granularity. Table 2 extends this

to 88 methods organized by primary contribution dimension, with year, divergence type, and key innovation. Table 3 complements both by listing concrete teacher-student pairings and benchmark results across six application domains.

Distributional observations. Self-distillation is the largest single category with 28 methods (32% of all entries), nearly all published in 2026. Training Dynamics follows with 17 methods (19%), reflecting concentrated recent work on stabilizing and accelerating on-policy rollouts. RL-augmented objectives, white-box signal innovations, and black-box methods each account for 11 entries (13%), while adaptive divergences contribute 4 (5%) and fixed-divergence objectives (the earliest OPD formulations from 2024) account for only 6 entries (7%). This distribution is consistent with a gradual shift from the foundational “which KL direction” question toward self-improvement and adaptive alternatives. These proportions also match a broader trend in LLM research in which the model’s own rollouts increasingly serve as the primary data source, whether or not an external teacher is present.

Table 1: Comparison of key OPD methods categorized by their primary contribution.

Method	Category	Core Contribution	Signal Source	Loss Granularity
GKD (Agarwal et al., 2024)	Objective (Fixed)	DAGger for LLMs, interpolation	White-Box Logits	Token
DistiLLM (Ko et al., 2024)		Skewed KL mixtures	White-Box Logits	Token
MiniLLM (Gu et al., 2024)		Sequence-level Reverse KL via RL	White-Box Logits	Sequence
ToDi (Jung et al., 2025)	Objective (Adaptive)	Log-ratio-routed token divergence	White-Box Logits	Token
EOPD (Jin et al., 2026a)	Objective (RL)	Entropy-based KL interpolation	White-Box Logits	Token
G-OPD (Yang et al., 2026c)		OPD equivalence to KL-RL	White-Box Logits	Sequence/Token
RLKD (Xu et al., 2025b)		KD augmented with reward model	Reward Model	Sequence
DSKD (Zhang et al., 2025b)	Signal (White-Box)	Cross-architecture alignment	White-Box Logits	Token
Delta-KD (Cao et al., 2025)	Signal (Black-Box)	Base-to-Instruct delta isolation	White-Box Logits	Token
Lion (Jiang et al., 2023)		Verbal feedback curriculum	Black-Box API	Sequence
GAD (Ye et al., 2025)		Adversarial distribution matching	Black-Box API	Sequence
OPSD (Zhao et al., 2026b)	Signal (Self)	Privileged Info (PI) conditioning	Self-Generated	Token
SPIN (Chen et al., 2024)		Self-play vs reference responses	Self-Generated	Sequence
TIP (Xu et al., 2026b)	Dynamics (Weight)	Token importance profiling via diff	White-Box Logits	Token
PACED (Xu et al., 2026a)		Beta-kernel curriculum pacing	White-Box Logits	Token

3.3 Method Selection Considerations

Table 3 suggests that the most successful deployments tend to make deliberate choices at each of the three stages rather than optimizing a single axis in isolation. The constraints and interactions among these axes appear to reduce the space of viable configurations substantially, and we summarize the key selection factors below.

Teacher access constraints. The available signal source determines which objectives are feasible. White-box logit access enables exact token-level divergence computation, supporting methods such as DistiLLM (Ko et al., 2024), G-OPD (Yang et al., 2026c), and cross-tokenizer approaches like DSKD (Zhang et al., 2025b) and ULD (Boizard et al., 2025). API-only access restricts supervision to generated text or scalar scores, favoring methods designed for sparse feedback such as OVD (Xiong et al., 2026), GAD (Ye et al., 2025), and LUFFY (Yan et al., 2025). When no external teacher is available, self-distillation methods operate either through privileged information (OPSD (Zhao et al., 2026b), GATES (Stein et al., 2026)) or through verifier-guided self-play (SD-ZERO (He et al., 2026), SSD (Zhang et al., 2026d)).

Task characteristics. Reasoning and mathematical tasks benefit from mode-seeking objectives (reverse KL or RL-augmented variants) that concentrate probability mass on correct solutions (Gu et al., 2024; Yang et al., 2026c; Xu et al., 2025a). Open-ended generation tasks instead require mode-covering objectives (forward KL or adaptive entropy-aware divergences) that preserve output diversity. Instruction-following tasks occupy a middle ground where moderate on-policy mixing ($\lambda \geq 0.5$) with JSD provides strong quality-compute

Table 2: Comprehensive comparison of on-policy distillation methods organized by primary contribution dimension.

Method	Year	Category	Divergence/Objective	Signal	Granularity	Key Innovation
<i>Objective: Fixed Divergences</i>						
GKD (Agarwal et al., 2024)	2024	Objective	F-KL / R-KL / JSD	White-box	Token	Unified OPD framework, DAgger for LLMs
MiniLLM (Gu et al., 2024)	2024	Objective	Reverse KL	White-box	Sequence	REINFORCE with teacher as reward
DistiLLM (Ko et al., 2024)	2024	Objective	Skew KL (SKL+SRKL)	White-box	Token	α -smoothing for KL stability
DistiLLM-2 (Ko et al., 2025)	2025	Objective	Contrastive SKL/SRKL	White-box	Token	Asymmetric loss per data source
KETCHUP (Fan et al., 2026)	2025	Objective	Reverse KL	White-box	Sequence	k-step Bellman return estimation
vOPD (Oh et al., 2026)	2026	Objective	Reverse KL + CV baseline	White-box	Token	Control variate variance reduction
<i>Objective: Adaptive Divergences</i>						
AKL (Wu et al., 2025)	2025	Objective	Adaptive FKL+RKL	White-box	Token	Head/tail gap-weighted blending
ToDi (Jung et al., 2025)	2025	Objective	Log-ratio-weighted FKL/RKL	White-box	Token	Sigmoid per-token blending
EOPD (Jin et al., 2026a)	2026	Objective	Continuous FKL+RKL blend	White-box	Token	Smooth entropy interpolation
AOPD (Jia et al., 2026)	2026	Objective	Asymmetric PG/F-KL switch	White-box	Token	Advantage-sign token-level switching
<i>Objective: RL-Augmented</i>						
G-OPD (Yang et al., 2026c)	2026	Objective	KL-Constrained RL	White-box	Token/Seq	OPD $\equiv$ dense KL-RL, reward extrapolation
Delta-KD (Cao et al., 2025)	2025	Objective	Base-to-Instruct delta	White-box	Token	Generative Structure Reward Model (GSRM)
KDRL (Xu et al., 2025a)	2025	Objective	Joint KD + RL	White-box	Token/Seq	On-policy KL regularizer during RL
RLAD (Zhang et al., 2026h)	2026	Objective	Trust Region Ratio	White-box	Token/Seq	PPO-style selective teacher following
AlignDistil (Zhang et al., 2025a)	2025	Objective	DPO + Reverse DPO	White-box	Token/Seq	Token-level RLHF via contrastive DPO reward
SuperCorrect (Yang et al., 2025b)	2025	Objective	Template + Cross-DPO	White-box	Sequence	Hierarchical thought scaffolding
SCoRe (Lyu et al., 2025)	2025	Objective	Earliest-error correction	White-box	Sequence	Root-cause error targeting
CMDP-KD (Zimmer et al., 2025)	2025	Objective	Constrained KL	White-box	Sequence	Unaugmented CMDP with hard KL budget
λ -KD (Cai & Yuan, 2026)	2026	Objective	Inverse RL + KL	White-box	Token/Seq	AVRIL-based reward recovery distillation
TSD-KD (Kim & Baek, 2026)	2025	Objective	Indirect + Direct KD	White-box	Sequence	Indirect re-ranking + selective direct KL
REOPOLD (Ko et al., 2026)	2026	Objective	Log-ratio token reward	White-box	Token	Mixture reward clipping + entropy-dynamic sampling
<i>Signal: White-Box Innovations</i>						
DSKD (Zhang et al., 2025b)	2025	Signal	Dual-Space KL	White-box	Token	Cross-vocabulary projection
ULD (Bozard et al., 2025)	2025	Signal	Probability-space W_1	White-box	Token	Wasserstein distance for tokenizer mismatch
Delta-KD (Cao et al., 2025)	2025	Signal	Base-to-Instruct delta	White-box	Token	Signal isolation
PromptKD (Kim et al., 2024)	2024	Signal	Prompt-based elicitation	White-box	Token	Input-side distillation
TAID (Shing et al., 2025)	2025	Signal	Temporal target interpolation	White-box	Token	Progressive target revelation
MiniPLM (Gu et al., 2025)	2025	Signal	Pre-training KD	White-box	Token	OPD during pre-training
MAD-OPD (Wang et al., 2026b)	2026	Signal	Multi-agent debate ensemble	White-box	Token	Confidence-weighted multi-teacher
Veto (Jiang et al., 2026)	2026	Signal	Geometric bridge KL	White-box	Token	Adaptive target reformation
SimCT (Sun et al., 2026)	2026	Signal	Multi-token continuation KL	White-box	Token	Cross-tokenizer supervision recovery
CSD (Kim et al., 2026b)	2025	Signal	Continuous score matching	White-box	Token	Bypasses discrete vocabulary mismatch
AdaKD (Xie et al., 2026)	2025	Signal	Gradient-weighted KL	White-box	Token	Per-token importance from loss gradient
<i>Training Dynamics</i>						
TIP (Xu et al., 2026b)	2026	Dynamics	Soft-OR token weighting	White-box	Token	Entropy $\times$ divergence quadrant scoring
SCOPE (Zheng et al., 2026)	2026	Dynamics	Rollout routing	White-box	Token	Redirect on flawed prefix
SelectKD (Huang et al., 2025)	2025	Dynamics	Confidence filtering	White-box	Token	Top-1 probability threshold
AdaSwitch (Peng et al., 2025)	2025	Dynamics	Exploration/guidance switch	White-box	Token	Binary mode switching
PACED (Xu et al., 2026a)	2026	Dynamics	Beta-kernel curriculum	White-box	Sample	Frontier difficulty sampling
FOPD (Zhang et al., 2026a)	2026	Dynamics	Prefix truncation	White-box	Token	Early-token focus
Lightning-OPD (Wu et al., 2026a)	2026	Dynamics	Offline caching	White-box	Token	4 $\times$ cost reduction
SKD (Xu et al., 2025c)	2025	Dynamics	Draft-verify	White-box	Block	Amortized teacher cost
NPD (Rang et al., 2026)	2026	Dynamics	Async gen-train decoupling	White-box	Token	8.1 $\times$ speedup via Δ -IFD filtering
Prune-OPD (Yang et al., 2026d)	2026	Dynamics	Top-k overlap drift detection	White-box	Token	Dynamic trajectory pruning, 37-68% time reduction
R-OPD (Fu et al., 2026)	2026	Dynamics	Top-p rollout + masking	White-box	Token	Empirical failure mode fixes
SOD (Zhong et al., 2026a)	2026	Dynamics	Step-level divergence reweighting	White-box	Step	Tool-call cascade attenuation
Stable-OPD (Luo et al., 2026)	2026	Dynamics	R-KL + reference div.	White-box	Token	Rollout mixing for stability
TCOD (Wang et al., 2026c)	2026	Dynamics	Temporal curriculum	White-box	Trajectory	Multi-turn KL stability
Uni-OPD (Hou et al., 2026)	2026	Dynamics	Dual-perspective data balancing	White-box (multi-teacher)	Token	Unified LLM+MLLM OPD
SSB (Mitra & Ulukus, 2025)	2025	Dynamics	In-context self-teacher KL	Self (ICL bootstrapped)	Token	Semantic bootstrapping via exemplar supervision
CaOPD (Zhang et al., 2026b)	2026	Dynamics	Calibration-aware OPD	White-box	Token	Diagnoses scaling law of miscalibration
<i>Self-Distillation</i>						
SPIN (Chen et al., 2024)	2024	Self-Play	Self-play DPO	Self	Sequence	Iterative self-vs-reference game
IRIS (Liao et al., 2026)	2026	Self-Play	Interpolative Rényi self-play	Self	Sequence	Unifies SPIN/SPACE/SPH
OPSD (Zhao et al., 2026b)	2026	Self (PI)	KL on GT-conditioned	Self	Token	Ground-truth as privileged info
CRISP (Sang et al., 2026)	2026	Self (PI)	Conciseness PI	Self	Token	57% token reduction, +9% accuracy
GATES (Stein et al., 2026)	2026	Self (PI)	Gated consensus KL	Self	Token	Document as PI, gated supervision
OPSDL (Zhang et al., 2026e)	2026	Self (PI)	Short-ctx $\rightarrow$ long-ctx	Self	Token	Long-context self-distillation
SD-ZERO (He et al., 2026)	2026	Self (EF)	Self-revision preference	Self + Verifier	Sequence	Binary reward $\rightarrow$ dense self-supervision
π -Play (Zhang et al., 2026g)	2026	Self (PI+SP)	Multi-agent co-evolution	Self	Token	Internally generated PI from self-play
SSD (Zhang et al., 2026d)	2026	Self (SP)	Temperature-sampled SFT	Self	Sequence	Minimalist self-play
UniSD (Jin et al., 2026b)	2026	Self (SP)	Multi-axis unified SD	Self (EMA + agreement)	Token	Systematic component analysis
SDPO (Hübner et al., 2026)	2026	Self (EF)	Textual feedback + DPO	Self + Verifier	Sequence	Credit assignment from outcomes
RLTF (Song et al., 2026)	2026	Self (EF)	Critique-conditioned	Self + Verifier	Sequence	Text feedback for self-play
SRPO (Li et al., 2026b)	2026	Self (EF)	GRPO + SDPO routing	Self + Verifier	Sequence	Sample-routed multi-objective
CoPD (Gu et al., 2026)	2026	Self (Mutual)	Co-evolving bidirectional OPD	Mutual (parallel)	Sequence	Multi-capability consolidation
PAINT (Tan & Hong, 2026)	2026	Self (PI)	Partial-solution masking	Self (re-score)	Token	Overlap-adaptive self-distillation
MSD (Qin et al., 2026)	2026	Self (PI)	Cross-lingual safety distill	Self (English PI)	Token	Multilingual safety alignment
GUI-SD (Zhang et al., 2026f)	2026	Self (PI)	Visual spatial PI for GUI	Self (Gaussian mask)	Token	Multimodal GUI grounding
PBSD (Yu et al., 2026)	2026	Self (PI)	Pref.-based reward-reg. KL	Self (ctx-augmented)	Sequence	DPO-style loss, preference margin
TF-OPD (Jeong, 2026)	2026	Self (PI)	Turn-level truncated KL	Self (EMA + outcome hints)	Token	Multi-turn agentic OPD
VISD (Lin et al., 2026)	2026	Self (PI)	Direction-magnitude decouple	Self (structured judge)	Token	Structured PI for video reasoning
π -Distill (Penaloza et al., 2026)	2026	Self (PI)	Unified PI framework	Self	Token	Theoretical foundations for PI self-distill
OPCD (Ye et al., 2026b)	2026	Self (PI)	Context-conditioned R-KL	Self (experience PI)	Token	Experiential knowledge internalization
OEL (Ye et al., 2026a)	2026	Self (PI)	Online experiential learning	Self (trajectory PI)	Token	Continuous post-deployment distillation loop
HDPO (Ding, 2026)	2026	Self (PI)	GRPO + JSD self-distill	Self (GT-conditioned)	Sequence	Cliff-prompt non-zero gradient patch
MTPSD (Kirchensbauer et al., 2026)	2026	Self (SP)	Multi-token span scoring	Self (frozen copy)	Span	Architectural self-distillation for > 3 $\times$ decoding
SDFT (Shenfeld et al., 2026)	2026	Self (SP)	In-context demo teacher R-KL	Self (ICL-conditioned)	Token	Forgetting resistance via trust-region implicit KL
OPSFT (Zhao et al., 2026a)	2026	Self (SP)	Generate-filter-SFT	Self + Verifier	Sequence	Correctness+conciseness filtering, no RL gradient
RLSD (Zhang et al., 2026a)	2026	Self (EF)	RLVR + self-distill decomp.	Self + Verifier	Token	Magnitude (self-distill) + direction (RLVR), 2 $\times$ efficiency
<i>Black-Box</i>						
GAD (Ye et al., 2025)	2025	Signal	Adversarial matching	Black-box	Sequence	Discriminator as proxy reward
Lion (Jiang et al., 2023)	2023	Signal	Verbal feedback	Black-box	Sequence	NL critique curriculum
OVD (Xiong et al., 2026)	2026	Signal	Verbal scores (0-9)	Black-box	Sequence	Score-based trajectory matching
DAIL (Mendes et al., 2026)	2026	Signal	Mixed-policy decoding	Black-box	Sequence	Didactic-to-constructive learning
LUFFY (Yan et al., 2025)	2025	Signal	Mixed-policy GRPO	Off-policy + RL	Sequence	Importance-weighted off-policy
ORPO-Distill (Singh et al., 2025)	2025	Signal	Preference optimization	Black-box	Sequence	Teacher-ranked pairs
PRISM (Wang et al., 2026d)	2026	Signal	Black-box pre-alignment	Black-box (API)	Sequence	Multimodal OPD before RLVR
ROPD (Fang et al., 2026a)	2026	Signal	Rubric-based scoring	Black-box (API)	Sequence	Structured semantic rubrics, 10 $\times$ sample efficiency
ThinkTuning (RRV et al., 2025)	2025	Signal	Corrective hints + GRPO	Black-box (same-size)	Sequence	Interleaved corrective episodes during RL
DASD (Yan et al., 2026)	2026	Signal	Sequence-level distribution match	Black-box	Sequence	Multi-trace distribution alignment
DDT (Zhang et al., 2026c)	2026	Signal	Theoretical RL+KD framework	Black-box	Sequence	On-policy generalization theory

tradeoffs (Agarwal et al., 2024). Multi-step agentic tasks demand sequence-level objectives with trajectory-aware credit assignment (Lyu et al., 2025; Zhang et al., 2026h).

Compute budget. Under constrained compute (<500 GPU-hours), prefix-truncated roll-outs (Zhang et al., 2026a) and offline caching strategies (Wu et al., 2026a) reduce on-policy overhead substantially. At medium scale, full on-policy training with adaptive divergences (ToDi (Jung et al., 2025), AKL (Wu et al., 2025)) becomes feasible. Large-scale deployments (>5000 GPU-hours) typically adopt multi-stage pipelines that begin with off-policy warmup,

Table 3: Representative experimental configurations of OPD methods across different application domains.

Method	Teacher	Student	Training Signal	Key Benchmarks
<i>Mathematical Reasoning</i>				
G-OPD (Yang et al., 2026c)	Qwen3-30B-A3B-Instruct	Qwen3-1.7B/4B	Dense logit + reward extrap.	MATH, GSM8K, Code
OPSD (Zhao et al., 2026b)	Self (GT-conditioned)	Qwen3-1.7B/4B/8B-Inst	GT-FI self-distillation	AME, HMMT
PACED (Xu et al., 2026a)	Qwen3-14B	Qwen3-8B	Best-kernel curriculum	MATH-500, AIME
KDRL (Xu et al., 2025a)	Skywork-OR1-Math-7B	R1-Distill-Qwen-1.5B	Joint KD + RL	AIME, MATH
LUFFY (Yan et al., 2025)	DeepSeek-R1 (off-policy traces)	Qwen2.5-Math-7B/1.5B	Mixed-policy GRPO	AIME, AMC, MATH-500
AOPD (Jia et al., 2026)	Qwen2.5-Math-7B (teacher)	Qwen2.5-Math-1.5B/7B	Asymmetric PG + F-KL	AIME24/25, MATH-500
vOPD (Oh et al., 2026)	Qwen3-1.7B-Inst (self-distill)	Qwen3-1.7B/4B-Base, OLMO-3-7B	CV baseline + single-sample R-KL	MATH-500, AIME24/25, Minerva
<i>Instruction Following & General</i>				
GKD (Agarwal et al., 2024)	T5-XL (3B)	T5-Small/Base	On-policy FKL/RKL/JSD	XSum, WMT
DnatLLM (Ko et al., 2024)	GPT-2-XL (1.5B)	GPT-2 (124M)	Skewed RL	ROUGE-L, Dolly
AlignDistil (Zhang et al., 2025a)	Self (contrastive DPO reward)	Qwen2/2.5-1.5B-Inst	Token-level RLHF distillation	AlpacaEval, MT-Bench
<i>Industrial Scale</i>				
Qwen3 (Yang et al., 2025a)	Qwen3-32B/235B-A22B	Qwen3-0.6B-14B	On-policy logit distillation	AIME, MATH, LiveCode
Gemma 2 (Gemma Team et al., 2024)	Gemma 2 27B	Gemma 2 9B / 2B	KD in pre-training	MMLU, HellaSwag
MMMo-V2 (Xiaomi LLM-Core Team et al., 2026)	Domain specialists	MMMo-V2-Flash (309B/15B MoE)	Multi-teacher logit + reward	AIME, MATH, GPQA
KAT-Coder-V2 (Li et al., 2026a)	5 domain specialists	Unified agentic coder	Specialize-then-Unify OPD	SWE-bench (79.6%)
Nemotron-Cascade 2 (Yang et al., 2026e)	Domain RL/SFT teachers	30B MoE (5B active)	Cascade RL + domain OPD	IMO, IOI, ICPC
DeepSeek-V4 (DeepSeek-AI, 2026)	10+ domain experts (1.6T)	DeepSeek-V4-Pro (1.6T MoE)	Full-vocab multi-teacher R-KL	AIME, LiveCode, HMMT
ORBIT (Liang et al., 2026)	Multi-mode RL experts	DeepSeek-Distill-Qwen-1.5B, Qwen3-4B, Openmath-Nemotron-7B	Multi-teacher mode fusion	AIME, GPQA, MMLU-Pro
<i>Multimodal & Domain-Specific</i>				
SCoRe (Lyu et al., 2025)	Qwen2.5-72B-Inst	Qwen2.5-7B/3B, Llama-3.1-8B	Earliest-error correction	12 Agent benchmarks
VOLD (Boussehjam et al., 2025)	Qwen3-8B (text-only)	Qwen2.5-VL-3B	GRPO + token-level KL	MMLU-Pro, MathVision
CORD (Hu et al., 2026)	Test-mode LLaLM	Audio-mode LLaLM	Cross-modal R-KL + GRPO	Audio QA/reasoning
SKD (Xu et al., 2025c)	Gemma-7B-IT, Qwen2-7B	Gemma-2B, Qwen2-0.5B	Block-verified on-policy KL	Translation, Dialogue
LUD (Boizard et al., 2025)	Llama-2-7B, Mistral-7B	OPT-350M, Pythia-160M-1B	Probability-space IW	QA, Summarization
DP-OPD (Khadem et al., 2026)	GPT-2 Large, 774M	DistilGPT-2-82M	GKD + DP-SCD (student only)	Yelp, BigFate
Qwen2-7B-IT	Qwen2-0.5B-IT	Qwen2-0.5B-IT	Geometric bridge KL	Reasoning, Generation
TAID (Shing et al., 2025)	Phi-3-mini (3.8B), Llama-2-7B	TinyLlama (1.1B)	Temporal interpolation	HellaSwag, PQQA, ARC
GLI-SD (Zhang et al., 2026f)	Self (visual PT)	Qwen3-VL-Inst-8B	Entropy-guided KL	ScreenSpot-v2/Pro, OSWorld
MAD-OPD (Wang et al., 2026b)	Multi-teacher debate	Qwen3/3.5 (1.7B-14B)	Confidence-weighted JSD/R-KL	Agentic, Code
MSD (Qin et al., 2026)	Self (English CoT PT)	Qwen2.5-7B, Llama-3-8B	DFSW cross-lingual distill	Multilingual Safety
NPD (Rang et al., 2026)	Teacher (parallel prefill)	openPangea-Embedder-1B	Async CT+KD + A-FD	GSM8K, MATH (8.1x speedup)
SimCT (Sun et al., 2026)	Qwen2.5-7B-Inst, Phi-4-mini	Phi-4-mini, Gemma-2-2B-IT	Multi-token continuation KL	Math, Code (cross-tokenizer)
Prune-OPD (Yang et al., 2026d)	R1-Distill-Qwen-7B, Qwen3-4B	R1-Distill-Qwen-1.5B, Qwen3-1.7B	Top-k overlap drift pruning	AMC, AIME, HMMT (37-68% time saved)
SDO (Zhong et al., 2026a)	Qwen3-4B	Qwen3-1.7B	Step-level divergence reweighting	AIME25 (26.13% @0.6B), SeCode
LiteGLI (Wu et al., 2026b)	Qwen3-VL-32B	SD 3.5 Medium (consolidated)	Guided OPD + dual-level GRPO	ScreenSpot-Pro, OS-World, Lite-Bench
Flow-OPD (Fang et al., 2026b)	GRPO domain specialists (SD3.5)	SD 3.5 Medium (consolidated)	Two-stage GRPO $\rightarrow$ OPD	GenEval (63 $\rightarrow$ 92), OCR (59 $\rightarrow$ 94)
VED (Lin et al., 2026)	Self (video-aware judge)	VideoLLM	Dir-mag, decoupled advantage + RL	OpenQ1-Video benchmarks
Uni-OPD (Hou et al., 2026)	Multi-teacher (LLM+MLLM)	Qwen3 (1.7B-4B), Qwen3-VL-4B	Dual-perspective OPD	Math, Code, MLLM (16 benchmarks)
<i>Self-Distillation</i>				
SPIN (Chen et al., 2024)	Self (iter 1-1)	Zephyr-7B-SFT	Self-play DPO	Open LLM, MT-Bench
CRSP (Sang et al., 2026)	Self (condse prompt)	Qwen3-8B/14B	Per-token R-KL	AIME, MATH-500
SD-ZERO (He et al., 2026)	Self (revision-conditioned)	Qwen3-4B-Inst, OLMO-3-7B	On-policy self-distill	AIME, MATH, Codeforces
π -Play (Zhang et al., 2026g)	Self (QCP-conditioned)	Qwen3-4B/8B	Multi-agent self-distill	NQ, TriviaQA, HotpotQA
SSD (Zhang et al., 2026d)	Self (temperature-sampled)	Qwen3-30B-Inst	On-policy SFT	LiveCodeBench v6
UnisD (Jin et al., 2026)	Self (EMA + multi-teacher agree.)	Qwen2.5-7B/0.5B-3B, Llama-3.1-8B, Gemma-3-4B	Agreement + EMA + contrastive + feat. match	SciQnQA, MBPP, CoS-E, ToolAlpaca
SDPO (Hubotter et al., 2026)	Self (textual feedback)	Qwen3-8B	Credit-assignment self-distill	Code, Math
SRPO (Li et al., 2026b)	Self (SDPO+GRPO)	Qwen3-8B	Sample-weighted dual-objective	Chem, Phys, Bio, Mat, ToolUse
PRSD (Yu et al., 2026)	Self (cb-augmented ref.)	Qwen3-1.7B/4B/8B-Inst	Prft-based reward-rog. KL	AIME24/25, HMMT25, Tool Use
TT-OPD (Jeong, 2026)	Self (EMA + outcome hints)	Qwen3-5-9B	Turn-level truncated KL	Healthcare AI Gym (18 benchmarks)
Kim & Lee (2026)	Self (GT-conditioned)	Qwen3-8B, R1-Distill-7B, AceReason-7B	Correct-only OPD	MATH-500, AIME24/25
<i>Black-Box</i>				
GAD (Ye et al., 2025)	GPT-5-Chat (API only)	Qwen2.5-14B-Inst	Adversarial minimax	LMSYS-Chat, Dolly
Lion (Jiang et al., 2023)	ChatGPT (API only)	LLaMA-7B/13B	NL critique curriculum	BIG-Bench Hard, AGIEval
OVD (Xiong et al., 2026)	QwQ-32B / Env. Agent	Qwen2.5-3B, Llama-3.2-3B	Verbal scores (0-9)	Web QA, MATH
ROFD (Fang et al., 2026a)	GPT-5.2, Qwen3-30B-A3B	Qwen3-4B, Gemma-3-4B	Rubric-based semantic scoring	AIME24/25, HMMT25, GPQA-D, IFEval

transition to on-policy refinement, and optionally add RL-based exploration, as reported for Qwen3 (Yang et al., 2025a).

Stability requirements. Common failure modes motivate specific stabilization choices. Length inflation during on-policy training can be mitigated by reference divergence constraints and rollout mixing (Luo et al., 2026). The flawed prefix trap, where teacher feedback on student-generated errors produces misleading gradients, is addressed by token-level reliability filtering (Xu et al., 2026b; Zheng et al., 2026). Thinking-pattern mismatch between teacher and student is typically resolved through an off-policy cold-start phase (Li et al., 2026d) before transitioning to on-policy training.

4 Objective Functions and Optimization

The objective function determines what mathematical quantity the student optimizes at each training step, governing gradient geometry, mode coverage, and ultimately the achievable performance ceiling. Methods in this category have progressed from fixed divergences that apply a single metric uniformly, through adaptive divergences that modulate the metric per-token based on local distributional geometry, to RL-augmented objectives that incorporate reward signals beyond the teacher’s output distribution.

4.1 Fixed Divergence Objectives

GKD (Agarwal et al., 2024) established the canonical OPD framework by drawing on the on-policy imitation learning paradigm (Ross et al., 2011) for autoregressive language models. It samples trajectories from a mixture policy $\pi_{\text{mix}} = \lambda p_{\theta} + (1 - \lambda) p_{\text{data}}$ and applies standard f -divergences (Forward KL, Reverse KL, or JSD) at each token position. The mixture coefficient λ interpolates between fully off-policy ($\lambda = 0$) and fully on-policy ($\lambda = 1$), serving as a single control for exposure bias mitigation. On instruction-following and summarization benchmarks, moderate on-policy mixing ($\lambda \geq 0.5$) consistently outperforms pure off-policy training, validating the distributional alignment hypothesis.




Figure 2: Forward KL vs. Reverse KL divergence for fitting a student distribution P_S to a bimodal teacher P_T . **(a)** The teacher places equal probability mass on two valid modes (e.g., two correct answers or two stylistic variants). **(b)** Forward KL minimizes $D_{\text{KL}}(P_T \parallel P_S)$ and is mode-covering (zero-avoiding): the student spreads a single broad Gaussian over both modes and places substantial probability in the inter-mode hallucination zone (red dashed boundary), corresponding to outputs that are a blend of the two valid answers and therefore neither. **(c)** Reverse KL minimizes $D_{\text{KL}}(P_S \parallel P_T)$ and is mode-seeking (zero-forcing): the student concentrates sharply onto a single peak (Mode 2), dropping Mode 1 entirely but producing crisp, coherent generations within the chosen mode. This trade-off motivates the adaptive-divergence methods in §4 (ToDi, AKL, EOPD), which interpolate between the two regimes based on per-token teacher entropy or log-ratio statistics.

Pure KL divergences, though elegant, exhibit numerical instability during on-policy exploration. When the student generates tokens to which the teacher assigns near-zero probability, Forward KL produces unbounded gradients. Reverse KL ignores these regions entirely, creating blind spots in the loss landscape. DistiLLM (Ko et al., 2024) tackles this through *Skewed KL* (SKL), which replaces q_θ in the standard KL with a skewed mixture $\tilde{p} = \alpha p_T + (1 - \alpha)p_\theta$ for $\alpha \in (0, 1]$. The token-level SKL loss is

$$\mathcal{L}_{\text{SKL}} = \mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{mix}}} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_T(\cdot | y_{<t}) \parallel \alpha p_T(\cdot | y_{<t}) + (1 - \alpha)p_\theta(\cdot | y_{<t})) \right] \quad (10)$$

where $\mathcal{D}_{\text{mix}}$ mixes ground-truth sequences with cached student rollouts according to DistiLLM’s adaptive off-policy scheduler. Skewing ensures the target density is bounded below by $(1 - \alpha)p_\theta$, which avoids gradient explosion without requiring RL policy gradients. DistiLLM further proposes *Skewed Reverse KL* (SRKL) for mode-seeking applications, creating a symmetric pair of stabilized divergences. Its successor, DistiLLM-2 (Ko et al., 2025), recognizes that teacher-generated and student-generated data carry qualitatively different learning signals and applies *asymmetric* objectives, using Forward SKL on teacher-generated data (where mode-covering teaches the student to explore the teacher’s full distribution) and Reverse SRKL on student-generated data (where mode-seeking teaches the student to sharpen its own best responses). This source-aware asymmetry produces consistent improvements over symmetric single-divergence baselines across instruction following and summarization tasks, consistent with the view that the appropriate divergence is data-source-dependent.

DistiLLM-2 applies different divergences to different data sources, but operates at a coarse granularity (entire sequences grouped by source). The fixed-divergence family faces a more fundamental limitation rooted in the heterogeneity of autoregressive token distributions *within* a single sequence. The choice between Forward KL (mode-covering) and Reverse KL (mode-seeking) is a global decision applied uniformly across all token positions (Figure 2). The teacher’s distributional structure, however, varies considerably within a single sequence. At a mathematical operator token, the teacher concentrates probability mass on one or two correct symbols, so mode-seeking is appropriate because the student must commit to

the correct answer. At a conjunction or filler word, dozens of tokens share comparable probability, so mode-covering is appropriate because any synonym suffices. Applying a single divergence to both positions therefore wastes gradient signal in one regime or discards coverage information in the other. This per-position mismatch between the geometry of the loss and the geometry of the target distribution motivates the adaptive methods of Section 4.2.

Sequence-level objectives offer a complementary perspective. MiniLLM (Gu et al., 2024) elevates Reverse KL to the full sequence level, where the student p_θ appears inside the sampling expectation. The central challenge is computing the gradient of an expectation over p_θ when p_θ itself is parameterized by θ . Starting from $D_{\text{KL}}(p_\theta \parallel p_T) = \mathbb{E}_{y \sim p_\theta} [\log p_\theta(y|x) - \log p_T(y|x)]$ and applying the policy gradient theorem, MiniLLM derives (Eq. 2 of Gu et al., 2024):

$$\nabla_\theta \mathcal{L}_{\text{MiniLLM}} = -\mathbb{E}_{y \sim p_\theta} \left[\sum_{t=1}^{|y|} (R_t - 1) \nabla_\theta \log p_\theta(y_t | y_{<t}) \right], \quad R_t = \sum_{t'=t}^{|y|} \log \frac{p_T(y_{t'} | y_{<t'})}{p_\theta(y_{t'} | y_{<t'})} \quad (11)$$

where the -1 arises from the $-\log p_\theta(y|x)$ entropy term of the reverse-KL objective. Minimizing sequence-level Reverse KL is thus mathematically equivalent to policy gradient RL with per-step reward $r_t = \log[p_T(y_t | y_{<t}) / p_\theta(y_t | y_{<t})]$ and an entropy-induced constant. MiniLLM further decomposes R_t into a single-step quality term and a future-return term (Eq. 3 of Gu et al., 2024), computing $\mathbb{E}_{y_t \sim q_\theta}[r_t]$ in closed form over the vocabulary instead of by Monte Carlo sampling. This decomposition reduces variance by isolating the dominant source of gradient noise (single-step quality) from the accumulated future return, and supports credit assignment at the token level. Per-step returns R_t attribute the trajectory’s deviation from the teacher to specific token decisions rather than spreading uniform blame across the sequence, placing the optimization within the RL framework with the teacher’s log-probability as a dense reward. This connection to policy optimization is consistent with MiniLLM’s empirical strength on reasoning tasks where sequence-level coherence matters more than token-level accuracy, at the cost of higher training compute because REINFORCE estimation in combinatorial output spaces requires more iterations and careful baseline subtraction to converge. KETCHUP (Fan et al., 2026) addresses this variance problem by replacing the full cumulative return G_t with a K -step Bellman return estimate derived from the optimality equation, reducing gradient variance at a power rate in K while preserving the sequence-level objective (Theorem 1 of Fan et al., 2026). CMDP-KD (Zimmer et al., 2025) takes a complementary approach, formulating distillation as reward maximization subject to a hard KL budget via constrained MDP. Their key insight is that the state-augmented formulation of Saute can be simplified for LLMs by removing state augmentation entirely, since the full history already determines the remaining budget, yielding constraint satisfaction guarantees without requiring teacher access at inference.

The token-level vs. sequence-level tradeoff. The choice between token-level and sequence-level objectives represents a core bias-variance tradeoff whose resolution foreshadows the RL-augmented methods of §4.3. Token-level methods (GKD, DistiLLM) yield $|V|$ -dimensional gradient signals at every position, yielding low variance but potentially biased gradients when the teacher’s token-level distribution is poorly calibrated (the flawed prefix trap of Section 6.1). Sequence-level methods (MiniLLM, KETCHUP) yield unbiased gradients through the policy gradient theorem but suffer from high variance due to single-sample REINFORCE estimation over exponentially large sequence spaces. The practical resolution is task-dependent. For instruction following where token-level teacher quality is high, GKD-style methods tend to dominate. For reasoning tasks where sequence-level coherence matters more, MiniLLM-style methods are often preferred despite their variance, and KETCHUP’s K -step Bellman returns offer an intermediate bias-variance point on this spectrum. Recognizing that MiniLLM is already performing policy gradient RL (with the teacher’s log-probability as reward) suggests that sequence-level distillation and RL-augmented distillation are not distinct paradigms but endpoints of a continuum parameterized by the reward source: pure teacher signal at one end, pure task reward at the other, and hybrid objectives (G-OPD, RLKD) interpolating between them.

Full-vocabulary vs. sampled-token KL. Orthogonal to the token-level vs. sequence-level choice, a practical implementation decision determines how much of the teacher’s distri-

bution is used at each position. Many RL-based OPD implementations approximate the per-position KL by computing $\log p_\theta(y_t|y_{<t}) - \log p_T(y_t|y_{<t})$ only at the *sampled* token y_t , treating this scalar as a per-token advantage within a policy gradient framework (Lu & Thinking Machines Lab, 2025). While computationally lightweight (requiring only one teacher log-probability per position), this single-sample estimate incurs high gradient variance and discards the distributional information across non-sampled tokens. DeepSeek-V4 (DeepSeek-AI, 2026) reports that *full-vocabulary* logit distillation, computing the exact $D_{\text{KL}}(p_\theta(\cdot|y_{<t}) \parallel p_T(\cdot|y_{<t}))$ over all $|V|$ tokens at each position, yields more stable gradients and more faithful knowledge transfer at the cost of transmitting or reconstructing the full teacher logit vector. Their engineering solution (caching teacher hidden states and reconstructing logits on-the-fly through the prediction head) makes this feasible even with 10+ trillion-parameter teachers over vocabularies exceeding 100K tokens, which positions full-vocabulary KL as an attractive choice for industrial-scale OPD when infrastructure permits. Lightning-OPD (Wu et al., 2026a) further reports that the full teacher forward pass can be largely avoided by precomputing teacher log-probabilities over SFT rollouts, recording $4\times$ speedup over standard OPD under a *teacher consistency* condition (using the same teacher for both SFT data generation and distillation).

vOPD (Oh et al., 2026) offers a third resolution to this variance-cost tradeoff by importing the *control variate baseline* from policy-gradient RL into the OPD estimator. Casting the single-sample OPD gradient (Eq. 3 of Oh et al., 2026) as a REINFORCE update with per-token reward $r_t = \log p_T(y_t|c_t) - \log p_\theta(y_t|c_t)$, vOPD subtracts a baseline $b_t(c_t)$ from r_t before multiplying by $\nabla_\theta \log p_\theta$. The key insight is that the natural value function for this reward admits a closed form as the per-token negative reverse KL between student and teacher, $V^{\pi_\theta}(c_t) = -D_{\text{KL}}(p_\theta(\cdot|c_t) \parallel p_T(\cdot|c_t))$, computable directly from the forward pass that already produces the OPD objective without requiring an auxiliary critic or additional rollouts. Because the baseline depends only on the context c_t and not on the sampled token y_t , subtracting it preserves the unbiasedness of the gradient while reducing variance. A top- k approximation of this baseline (restricting the KL summation to the student’s k most probable tokens) further lowers compute, and because the approximation remains independent of the sampled action, unbiasedness still holds regardless of k . Across six reasoning benchmarks on Qwen3-1.7B and 4B, vOPD yields an average +3% absolute gain over vanilla single-sample OPD (up to +6.2% on MATH500), matches the accuracy of full-vocabulary OPD while reducing wall-clock time by up to 57.7%, and substantially outperforms top- k OPD, which adds compute but introduces bias. The result positions vOPD at a favorable point on the bias-variance-compute Pareto frontier for on-policy distillation of reasoning models.

4.2 Adaptive Divergence Objectives

Adaptive methods tackle this per-position mismatch by selecting or interpolating divergences based on local distributional geometry. When the teacher is certain (low entropy), mode-seeking Reverse KL preserves precision, and when uncertain (high entropy), mode-covering Forward KL preserves diversity.

ToDi (Jung et al., 2025) implements this insight through per-token, per-vocabulary-entry weighting derived from the teacher-student probability log-ratio. For each vocabulary entry v_i at position t , a sigmoid-based weight $\omega_{t,i} = \sigma\left(\text{sg}\left[\log \frac{p(v_i|y_{<t},x)}{q_\theta(v_i|y_{<t},x)}\right]\right)$ blends Forward KL and Reverse KL at the finest possible granularity:

$$\mathcal{L}_{\text{ToDi}} = \sum_{t=1}^{|y|} \sum_{i=1}^{|V|} \omega_{t,i} \cdot D_{\text{FKL}}^{(t,i)}(p, q_\theta) + (1 - \omega_{t,i}) \cdot D_{\text{RKL}}^{(t,i)}(p, q_\theta) \quad (12)$$

where $\text{sg}[\cdot]$ denotes stop-gradient. The weight $\omega_{t,i}$ exceeds 0.5 (emphasizing Forward KL) when the teacher assigns higher probability than the student to v_i , boosting underestimated tokens, and falls below 0.5 (emphasizing Reverse KL) when the student overestimates, suppressing excess probability. This per-vocabulary-entry granularity is the defining distinction from AKL, which applies a single interpolation weight per position. By maintaining independent weights across the full vocabulary at each step, ToDi can simultaneously boost some

tokens and suppress others within the same position, capturing fine-grained distributional mismatches that position-level methods miss. EOPD (Jin et al., 2026a) takes an alternative approach, using the teacher’s per-token entropy $H(p_T(\cdot|y_{<t}))$ as a gating signal. The base objective is Reverse KL at all positions, with Forward KL *additively* activated at high-entropy tokens where the teacher distributes probability across multiple valid continuations:

$$\mathcal{L}_{\text{EA}} = \mathbb{E}_{y \sim p_\theta} \left[\sum_{t=1}^{|y|} D_{\text{KL}}(p_\theta \parallel p_T) + \mathbb{I}[H_t > \tau] \cdot D_{\text{KL}}(p_T \parallel p_\theta) \right] \quad (13)$$

Both methods consistently outperform any fixed divergence on math reasoning (MATH, GSM8K) and open-ended generation (AlpacaEval), consistent with the view that the appropriate divergence is position-dependent. The complementarity between these two methods is instructive. ToDi conditions adaptation on the teacher-student gap (a pairwise property), while EOPD conditions on teacher uncertainty alone (a marginal property). The former adapts more aggressively when the student is far from the teacher, while the latter delivers a stable routing signal independent of student quality, making it more suitable for early training when the student’s distribution is less reliable.

AKL (Wu et al., 2025) synthesizes elements of both approaches, dynamically interpolating between Forward KL and Reverse KL based on a head/tail vocabulary partition. At each position, it splits tokens into head (high-probability) and tail sets by a threshold μ , computes the cumulative absolute gap $g_{\text{head}} = \sum_{v \in V_h} |p_T(v) - p_\theta(v)|$ (and similarly g_{tail}), and weights Forward KL by $g_{\text{head}} / (g_{\text{head}} + g_{\text{tail}})$ and Reverse KL by $g_{\text{tail}} / (g_{\text{head}} + g_{\text{tail}})$. Unlike entropy-based routing, which depends only on the teacher’s distributional shape, AKL considers the *relative* geometry of teacher and student via partition-level discrepancy. All three adaptive methods admit a natural interpretation through information geometry. The Fisher information metric on the statistical manifold of autoregressive distributions varies across token positions. At reasoning-critical positions, the manifold is highly curved (the teacher concentrates probability on few tokens), making Reverse KL’s mode-seeking behavior locally appropriate because it follows the manifold’s natural gradient. At generation-flexible positions, the manifold is nearly flat (many synonyms are equally valid), and Forward KL’s mode-covering behavior preserves useful entropy. Adaptive divergences implicitly navigate this varying curvature by selecting the locally best metric, which explains their consistent empirical superiority over any fixed choice.

AOPD (Jia et al., 2026) takes the per-token adaptive principle to its logical conclusion by diagnosing *why* standard OPD gradients fail at specific token positions. In non-positive advantage regions ($A_t \leq 0$), standard advantage-weighted policy gradients exhibit three pathologies that intensify under weak initialization. Heavy-tailed variance from unbounded negative advantages destabilizes updates, vanishing gradients at zero-advantage positions waste teacher signal, and “exploration black holes” trap learning because suppressing a wrong token only redistributes mass proportional to the student’s own prior, unable to boost low-probability correct alternatives. AOPD resolves this by switching the objective at the token level. For positive-advantage tokens ($A_t > 0$), standard policy gradient exploitation is preserved. For non-positive tokens, the framework substitutes localized forward-KL guidance on the teacher’s top- K support, creating an asymmetric exploitation-imitation mechanism. A threshold parameter τ (set to zero by default, corresponding to the probability difference between teacher and student) determines the switching point, with $\tau = -1$ recovering standard OPD (intervention never triggers because $P_T - P_S \leq -1$ is almost never satisfied) and $\tau = 1$ recovering GKD (intervention everywhere because $P_T - P_S \leq 1$ always holds). On competition-level mathematics (AIME 2024/2025, MATH-500), AOPD attains average gains of +4.09 with strong initialization and +8.34 with weak initialization over standard on-policy baselines, with the weak-initialization gap suggesting that the exploration black hole is among the dominant failure modes when the student starts far from the teacher. The relationship between AOPD and the earlier adaptive methods is one of complementary scope. ToDi, AKL, and EOPD adapt the divergence *type* based on distributional geometry (teacher entropy, teacher-student gap), while AOPD adapts the *objective class itself* (policy gradient vs. divergence minimization) based on the gradient’s structural health at each position.

The field has moved noticeably from “which divergence?” to “which divergence *where*?” The evidence from ToDi, AKL, EOPD, and AOPD points to a consistent finding across evaluated settings, namely that no single fixed divergence dominates uniformly, and that even the simplest adaptive routing (a hard entropy threshold) often outperforms the best-tuned fixed alternative. Yet adaptive divergences, however sophisticated their per-token routing, share a structural limitation with their fixed counterparts. They optimize a distance to the teacher distribution. Perfect optimization under any f -divergence yields a student that matches the teacher but rarely exceeds it. When the teacher itself makes errors or lacks reasoning paths that a well-guided student could discover, divergence minimization can act more as a ceiling than as a floor.

4.3 RL-Augmented Objectives

Divergence matching, whether fixed or adaptive, tends to bound the student near the teacher’s capability ceiling. In the limit of perfect optimization, the student matches the teacher’s distribution and rarely exceeds it. RL-augmented objectives target this limitation by injecting an external performance signal that can steer the student toward solutions the teacher itself would be unlikely to produce.

OPD as KL-constrained RL. G-OPD (Yang et al., 2026c) formalizes an equivalence between standard OPD and dense KL-constrained reinforcement learning. It formulates distillation as maximizing a token-level reward defined as the log-probability ratio between teacher and reference policy, penalized by KL divergence from the reference:

$$\max_{\theta} \mathbb{E}_{y \sim p_{\theta}} \left[\sum_{t=1}^{|y|} \alpha \log \frac{p_T(y_t | y_{<t})}{p_{\text{ref}}(y_t | y_{<t})} - D_{\text{KL}}(p_{\theta}(\cdot | y_{<t}) \parallel p_{\text{ref}}(\cdot | y_{<t})) \right] \quad (14)$$

When $\alpha = 1$, this reduces to standard Reverse KL distillation. When $\alpha > 1$, the objective forces the student to *extrapolate* beyond the teacher’s probability mass, allowing a phenomenon documented as *Reward Extrapolation*, whereby the student discovers previously unseen, structurally sound reasoning paths in regions where the teacher’s generation probability $p_T(y|x)$ is low but the outcome reward remains high. In the multi-teacher setting where several domain-specific RL experts share the same base model, reward extrapolation (ExOPD) produces a unified student that surpasses all its same-size domain teachers, showing that OPD need not be a lossy compression step.

This equivalence has a deeper theoretical consequence that unifies the entire section. It bridges the distillation and RL communities by showing that all objectives discussed here (GKD’s token-level KL, MiniLLM’s sequence-level Reverse KL, G-OPD’s reward-augmented formulation) are special cases of KL-constrained policy optimization with different reward definitions and constraint strengths. Any advance in KL-constrained RL (better trust regions, adaptive penalty coefficients, variance reduction) immediately transfers to OPD, and vice versa.

Gradient decomposition and the KD+RL complementarity. The unified KD+RL framework of Li et al. (2025) decomposes the hybrid gradient into two complementary components, a dense KD gradient for token-level imitation and a Monte Carlo RL gradient for reward maximization. The KD component dampens the high variance inherent in policy gradient estimation, while the RL component prevents the student from collapsing onto teacher modes that are suboptimal for the downstream task. REOPOLD (Ko et al., 2026) instantiates this insight by interpreting on-policy distillation as policy optimization where the teacher-student log-likelihood ratio acts as a token reward. It stabilizes training through three mechanisms: mixture-based reward clipping to prevent over-trust of extreme teacher signals, entropy-based dynamic sampling to selectively leverage teacher feedback at high-uncertainty tokens, and a unified exploration-to-refinement schedule. Empirically, REOPOLD attains $6.7\text{--}12\times$ greater sample efficiency than recent RL approaches and allows a 7B student to match a 32B teacher in visual reasoning with $\sim 3.3\times$ inference speedup.

Structure-aware reward as an alternative to dense KL. While G-OPD and REOPOLD retain token-level KL as the primary distillation signal (differing only in how they augment or stabilize it), RLKD (Xu et al., 2025b) departs from this template entirely by replacing

token-level KL with a Generative Structure Reward Model (GSRM). The GSRM parses both teacher and student reasoning paths into sequences of meta-reasoning and solving steps, then scores structural alignment between the two sequences. This structural reward is combined with a task-level outcome reward (e.g., math accuracy) and optimized via GRPO:

$$R_{\text{total}}(x, y) = \alpha R_{\text{GSRM}}(y, y^{\text{teach}}) + (1 - \alpha) R_{\text{outcome}}(x, y) \quad (15)$$

The main idea is that reasoning distillation need not operate at the token level. By supervising at the *step structure* level, RLKD avoids the vocabulary-coupling assumption of logit-based methods while still transferring the teacher’s multi-branch reasoning strategy. Empirically, RLKD trained on only 0.1% of the data under a pure RL regime surpasses standard SFT-RL pipelines, showing that structural reward can substitute for dense KL when the teacher’s value lies in its reasoning organization rather than its token-level distribution. This progression from dense token-level KL (G-OPD) through variance-reduced token rewards (REOPOLD) to step-level structural matching (RLKD) traces an increasing abstraction of what “teacher knowledge” means, moving from distributional fidelity to behavioral organization.

Joint vs. sequential optimization. Given that KD and RL supply complementary gradient signals (dense imitation vs. sparse reward maximization), a persistent engineering challenge is *how* to combine them within a single training pipeline without mutual interference. Sequential approaches (distill-then-RL or RL-then-distill) suffer from catastrophic interference, as the RL stage erases distilled knowledge through policy drift, while the KD stage suppresses previously unseen solutions discovered through exploration. KDRL (Xu et al., 2025a) implements joint optimization by adding an on-policy KL regularizer between student and teacher *during* RL training, preventing policy drift while still allowing reward-driven exploration. RLAD (Zhang et al., 2026h) refines this further by recognizing that unconditional KL regularization is itself suboptimal. Its Trust Region Ratio Distillation (TRRD) replaces the standard KL penalty with a PPO-style likelihood-ratio objective anchored to a teacher-old-policy mixture. This selective strategy follows the teacher only when its signal improves the policy update, preventing blind KL minimization from dragging the student toward suboptimal trajectories on problems the teacher itself cannot solve. The progression from KDRL’s uniform regularization to RLAD’s selective trust region mirrors the fixed-to-adaptive evolution in divergence objectives (Section 4.1→4.2), echoing a recurring pattern across this survey that conditioning the learning signal on local geometry often works better than applying it uniformly.

Preference-based objectives as implicit RL. When logit access is unavailable or when sequence-level quality judgments are more natural than token-level distributions, several methods map the KD+RL problem onto Direct Preference Optimization (DPO). AlignDistil (Zhang et al., 2025a) derives token-level rewards from contrastive DPO and reverse DPO models, converting sequence-level RLHF into per-token policy distillation that applies both standard and reverse preference objectives at the token granularity. OVD (Xiong et al., 2026) queries the teacher for pairwise preferences over student-generated responses, avoiding continuous probability extraction entirely. These preference-based objectives share a deep connection to token-level KL methods. Under DPO’s closed-form solution, the optimal policy satisfies $\pi^*(y|x) \propto \pi_{\text{ref}}(y|x) \exp(r(y, x)/\beta)$, and substituting the teacher’s log-probability as the implicit reward yields the geometric mixture distribution targeted by GKD with λ -mixing. The practical implication is that DPO-style methods are most useful when constructing preference pairs is cheaper than computing full teacher logits. PBSD (Yu et al., 2026) pushes this connection further in the self-distillation regime by deriving the reward-regularized optimum $\pi^* \propto \pi_{\text{teach}} \exp(r/\beta)$ against a context-augmented teacher rather than a static reference policy, and showing that this reward-tilted target provably improves over the teacher itself under the same objective whenever the reward is non-degenerate (Section 5.3.1). The spectrum exposes a unifying gradient. Token-level KL (G-OPD, REOPOLD), step-level structure matching (RLKD), and sequence-level preference ranking (AlignDistil, OVD) all implement KL-constrained policy optimization at different abstraction levels, with the appropriate level determined by signal availability and the granularity at which teacher expertise manifests.

Fine-grained credit assignment. All methods above, whether operating at the token, step, or sequence level, share one limitation: they apply supervision uniformly across the student’s trajectory. When the student makes a single error in step 3 of a 10-step proof, wholesale imitation retrains all 10 steps, diluting the corrective gradient on the actual failure point. SuperCorrect (Yang et al., 2025b) presents hierarchical thought templates with cross-model DPO for error localization, identifying which reasoning step diverged and concentrating feedback there. SCoRe (Lyu et al., 2025) pushes targeting to its logical extreme, correcting *only the earliest error* in the student’s trajectory, concentrating all supervision at the single point of maximum leverage. This exploits the cascading nature of reasoning errors, where correcting the root cause can remove downstream consequences. The contrast with RLKD’s structural approach is instructive. RLKD rewards trajectories whose overall step organization matches the teacher’s, while SCoRe surgically corrects the first point of divergence. The two are complementary, and combining global structural reward with local error correction remains unexplored.

From signal matching to environment reconstruction. The most radical extension asks whether the student can recover the teacher’s *underlying optimization landscape*. $\mathcal{X}$ -KD (Cai & Yuan, 2026) adopts the AVRIL framework to jointly model the teacher’s latent reward function and perform policy distillation, encouraging the student to optimize in a reconstructed version of the teacher’s training environment instead of mimicking its outputs. TSD-KD (Kim & Baek, 2026) combines indirect distillation (student generates candidates that the teacher re-ranks via preference feedback) with direct distillation (KL matching applied selectively to high-entropy tokens where teacher-student gaps are largest). These methods represent the current frontier, moving from matching the teacher’s *behavior* to understanding its *incentives*.

Across objective types, a rough progression is visible. Fixed divergences deliver stable, well-understood optimization but tend to be bounded by the teacher’s capability ceiling. Adaptive divergences can improve efficiency by matching the local geometry of the loss landscape. RL-augmented objectives offer a path around the teacher ceiling at the cost of added variance and instability. The appropriate choice depends on the specific application. For instruction following where teacher quality is high, fixed or adaptive divergences are often sufficient. For reasoning tasks where the student is expected to exceed the teacher, RL augmentation becomes attractive. For tasks where the teacher’s training process is accessible, environment reconstruction offers a more complete form of knowledge transfer.

5 Signal Source and Teacher Architecture

With the objective function fixed (Section 4), the next design axis concerns *where the teacher signal comes from*: full logit access (white-box), API-only text outputs (black-box), or the model’s own internal asymmetries (self-distillation). Each signal source imposes distinct constraints on the applicable objectives and achievable performance.

Historically, signal sources have evolved in step with model access constraints. Early OPD methods (GKD, DistiLLM) assumed full white-box access to model internals, a requirement that limits distillation to same-organization deployments. The rise of proprietary API-only models (GPT-4, Claude) necessitated black-box methods that can extract knowledge from text outputs alone. Most recently, self-distillation methods eliminate the external teacher entirely, permitting continuous improvement without access to any stronger model. This progression from full access through limited access to no external access represents increasing autonomy at the cost of signal density, a trade-off between autonomy and signal density that recurs throughout the design space.

Complex reasoning is among the settings where OPD shows its largest gains over off-policy alternatives. The reason is structural, because reasoning is *path-dependent* (Agarwal et al., 2024). A proof that begins “assume for contradiction...” requires entirely different subsequent tokens than one beginning “we proceed by induction...” If the student has never generated the inductive approach during training, it is unlikely to recover at inference time when that path would be optimal. Off-policy training over a fixed dataset of teacher reasoning traces only covers a finite set of solution paths, while the student needs to navigate

its own terrain of partial solutions, dead ends, and recovery strategies. This observation motivates the on-policy formulation for reasoning transfer. The student generates its own intermediate reasoning steps $r^S \sim P_\theta(\cdot|x)$, and the teacher evaluates these student-generated chains instead of dictating its own:

$$\mathcal{L}_{\text{CoT-OPD}}(\theta) = \mathbb{E}_{x \sim \mathcal{D}, r^S \sim P_\theta(\cdot|x)} \left[\sum_{t=1}^{|r^S|} D_{\text{KL}} \left(p_T(\cdot|x, r_{<t}^S) \parallel p_\theta(\cdot|x, r_{<t}^S) \right) \right] \quad (16)$$

The Forward KL direction encourages mode-covering, so that the student allocates probability wherever the teacher does on the student’s own prefixes, receiving corrective feedback on its actual generation patterns instead of memorizing the teacher’s preferred trajectories. This self-consistent learning dynamic, where the student practices its own reasoning style under expert supervision, has made on-policy CoT distillation a common choice for transferring multi-step reasoning to smaller models. The principle that rationale transfer outperforms answer-only supervision was first established in the off-policy setting by Hsieh et al. (2023), and subsequent work showed that generating rationales on-policy (Agarwal et al., 2024) can yield further gains by reducing exposure bias, an effect that has also been observed at scale in DeepSeek-R1 (DeepSeek-AI et al., 2025) and its successors.

5.1 White-Box Logit Supervision

White-box access yields the densest possible signal, namely the full probability distribution across $|V|$ tokens at every decoding step. This exposes the teacher’s implicit decision boundaries, including the relative probabilities among incorrect tokens (Hinton’s “dark knowledge”), which carry information about similarity structure that binary labels cannot capture.

5.1.1 Same-Family Distillation

When teacher and student share the same tokenizer (as within a single model family), white-box logit transfer is straightforward and all methods can directly compute per-token KL divergence without alignment overhead.

Standard logit distillation. The majority of methods surveyed in Section 4 assume white-box access. GKD, DistiLLM, ToDi, G-OPD, and their variants all require computing $p_T(\cdot|x, y_{<t})$ at each student-generated token position. The computational cost is dominated by the teacher forward pass, executed once per student-generated token (or once per sequence when caching is enabled).

Token-level adaptive supervision. Within the white-box setting, several methods refine *which* tokens receive supervision, and they can be organized along the axis of what drives the selection decision. SelecTKD (Huang et al., 2025) routes based on student-teacher agreement, applying a propose-and-verify mechanism where the student proposes tokens that the teacher verifies, and only positions where agreement holds receive full loss while rejected tokens are masked or down-weighted. AdaSwitch (Peng et al., 2025) routes based on student prefix reliability, maintaining a running error estimate and switching between exploration (student generates freely) and guidance (teacher supplies dense supervision) so that guidance kicks in only when the trajectory has drifted far enough to warrant teacher intervention. AdaKD (Xie et al., 2026) routes based on gradient magnitude, computing per-token importance weights from the KL loss itself and concentrating compute where the teacher-student gap is largest. The three selection criteria, agreement, prefix quality, and gradient magnitude, are complementary rather than competing, each capturing a different signal about where supervision is most informative. They also compose naturally with entropy-based divergence routing (ToDi, EOPD), which selects *which loss function* applies at each position while token-level selection determines *how much gradient signal* propagates. The combination yields an orthogonal two-dimensional design space that current methods have only begun to explore jointly. These token-selection mechanisms complement the sample-level and token-level weighting schemes discussed in Section 6.

The temporal scope of white-box distillation is not limited to post-training. MiniPLM (Gu et al., 2025) extends on-policy distillation to the *pre-training* stage, showing that distilling during pre-training yields stronger foundations for downstream fine-tuning. This represents a qualitative shift, because instead of being a post-hoc compression step, OPD becomes an integral part of the model creation pipeline. Gemma 2 (Gemma Team et al., 2024) validates this approach at industrial scale, embedding KD directly into the pre-training loop for its 9B and 2B models, which are distilled in parallel from a larger 27B teacher. The success of pre-training distillation raises a deeper question, namely whether OPD should be viewed as a separate post-processing step at all, or as an intrinsic component of the training objective alongside next-token prediction.

5.1.2 Cross-Family Distillation

The methods above assume that teacher and student share a common tokenizer, which holds when both belong to the same model family (e.g., Qwen3-30B→Qwen3-4B). A critical challenge arises when teacher and student use different tokenizers or vocabulary sizes, as is common in cross-family transfer (e.g., Llama teacher→Qwen student). Direct KL divergence over mismatched vocabularies is undefined. DSKD (Zhang et al., 2025b) addresses this through *dual-space* knowledge distillation, employing projectors that map between teacher and student representation spaces:

$$\mathcal{L}_{\text{DSKD}} = D_{\text{KL}}(P_{T \rightarrow S} \parallel P_S) + D_{\text{KL}}(P_T \parallel P_{S \rightarrow T}) \quad (17)$$

where $P_{T \rightarrow S}$ denotes the teacher’s hidden states projected into the student’s space and decoded through the student’s prediction head. An exact token alignment algorithm handles vocabulary mismatches by identifying identical tokens across tokenizers. Vocabulary tokens with similar semantic embeddings should receive similar probability mass even when the tokenization schemes differ, allowing distillation across different architectures (e.g., Llama teacher to Qwen student) that share no vocabulary overlap.

ULD (Boizard et al., 2025) takes an alternative approach based on optimal transport, operating directly in probability space rather than hidden-state space. The method augments the standard cross-entropy loss with a Wasserstein-1 distance term that measures the cost of transporting probability mass between the teacher’s and student’s token distributions at each time step:

$$\mathcal{L}_{\text{ULD}} = \sum_{t=1}^{|x|} \text{CE}(t) + \lambda \cdot W_1(p_{\theta_S}(\cdot | x_{<t}^S), q_{\theta_T}(\cdot | x_{<t}^T)) \quad (18)$$

where λ controls the trade-off between the two terms. Under a uniform transport cost assumption (all inter-token costs equal to 1), the W_1 distance admits a closed-form solution that sorts both probability vectors in decreasing order and computes the sum of absolute differences between corresponding entries, yielding $O(n \log n)$ complexity instead of the $O(n^3 \log n)$ of general optimal transport. This closed-form avoids any learnable projection matrix or coupling optimization, making the method lightweight and applicable to arbitrary teacher-student vocabulary pairs.

The distinction between hidden-state alignment (DSKD) and probability-space alignment (ULD) exposes a central design trade-off. Hidden-state alignment preserves richer representational structure but requires maintaining two projectors during training. Probability-space alignment is more lightweight and operates without learnable alignment parameters, but at a level where information has already been compressed through the vocabulary bottleneck. DSKD’s dual projectors must be co-trained with the student because the student’s representation space shifts during on-policy training as its policy evolves, a form of representation drift absent in off-policy settings. CSD (Kim et al., 2026b) offers another solution, distilling via a continuous score function that bypasses the discrete vocabulary mismatch entirely.

SimCT (Sun et al., 2026) takes a complementary approach that operates entirely within the OPD loss form itself. Rather than projecting representations or transporting probability mass, SimCT enlarges the supervision space by comparing teacher and student over short multi-token continuation units that both tokenizers can realize. These units represent the

finest jointly tokenizable supervision interface, recovering teacher signal that exact shared-token matching silently discards at positions where vocabularies disagree. Across three heterogeneous teacher-student pairs (Qwen2.5-7B→Phi-4-mini, Qwen2.5-7B→Gemma-2-2B, Phi-4-mini→Gemma-2-2B) on math reasoning and code generation, SimCT consistently outperforms both shared-vocabulary OPD and coarser cross-tokenizer baselines, with the largest gains at positions of maximal vocabulary disagreement. The result suggests that a substantial fraction of the cross-tokenizer supervision loss documented in DSKD and ULD stems not from inherent incompatibility but from an unnecessarily restrictive matching granularity.

Together, these methods move OPD toward being more architecture-agnostic, allowing organizations to distill from teachers in one model family (e.g., Llama-based) into students in another (e.g., Qwen-based) without the historically restrictive requirement of shared vocabulary, and supporting heterogeneous teacher-student pairs that were previously difficult to handle.

Beyond selecting *which tokens* to supervise, a distinct question is *what information* the teacher should convey at each token. Not all content in the teacher’s logits is equally useful. Delta-KD (Cao et al., 2025) formalizes the principle of *signal isolation*, distilling only the differential shift between a base model and its instruction-tuned variant. Since the base model’s distribution captures generic language modeling knowledge the student already possesses through pre-training, distilling it is redundant at best and harmful at worst because it wastes gradient budget on information the student has already internalized. By extracting only the instruction-following delta, Delta-KD delivers a cleaner supervisory signal that transfers alignment without overwriting the student’s foundation.

A related but distinct challenge is teacher *overconfidence*. When the teacher assigns extreme probability to tokens outside the student’s reachable space, the resulting optimization targets are simultaneously precise and useless. Two solutions tackle this from opposite directions. Veto (Jang et al., 2026) (Adaptive Target Reformulation) refines the teacher’s output *post-hoc*, constructing a geometric bridge in logit space with a parameter β that serves as both an Adaptive Gradient Veto (suppressing pathological gradients on low-confidence tokens) and a Decisiveness Knob (balancing performance with diversity), dynamically reformulating the teacher’s target distribution based on the student’s current capability. PromptKD (Kim et al., 2024) takes the opposite approach, adapting the teacher *pre-hoc* by prepending learnable prompt tokens (adding only 0.0007% of the teacher’s parameters) that cause the teacher to produce “student-friendly” distributions more accessible to the student’s limited capacity. Post-hoc methods (Veto, Delta-KD) operate on the supervisory signal after it is produced, while pre-hoc methods (PromptKD) modify the teacher’s generation process itself. Both are effective, and they compose without interference.

TAID (Shing et al., 2025) adds yet another dimension of adaptation, namely temporal interpolation of the target itself. It constructs a time-dependent intermediate distribution via *logit-space* interpolation, $p_t(\cdot|y_{<s}) = \text{softmax}((1 - \lambda_t) \text{logit}_{q_\theta}(\cdot|y_{<s}) + \lambda_t \text{logit}_p(\cdot|y_{<s}))$, that starts close to the student and gradually approaches the full teacher distribution as training proceeds. This temporal curriculum over the target (as opposed to the input prompts as in PACED) offers an additional axis of control. PACED asks “which problems to train on,” while TAID asks “how much of the teacher’s knowledge to reveal at each stage.” The progression from PACED (problem-level curriculum) through AdaSwitch (token-level timing) to TAID (target-level interpolation) and PromptKD (teacher-level adaptation) traces an important design spectrum, showing that curriculum can operate at four distinct granularities (prompt, token, target, teacher), all of which target the same underlying bottleneck of keeping the difficulty gap between teacher signal and student capability in the productive learning range. Because these four granularities are independent, they can in principle be combined within a single training loop, though no existing system integrates all four simultaneously.

5.2 Black-Box and API-Constrained Distillation

When only API access is available, the student observes the teacher’s top-1 text output (or at best, top- k log-probabilities) rather than the full distribution. This information bottleneck rules out token-level divergence matching and forces methods to operate at the sequence level. Despite this limitation, black-box methods have been reported to work well in practice, suggesting that a substantial portion of the teacher’s knowledge can be recovered from text-only observations.

Distilling Step-by-Step. Hsieh et al. (2023) demonstrated the foundational black-box approach, extracting not just the teacher’s final answers but its intermediate reasoning steps (rationales) as additional supervision. Although technically off-policy (the student fine-tunes on static teacher-generated rationales rather than its own), this work is the conceptual ancestor of on-policy CoT distillation and established a key principle. The student is trained with a multi-task objective combining answer prediction and rationale generation, allowing a 770M-parameter model to outperform a 540B-parameter teacher with over $700\times$ fewer parameters and substantially less training data (50% fewer examples on average, up to 85% reduction on some tasks). This demonstrated that *how* the teacher reasons is often more valuable than *what* it concludes, a principle that pervades modern OPD and motivates the on-policy extensions below.

If extracting textual rationales recovers partial teacher reasoning, the natural next question is whether a richer proxy for the teacher’s full distribution can be reconstructed from black-box outputs alone. The following methods attempt increasingly aggressive reconstructions, from adversarial discrimination through verbal scoring to preference ranking.

Adversarial distribution matching. GAD (Ye et al., 2025) reformulates the distillation problem as a generative adversarial game:

$$\max_G \min_D V(G, D) = \mathbb{E}_{(x, y^T) \sim \mathcal{T}} \left[-\log \sigma \left(D(y^T) - D(G(x)) \right) \right] \quad (19)$$

A discriminator D (initialized from the student with a scalar prediction head) distinguishes student rollouts from teacher API outputs via a Bradley-Terry preference model, while the generator G maximizes V via policy gradient (implemented as GRPO) using the discriminator score as an evolving on-policy reward signal. Unlike white-box methods that match full distributions, GAD extracts knowledge through a scalar quality signal, trading distributional fidelity for API compatibility. It inherits the training instability of adversarial objectives, as the discriminator can overfit to stylistic cues rather than semantic quality, but produces consistent gains over SFT baselines.

Verbal feedback and curriculum. Lion (Jiang et al., 2023) exploits the teacher’s capacity as a natural language critic through a three-stage adversarial loop. First, in *imitation*, the student is fine-tuned on teacher-generated responses. Second, in *discrimination*, the teacher identifies instructions on which the student still falls short. Third, in *generation*, the teacher creates harder instructions targeting the student’s identified weaknesses. This closed-loop curriculum steadily narrows the teacher-student gap without requiring any gradient information from the teacher, making it applicable to any instruction-following API. Lion-13B reaches competitive performance on BIG-Bench Hard and AGIEval using only 70K training examples.

Didactic-to-constructive learning. DAIL (Mendes et al., 2026) tackles the gap between high-quality expert solutions (often out-of-distribution) and the student’s actual reasoning space. Fine-tuning on expert solutions yields poor transfer because the student rarely learns to navigate its own reasoning space. DAIL uses a two-step process, transforming expert solutions into detailed, in-distribution reasoning traces via mixed policy decoding (the student generates while conditioned on the ground-truth solution), and applying a contrastive objective to focus learning on expert methodologies rather than surface patterns. Using fewer than 1,000 expert solutions, DAIL records 10–25% pass@ k gains, consistent with the recurrent observation that the alignment of the supervision signal with the student’s current distribution often matters more than its absolute quality.

Distribution-Aligned Sequence Distillation (DASD) (Yan et al., 2026) reexamines the dominant practice of SFT on teacher-generated reasoning traces, identifying three limitations. First, inadequate representation of the teacher’s sequence-level distribution from single best-of- n samples. Second, misalignment between teacher output distribution and student capacity. Third, exposure bias from teacher-forced training. Its distribution-aligned pipeline resolves all three by generating multiple diverse teacher traces per prompt and aligning the student’s generation distribution with the teacher’s through sequence-level matching, reaching state-of-the-art performance using only 448K training samples. DDT (Zhang et al., 2026c) complements DASD’s empirical findings with a theoretical framework explaining why RL generalizes better than SFT for reasoning. RL’s on-policy data generation keeps training aligned with the model’s evolving distribution instead of a static snapshot.

Verbal score distillation. OVD (Xiong et al., 2026) replaces token-level probability matching with trajectory matching using discrete verbal scores (0–9) from teacher models. This removes the requirement for token-level alignment between vocabularies, substantially reduces memory consumption by avoiding logit storage, and reports up to +12.9% absolute EM improvement on web QA and +25.7% on math benchmarks. The effectiveness of such coarse supervision suggests that on-policy exploration itself supplies a substantial fraction of the learning signal, with the teacher’s role being closer to *selecting among student trajectories* than to *correcting individual tokens*.

These methods form a spectrum of signal richness versus access requirements. Adversarial methods recover the densest proxy for the teacher’s distribution but require the most complex training dynamics (discriminator co-training, generator-discriminator instability). Score-based methods are simpler to train but lose distributional information, recovering only a scalar quality signal. Preference-based methods occupy a principled middle ground, recovering relative quality without requiring calibrated absolute scores, and their connection to DPO-style objectives yields well-understood convergence guarantees. The signal type therefore maps naturally to API access tier, with full logit access enabling KL methods (Section 4), unavailable token-level logits motivating adversarial or score-based approaches, and ranking-only access favoring preference-based methods.

Preference-based formulations. ORPO-Distill (Singh et al., 2025) maps black-box distillation onto preference optimization, generating candidate responses from the student and querying the teacher API for pairwise rankings. The mixed-policy strategy, combining on-policy student traces with off-policy teacher traces, outperforms both pure off-policy and pure on-policy alternatives, suggesting that the contrastive structure of preference learning (learning what *not* to do alongside what to do) supplies complementary signal to the positive-only feedback of score-based methods. ThinkTuning (RRV et al., 2025) occupies an intermediate position between black-box distillation and RL. Rather than scoring after the fact, a same-size teacher reviews a random fraction of the student’s on-policy rollouts *during* the GRPO training loop and delivers short corrective hints (rather than complete solutions). This interleaving of corrective feedback episodes with RL training acts as an implicit privileged signal that shapes policy gradient updates and is particularly effective for models that lack strong prior reasoning behaviors, reaching +3.85% over zero-shot baselines across benchmarks (+2.08% on MATH-500, +2.23% on AIME, +3.99% on GPQA-Diamond over vanilla GRPO).

MoE discriminator for pipeline-staged OPD. PRISM (Wang et al., 2026d) contributes a distinct architectural innovation to black-box distillation by inserting an OPD-based *pre-alignment* stage between SFT and RLVR, tackling the observation that direct RLVR from SFT checkpoints suffers from cold-start instability due to the large distributional gap between SFT outputs and reward-aligned behavior. Rather than using a monolithic discriminator (as in GAD), PRISM employs a Mixture-of-Experts discriminator with specialized perception and reasoning expert heads, delivering *disentangled* corrective signals that distinguish visual grounding errors from logical reasoning failures. The adversarial game operates entirely at the response level without requiring teacher logits. The student generates on-policy responses, a black-box teacher (Gemini 3 Flash) generates reference demonstrations, and the MoE discriminator learns to distinguish them while delivering differentiated feedback through its expert routing. This disentangled feedback is critical for multimodal models where failure modes are heterogeneous, since a model may correctly perceive an image

but reason incorrectly about it, or vice versa. On Qwen3-VL at 4B and 8B scales, the pre-alignment stage alone yields +4.4 and +6.0 points over direct SFT-to-RLVR pipelines, and the subsequent RLVR training converges faster and more stably, suggesting that black-box OPD can serve as a structured warm-up that shapes the policy landscape for downstream RL optimization.

Rubric-based semantic distillation. ROPD (Fang et al., 2026a) challenges the assumption that effective OPD requires token-level logit access by proposing structured semantic rubrics as a scalable alternative to teacher logits. A Rubricator induces prompt-specific evaluation rubrics by contrasting teacher and student rollouts, and a Verifier scores student rollouts against these rubrics to drive on-policy optimization. The teacher model assumes both roles, and its independence from the training loop enables offline rubric generation that substantially lowers GPU requirements. Across diverse benchmarks (AIME 24/25, HMMT 25, GPQA-Diamond, HealthBench, IFEval) and configurations (Qwen3-4B and Gemma3-4B students with GPT-5.2 and Qwen3-30B teachers), ROPD matches or surpasses logit-based OPD methods while achieving approximately $10\times$ sample efficiency. On AIME 2025 with thinking mode, the Qwen3-4B student (68.75%) surpasses its GPT-5.2 teacher (67.08%), indicating that rubric-augmented optimization can facilitate student-exceeds-teacher performance even without gradient-level access. ROPD occupies a distinctive position in the black-box spectrum. Where GAD recovers a scalar quality signal through adversarial discrimination and OVD through verbal scores, ROPD recovers *structured multi-dimensional* evaluation criteria that decompose quality into interpretable rubric dimensions, offering richer feedback than scores while remaining fully black-box compatible.

Off-policy knowledge integration. Not all training signal must originate on-policy. Incorporating knowledge from off-policy sources (e.g., pre-existing teacher-generated datasets) into an on-policy framework presents its own difficulties. LUFFY (Yan et al., 2025) handles this by extending GRPO to a mixed-policy objective that incorporates off-policy reasoning traces alongside the student’s on-policy rollouts. The key design challenge is preventing *superficial imitation*, because without correction the student tends to copy high-probability tokens from off-policy traces while ignoring low-probability but important reasoning steps. LUFFY’s policy shaping via regularized importance sampling reweights gradients to emphasize learning from unfamiliar yet effective actions, posting +6.4 average points over standard RLVR.

5.3 Self-Distillation

Self-distillation removes the need for an external teacher entirely. Instead, the model exploits asymmetries *within itself* to construct a training signal. Three families of methods have emerged, distinguished by the source of this internal asymmetry, namely privileged information that is available at training but not inference time, self-play between successive model checkpoints, and external feedback from verifiers or environments that the model itself has generated.

5.3.1 Privileged Information

The first family creates asymmetry through *privileged information* (PI), meaning training-time access to context unavailable at inference. Conditioning the same model on additional context produces a strictly stronger policy that can supervise the unconditioned version, turning a single network into both teacher and student.

Ground-truth as PI. OPSD (Zhao et al., 2026b) constructs the teacher policy by conditioning on both the problem x and the ground-truth answer y^* , while the student conditions only on x . The student generates on-policy rollouts, and the teacher delivers dense token-level supervision.

$$\mathcal{L}_{\text{OPSD}}(\theta) = \mathbb{E}_{(x, y^*) \sim \mathcal{S}} \mathbb{E}_{\hat{y} \sim p_{\theta}(\cdot | x)} \left[\sum_{n=1}^{|y|} D(p_{\theta}(\cdot | x, y^*, \hat{y}_{<n}) \parallel p_{\theta}(\cdot | x, \hat{y}_{<n})) \right] \quad (20)$$

On competition-level math benchmarks (AIME 2024/2025, HMMT), OPSD matches or exceeds GRPO across all three scales (1.7B, 4B, 8B) while using only a single rollout per

problem versus GRPO’s eight, yielding a significant computational advantage. The gains are largest at 1.7B (+5.7 average points) and more modest at 4B and 8B (+0.9 and +0.8), likely because GRPO’s reward diversity collapses faster at smaller scales (over half of its batches yield zero reward standard deviation within 100 steps on Qwen3-1.7B), giving OPSD’s dense token-level signal a larger relative advantage. The authors note a theoretical limitation, namely that when problems exceed a model’s comprehension threshold, even the ground-truth-conditioned teacher cannot supply meaningful supervision, suggesting that the PI mechanism has a problem-difficulty ceiling that scales with model capacity.

A complementary limitation arises when all rollouts for a given prompt fail (the “cliff prompt” problem), causing the policy gradient to vanish entirely. This failure mode directly parallels PACED’s finding (Section 6.1) that gradient SNR vanishes when student pass rate approaches zero, but with an important difference. PACED tackles this through curriculum selection (avoiding problematic prompts), while HDPO (Ding, 2026) tackles it through signal augmentation. HDPO augments GRPO with a JSD-based self-distillation term on ground-truth-conditioned rollouts, providing non-zero gradients even at the boundary of the student’s competence, directly patching the cliff-prompt vulnerability without restricting the problem distribution.

Relaxing the ground-truth requirement. Ground-truth labels are not always available. GATES (Stein et al., 2026) shows that *any* training-time context can serve as PI by having a single model act as tutor (conditioned on a source document) and student (answering from the question alone). Because source documents do not guarantee correct answers, GATES presents a consensus-based gating mechanism that suppresses the gradient when tutor uncertainty is high, preventing the model from reinforcing unreliable self-supervision. π -Distill (Penaloza et al., 2026) generalizes the OPSD and GATES paradigms into a unified theoretical framework, formalizing the conditions under which privileged self-distillation provably improves over standard training.

Behavioral PI: context distillation. The PI paradigm extends beyond factual ground truth to behavioral instructions. OPCD (Ye et al., 2026b) treats accumulated experiential knowledge (historical solution traces) and optimized system prompts as privileged context, training the student on its own rollouts while minimizing reverse KL against a context-conditioned teacher, allowing a model to internalize behaviors that would otherwise require expensive in-context demonstrations at inference time. OPSDL (Zhang et al., 2026e) extends this to long-context settings, where the model’s own short-context capability serves as the self-teacher to supervise its on-policy long-context generation via token-level reverse KL, enhancing long-context capabilities without relying on high-quality supervision data or sparse sequence-level rewards. OEL (Ye et al., 2026a) extends to continuous post-deployment learning, extracting transferable experiential knowledge from interaction trajectories as dynamic PI and iterating knowledge extraction with on-policy consolidation to form an online learning loop, closing a loop that OPCD opened.

Reasoning compression via PI. CRISP (Sang et al., 2026) finds that a model prompted to “be concise” can serve as a teacher for its own verbose version, reducing chain-of-thought token count by 57–59% on MATH-500 while *improving* accuracy by 9–16 percentage points. Much of the verbosity in distilled reasoning models is trainable inefficiency rather than a necessary feature of competent reasoning.

Visual PI for GUI grounding. GUI-SD (Zhang et al., 2026f) extends the privileged information paradigm to multimodal GUI grounding, where the task is to predict click coordinates on a screen given a natural language instruction. The privileged context is a Gaussian soft mask centered on the target element’s bounding box, giving the teacher version explicit spatial localization that the student must learn to infer from raw screenshots alone. This represents a qualitatively different form of PI from textual ground truth (OPSD) or document context (GATES), one that is *visual* and *spatial*, suggesting that the PI mechanism extends across modalities. Beyond the new PI type, GUI-SD proposes entropy-guided token weighting via a dual-criterion, where tokens receive higher loss weight based on both their digit significance (coordinate digits carry more semantic weight than filler tokens) and the teacher’s confidence at that position. With only a single on-policy rollout per instance, GUI-SD outperforms GRPO (which requires multiple rollouts for reward estimation) on

six GUI grounding benchmarks including ScreenSpot-v2, ScreenSpot-Pro, and OSWorld. The success of single-rollout PI distillation in the visual domain mirrors OPSD’s finding in mathematical reasoning, offering further evidence that privileged information can deliver sufficiently dense gradient signal to reduce reliance on the sample-intensive reward estimation required by RL-based alternatives.

Cross-lingual safety via PI. MSD (Qin et al., 2026) extends the PI paradigm to multilingual safety alignment, where the same model’s English-language reasoning serves as privileged information for its low-resource-language student, transferring safety boundaries without target-language safety data. Its Dual-Perspective Safety Weighting (DPSW) concentrates distillation gradients on safety-critical tokens by combining teacher confidence with student disagreement. We discuss MSD’s broader implications for safety-aware OPD in Section 9.

Beyond KL matching via reward regularization. A recurring concern across the PI literature above is that directly matching the privileged teacher’s distribution can destabilize training and, in the worst case, degrade reasoning capability as training proceeds. PBSO (Yu et al., 2026) diagnoses this pathology through a reward-regularized lens, replacing pure KL matching with the objective $\max_{\pi} \mathbb{E}[r(x, y)] - \beta D_{\text{KL}}(\pi \parallel \pi_{\text{teach}})$, whose analytic optimum is a reward-reweighted teacher $\pi^* \propto \pi_{\text{teach}}(y|x) \exp(r(x, y)/\beta)$ that provably dominates the teacher itself whenever the reward is informative. Practically, PBSO realizes this optimum through a DPO-style pairwise loss where the preferred sample y^+ is drawn from the context-augmented teacher and the dispreferred sample y^- from the current student, preserving on-policy sampling while avoiding the forced imitation of every teacher token. On Qwen3-1.7B/4B/8B, PBSO matches OPSD’s token efficiency while surpassing its peak accuracy and avoiding the post-peak decline documented in Figure 2 of Yu et al. (2026), suggesting that preference-based targets can serve as a more stable substitute for direct KL matching than curriculum or signal-augmentation patches alone. This post-peak decline is the same pathology that CaOPD (Zhang et al., 2026b) attributes to the training-deployment context mismatch and that Kim et al. (2026a) diagnose as epistemic suppression (Section 7.2), so PBSO can also be read as a self-distillation mechanism that converts these diagnostic insights into a corrective optimization target rather than a post-hoc recalibration step.

Outcome-conditioned PI at the turn level for agentic reasoning. The PI methods above all target single-turn generation. Extending PI to multi-turn agentic settings opens a separate question of granularity: at which temporal level should privileged information intervene? Existing agentic OPD work has explored the *trajectory* granularity (TCOD’s temporal curriculum over turn prefixes, Section 6.2), the *step* granularity (MAD-OPD’s step-level decomposition, Section 8), and the *skill* granularity (Skill-SD’s skill-conditioned updates, Section 8), but none of them supply PI at the *turn* level itself. TT-OPD (Jeong, 2026) fills this gap. A gradient-free EMA teacher receives correctness-dependent cues (reinforcing phrases for trajectories scored correct, corrective phrases for those scored incorrect) that are inserted at the prompt-response boundary and then stripped from the teacher’s output log-probabilities, so the student never observes the hints directly but still receives an outcome-aware KL signal at every conversation turn. Two aspects distinguish this instance from prior PI mechanisms. *First, the privileged context encodes trajectory-level correctness rather than problem-level ground truth*, which is the first PI realization we are aware of where the teacher is conditioned on a derived property of the student’s own rollout (its outcome correctness) rather than on the problem-level ground truth directly. This generalizes the PI paradigm from single-response answer-conditioned supervision (GATES, OPSD, GUI-SD) to multi-turn behavioral regularization. *Second, the ablation isolates why turn-level PI matters here:* EMA teacher updates alone eliminate the KL-collapse pathology but still allow the multi-turn structure to erode, and outcome hints without length control trigger response explosion, so neither EMA nor hints in isolation suffice. These pathologies are among the failure modes documented in Section 7.2, and the interaction between teacher dynamics and outcome-conditioned PI is what appears to keep agentic OPD stable in this setup.

Viewed together with TCOD, Skill-SD, and MAD-OPD, TT-OPD completes a granularity spectrum for agentic distillation in which the choice of temporal unit (trajectory, turn, step, skill) is matched to the level at which errors compound in the target task.

Structured PI for video reasoning. VISD (Lin et al., 2026) extends the privileged information paradigm from generic binary signals (OPSD’s correctness conditioning) to *structured multi-dimensional* feedback for VideoLLM reasoning. A video-aware judge model evaluates rollouts along answer correctness, logical consistency, and spatio-temporal grounding quality, producing structured feedback that decomposes reasoning errors into diagnostically meaningful subtypes. The central mechanism is direction-magnitude decoupling. Rollout-level advantages from environmental rewards determine the *update direction* (which trajectories to reinforce or penalize), while the structured teacher-student discrepancy conditioned on judge feedback modulates *token-level update magnitudes*. This ensures RL stability while allowing fine-grained credit assignment aligned with specific reasoning error types rather than undifferentiated token-level matching. Combined with curriculum scheduling (distillation→RL transition) and EMA teacher stabilization for long video sequences, VISD attains approximately $2\times$ faster convergence than RLVR baselines with consistent accuracy improvement on video reasoning benchmarks. The direction-magnitude decoupling parallels Self-Distilled RLVR’s decomposition (Section 5.3.2), where self-distillation sets magnitudes and RL sets directions, but VISD’s structured judge provides richer magnitude signals by distinguishing *why* a token matters (grounding error vs. logical inconsistency vs. factual mistake). VISD thus extends the PI spectrum from binary (OPSD) through spatial (GUI-SD) to multi-dimensional structured feedback, with domain-specific error taxonomies determining what the PI encodes.

5.3.2 Self-Play

A more radical question is whether a model can improve using nothing but its own rollouts, with no PI, no external feedback, and no additional context. The answer is a qualified yes, subject to important saturation limitations.

Self-play as a two-player game. SPIN (Chen et al., 2024) casts self-improvement as an adversarial game between successive checkpoints, training the updated model to distinguish its previous iteration’s generations from human-written responses:

$$\mathcal{L}_{\text{SPIN}} = \mathbb{E}_{x, y \sim p_{\text{data}}, y' \sim p_{\theta_t}} \left[\ell \left(\lambda \left(\log \frac{p_{\theta_{t+1}}(y|x)}{p_{\theta_t}(y|x)} - \log \frac{p_{\theta_{t+1}}(y'|x)}{p_{\theta_t}(y'|x)} \right) \right) \right] \quad (21)$$

SPIN carries a theoretical convergence guarantee. The global optimum of $\mathcal{L}_{\text{SPIN}}$ is achieved if and only if $p_{\theta_{t+1}} = p_{\text{data}}$, i.e., the student’s distribution exactly matches the reference distribution. At this fixed point, the loss is zero and the gradient vanishes, so the iterative process is *stable*. Once the student reaches the target distribution, further iterations cannot destabilize it. The game-theoretic interpretation is that p_{data} is the unique Nash equilibrium of the two-player game, and SPIN’s iterative updates correspond to fictitious play converging to this equilibrium. Starting from Zephyr-7B-SFT, it improves MT-Bench from 5.94 to 6.78 over 3 iterations, with diminishing returns at each round. This convergence guarantee comes with an inherent limitation, namely that SPIN converges to p_{data} , not to the optimal policy. If the reference data is sub-optimal (as human-written responses often are), SPIN will faithfully reproduce those sub-optimality, which is consistent with the view that pure self-play largely recycles existing knowledge rather than generating new knowledge.

IRIS (Liao et al., 2026) offers a unifying framework for this family by replacing the implicit KL divergence in SPIN with an interpolative Rényi divergence parameterized by $\alpha \in (0, \infty)$, recovering SPIN ($\alpha \rightarrow 1$, KL regime), SPACE (Jensen-Shannon regime), and SPIF ($\alpha = 2$, χ^2 -regularized regime) as special cases. The Rényi parameter controls the exploration-exploitation tradeoff, with larger α encouraging mode-covering exploration (broader distributional coverage) while smaller α enables mode-seeking refinement near convergence (concentrating updates on the most likely improvements). This unification shows that the various self-play methods in the literature differ only in their implicit divergence choice, not in their underlying mechanism, and suggests that α -tuning can be used as a meta-parameter for adapting self-play behavior to task requirements.

The trajectory from SPIN through IRIS to π -Play traces an escalating response to the saturation ceiling. SPIN accepts the ceiling (converging to p_{data}). IRIS pushes the ceiling by

tuning the divergence’s exploration-exploitation balance via the Rényi parameter. π -Play relaxes it further by injecting internal privileged information that generates genuinely new training signal. Each step suggests that self-play need not be inherently limited, but that breaking saturation tends to require increasingly creative sources of asymmetry. When internal asymmetry is exhausted, a natural remaining option is to inject external ground truth, motivating the external grounding mechanisms of Section 5.3.3.

Beyond adversarial games. The self-play paradigm extends beyond explicit two-player games to encompass several lighter-weight mechanisms. SDFT (Shenfeld et al., 2026) positions on-policy self-distillation as a mechanism for forgetting resistance: it uses the model’s own in-context learning ability to construct a demonstration-conditioned teacher from itself, then trains the student via reverse KL on its own rollouts, which is mathematically equivalent to trust-region RL with implicit KL regularization, converting continual learning from demonstrations into an on-policy distillation problem. MTP-SD (Kirchenbauer et al., 2026) targets *architectural* rather than behavioral improvement, converting a pre-trained autoregressive model into a standalone multi-token predictor by scoring student-generated token spans against a frozen copy of itself as teacher, delivering $> 3\times$ faster decoding at typically 3–7% accuracy drop (model-dependent) without requiring any auxiliary module at inference time, and expanding the scope of what self-distillation can optimize.

OPSFT (Zhao et al., 2026a) offers an instructive simplification. By filtering self-generated responses for correctness (via outcome verification) and conciseness (via a fixed length threshold), standard SFT on this filtered on-policy data matches or exceeds full RL training. The combination of on-policy generation (which naturally produces diverse solution strategies) with outcome filtering (which selects high-quality strategies) implicitly performs a form of evolutionary optimization that requires no gradient through the reward function. This “generate-filter-train” recipe substantially reduces implementation complexity compared to full RL-based distillation while retaining the core benefit of on-policy data, namely distributional alignment. Self-Distilled RLVR (Yang et al., 2026a) further decomposes the gradient update into two components. Self-distillation determines the *magnitude* of per-token policy updates (how much to change each token’s probability), while environmental RLVR feedback anchors the *direction* (which tokens should increase vs. decrease). At just 200 training steps, RLSD surpasses GRPO trained for 400 steps on Qwen3-VL-8B-Instruct, exhibiting $2\times$ sample efficiency from the dual signal structure.

Minimalist self-distillation. SSD (Zhang et al., 2026d) requires no RL, no verifier, no external teacher, and no code execution. It samples solutions at training-time temperature, fine-tunes via standard SFT, and reshapes token distributions in a context-dependent manner. On LiveCodeBench v6, SSD improves Qwen3-30B-Instruct from 42.4% to 55.3%, showing that when the base model is sufficiently capable, even the simplest self-play yields substantial gains. The requirement is not a sophisticated algorithm but sufficient *diversity* in the model’s sampling distribution to produce high-quality variants that can serve as training targets.

Systematic unification of self-distillation components. While the methods above each investigate individual mechanisms (temperature sampling, privileged information, preference optimization), the interaction effects across these design choices remain unclear. UniSD (Jin et al., 2026b) addresses this gap by proposing a unified framework that systematically decomposes LLM self-distillation along three complementary axes: *supervision reliability* (multi-teacher agreement to identify trustworthy self-derived signals, token-level contrastive learning to separate correct from plausible-but-wrong alternatives), *representation alignment* (feature matching to transfer internal structure beyond output distributions), and *training stability* (EMA teacher for temporal smoothing, divergence clipping to prevent rare high-divergence tokens from dominating updates). Evaluated across six benchmarks and six models spanning three families (Qwen2.5, Llama-3.1, Gemma-3), UniSD finds that on-policy trajectories combined with multi-teacher agreement and EMA stabilization form the core ingredients of effective self-distillation. The integrated variant UniSD* improves over the base model by +5.4 points and over the strongest baseline (GKD) by +2.8 on Qwen2.5-7B, without relying on any external teacher. UniSD’s modular analysis complements the minimalist philosophy of SSD by showing *which* design choices matter most when pure temperature

diversity is insufficient, and its three-axis decomposition provides a design vocabulary for future self-distillation methods.

Yet this also exposes the intrinsic boundary of pure self-play. When diversity is achieved through temperature sampling alone, the model’s improvement is ultimately bounded by the information already present in its pre-training distribution. Bridging this boundary requires either connecting the model to external reality (Section 5.3.3) or internally generating privileged information through the self-play process itself (π -Play below).

Unifying PI and self-play. π -Play (Zhang et al., 2026g) unifies privileged information and self-play within a single multi-agent self-evolution framework. Self-play naturally produces a byproduct during task generation, the examiner agent constructs a question construction path (QCP) that provides high-quality PI at no additional cost. Three co-evolving agents (examiner π_ϕ^E , teacher π_ψ^T conditioned on QCP, student π_θ^S without it) are jointly optimized through alternating updates. As the student strengthens, the examiner generates harder tasks, and the teacher’s behavior is softly constrained to track the student via a KL penalty, creating a self-adjusting curriculum that requires no external data. On Qwen3 (4B, 8B), data-free π -Play surpasses fully supervised search agents such as Search-R1 and delivers markedly higher evolutionary efficiency than conventional self-play (Dr.Zero), with the largest gains on multi-hop benchmarks where the teacher’s token-level credit assignment delivers the most leverage.

Partial-solution PI. PAINT (Tan & Hong, 2026) occupies a distinctive position between pure self-play (no external information) and full PI (complete answer revealed). Rather than conditioning on the entire ground-truth solution (as in OPSD), PAINT exposes only a *partial* solution whose extent is adaptively controlled by the overlap between the student’s rollout and the reference. It computes a recall-style overlap score based on the fraction of reference anchors (boxed answers, formulas, key numbers) that appear in the student’s rollout, hiding more of the reference when the student is close (encouraging broader generalization since the rollout already captures the key structure) and revealing more when far (providing stronger corrective guidance when the student has drifted). This overlap-adaptive partial-solution masking creates a natural curriculum within the PI mechanism itself, where easy problems where the student nearly solves them receive minimal additional context to encourage diverse reasoning paths, while hard problems where the student diverges completely receive maximal solution exposure to anchor the learning signal. PAINT further employs energy-space interpolation that applies the distillation loss only at positions where the teacher-student entropy mismatch exceeds a threshold, producing a sparse loss that concentrates gradient signal on the most informative positions. On competition-level mathematics (AIME 2024/2025 and HMMT 2025 average), PAINT records +2.1 over the OPSD baseline and +2.9 over GRPO, suggesting that partial revelation of privileged information can outperform both full revelation and no revelation, an information disclosure curve whose theoretical characterization remains open.

Together, π -Play and PAINT show that the self-play saturation problem can be mitigated *without* external feedback by internally generating privileged information, whether through multi-agent task construction or partial-solution conditioning. This mechanism is orthogonal to the external-verification strategies below, and combining both internal PI and external grounding remains an unexplored direction.

5.3.3 External Feedback

The self-play methods above require no external input but face an inherent ceiling. Without ground-truth or verifier access, the model cannot distinguish confident-but-wrong outputs from genuinely correct ones, leading to the saturation problem analyzed in Section 7.2. External feedback methods retain the self-distillation architecture (no separate teacher model) while injecting non-differentiable truth signals, typically code execution, symbolic verification, or outcome rewards, that anchor the self-play loop and prevent it from collapsing onto self-reinforcing errors. Unlike the internal-PI mechanisms of π -Play and PAINT, external feedback bridges self-distillation and reinforcement learning by treating environmental signals as the ultimate arbiter when internal asymmetries are exhausted.

From binary rewards to dense supervision. SD-ZERO (He et al., 2026) converts sparse verifier signals into dense self-supervision through a dual-role architecture. The Generator produces on-policy responses while the Reviser, conditioned on the Generator’s output and its binary correctness signal, produces an improved version. The Generator then distills from the Reviser’s token-level distribution, converting binary feedback into dense self-supervision. On Qwen3-4B-Instruct, SD-ZERO records 68.3% avg@8 on AIME 2024, outperforming GRPO (62.5%), showing that self-revision can be more effective than standard RL for converting sparse rewards into learning signal. SDPO (Hübotter et al., 2026) extends this beyond binary feedback to *structured textual* feedback, including runtime errors, failing unit tests, and LLM judge evaluations that construct fine-grained credit assignment attributing success or failure to specific reasoning steps rather than entire trajectories. RL from Text Feedback (RLTF) (Song et al., 2026) pushes further with free-form natural language critiques from an automated judge, training the model’s single-turn policy to match its own feedback-conditioned second-turn generations.

From SD-ZERO’s binary signal through SDPO’s structured errors to RLTF’s free-form critiques, a clear trajectory emerges. As external feedback becomes richer, the self-distillation objective becomes more precisely targeted, allowing the model to correct specific failure modes rather than globally adjusting its policy. Sample-Routed Policy Optimization (SRPO) (Li et al., 2026b) shows that these approaches can be composed by routing. Correct samples follow GRPO’s reinforcement path (reward-weighted likelihood maximization) while failed samples follow SDPO’s targeted logit-level correction. The insight is that correct and incorrect samples carry distinct types of learning signal, and treating both identically wastes the information each carries. SRPO raises the five-benchmark average on Qwen3-8B by 3.4% over GRPO and 6.3% over SDPO alone across science and tool-use tasks, consistent with the view that a useful external-feedback strategy involves routing each sample to its most informative objective.

Environmental grounding. RLTF and related methods use compiler outputs, unit test results, and symbolic math verifiers as external feedback that breaks self-reinforcing echo chambers. Even if the model becomes highly confident in a flawed derivation, the verifier assigns $R = 0$, delivering an unambiguous correction signal unavailable to pure self-play. Self-Distilled RLVR (RLSD) (Yang et al., 2026a), discussed in Section 5.3.2 for its self-play decomposition structure, also belongs here. Its RLVR environmental feedback anchors the direction of per-token updates while self-distillation calibrates their magnitude, achieving $2\times$ sample efficiency over GRPO (Section 5.3.2).

Saturation analysis. Kim et al. (2026a) trace self-distillation degradation in mathematical reasoning to the suppression of *epistemic verbalization*, the model’s expression of uncertainty during reasoning. When self-distillation shortens reasoning traces, it disproportionately removes hedging phrases and uncertainty markers that serve as implicit “attention flags” for difficult sub-problems. The result is shorter but less calibrated reasoning, where the model confidently commits to flawed steps. This connects to the teacher uncertainty challenge in Section 7, where an overconfident student produces unreliable self-play targets, and the saturation problem ultimately requires either external grounding or the internally generated PI mechanism of π -Play to overcome.

Across signal sources, from white-box logits through black-box APIs to self-distillation, a fundamental trade-off between signal density and autonomy governs the design space. White-box methods supply the densest supervision ($|V|$ -dimensional distributions at every token) but require same-organization deployment and co-location of teacher and student. Black-box methods sacrifice distributional information but operate across organizational boundaries, with recent methods (GAD, OVD) recovering unexpectedly effective proxy signals. Self-distillation removes external dependencies but is bounded by the model’s pre-existing capabilities unless augmented by external verification or internally generated PI. This continuum also reflects the field’s maturity gradient. The white-box extreme is well-explored (GKD, DistiLLM, DSKD are mature), while the self-distillation extreme still lacks principled convergence guarantees for complex reasoning. Recent activity concentrates at the boundaries between categories, with methods like ThinkTuning blurring the line between black-box feedback and RL, and π -Play blurring the line between self-play and externally grounded verification.

6 Training Dynamics and Efficiency

Given an objective function (Section 4) and a signal source (Section 5), the final engineering decision is *how to stabilize and accelerate the on-policy training loop*. On-policy generation brings unique optimization challenges absent from standard supervised fine-tuning, including non-stationary data distributions (the student’s policy shifts during training, making older rollouts stale), gradient signal-to-noise ratio (SNR) collapse on hard prompts (where all rollouts fail and no useful gradient emerges), and the computational overhead of autoregressive student rollouts followed by teacher scoring (several times more expensive than off-policy training).

These challenges are not independent. SNR collapse on hard prompts wastes compute, motivating curriculum strategies that avoid such prompts. Non-stationarity invalidates cached teacher signals, motivating efficient online scoring. And the compute overhead of autoregressive generation creates a systems-level bottleneck that requires architectural solutions. The methods in this section tackle these challenges through three complementary mechanisms.

6.1 Token and Sample Weighting

Not all teacher supervision is equally reliable in the on-policy setting. The core problem is the *flawed prefix trap*, where a student generating an erroneous token early in a sequence causes all subsequent teacher predictions to be conditioned on this out-of-distribution prefix. Because the teacher has never been trained on such prefixes, its conditional distribution $p_T(\cdot|x, \hat{y}_{<t}^{\text{bad}})$ becomes noisy and unreliable. Naively matching this corrupted signal propagates errors instead of correcting them.

Teacher reliability filtering. TIP (Xu et al., 2026b) (Token Importance Profiling) offers a principled framework for token weighting by organizing importance along two axes, student entropy h_t and teacher-student divergence δ_t . Crossing these axes yields four quadrants with qualitatively distinct learning roles. High-entropy, high-divergence tokens (Q1) carry the dominant corrective signal, high-entropy, low-divergence tokens (Q2) indicate agreement on uncertainty, low-entropy, high-divergence tokens (Q3) represent *overconfident errors* where the student is certain but wrong, and low-entropy, low-divergence tokens (Q4) are genuinely mastered. The theoretical result shows that Q3 tokens are structurally invisible to any entropy-only weighting scheme, because no non-decreasing function of entropy with $f(0) = 0$ can distinguish them from Q4. TIP’s parameter-free Soft-OR score $s_t = \hat{h}_t + \delta_t - \hat{h}_t \cdot \delta_t$ recovers Q3 coverage while preserving Q1/Q2 sensitivity, reaching consistent improvements at 50% token retention across three model families with capacity gaps from $2\times$ to $9\times$.

SCOPE (Zheng et al., 2026) elevates this analysis from tokens to *rollouts*. Its dual-path architecture routes rollouts by correctness. Incorrect trajectories receive teacher-perplexity-weighted KL distillation (down-weighting unreliable guidance on flawed prefixes), while correct trajectories receive student-perplexity-weighted MLE (concentrating reinforcement at the capability boundary). This counters diversity collapse, the counterintuitive phenomenon where Pass@1 improves at the expense of Pass@k, yielding a 7.3% relative Pass@32 gain over competitive baselines across six reasoning benchmarks.

SelecTKD (Huang et al., 2025) applies a propose-and-verify mechanism at the token level, where the student proposes tokens that the teacher verifies through greedy Top-k or non-greedy Spec-k matching. Accepted tokens receive full loss while rejected tokens are masked or down-weighted, inducing an implicit curriculum quantified by Token Acceptance Rate (TAR). Positions where teacher-student agreement fails indicate either genuine ambiguity or distributional divergence, and in neither case does the unfiltered gradient carry reliable information.

AdaSwitch (Peng et al., 2025) presents a dynamic switching mechanism between pure exploration and guided distillation. Unlike static filtering that simply attenuates unreliable signals, AdaSwitch maintains a sliding window of recent token-level divergences (KL or JSD

between student and teacher logits) and computes a running average $\bar{d}_{i-1} = \frac{1}{L} \sum_{j=i-L}^{i-1} d_j$ over a window of length L . When the current token’s divergence exceeds a context-adaptive threshold $\tau_i = K \cdot \bar{d}_{i-1}$ (where K is a sensitivity multiplier), the method switches from on-policy exploration to off-policy teacher guidance. This single-switch design preserves semantic coherence (avoiding the jarring token-level alternation of prior mixing methods) while guaranteeing high-quality supervision precisely when the student’s generation drifts into high-divergence regions. The threshold’s adaptive nature matters here, as a fixed threshold tends to be either too aggressive (switching on easy prefixes where minor divergence is harmless) or too permissive (failing to intervene when the student enters genuinely problematic states). By conditioning on recent divergence history, AdaSwitch calibrates its intervention threshold to the local sequence difficulty.

These four methods span a design spectrum from fine-grained to coarse-grained intervention. TIP applies continuous weighting over every token, SCOPE routes entire rollouts into distinct loss pathways, SelecTKD makes binary accept/reject decisions per token, and AdaSwitch commits to a single trajectory-level mode switch. The progression exposes an underlying tradeoff between signal fidelity (finer granularity preserves more information) and optimization stability (coarser decisions reduce gradient variance). This tradeoff is not merely theoretical. Fine-grained methods like TIP require reliable token-level teacher signals, which degrade precisely when the flawed prefix problem is severe. Coarse-grained methods sacrifice per-token optimality but are robust to local signal corruption because a single bad token cannot contaminate the entire trajectory’s gradient. The choice along this spectrum correlates with the dominant failure mode. Token-level methods excel when teacher reliability varies sharply within a sequence (e.g., long reasoning chains where some steps are easy and others require critical insight), while trajectory-level methods are preferred when the primary concern is prefix-quality drift across the generation (e.g., multi-turn dialogue where early errors shift the entire conversation topic).

Sample-level weighting. Beyond token-level filtering, several methods weight entire samples based on their informativeness. The flawed prefix trap manifests most severely on *hard* prompts where the student’s pass rate approaches zero. On these prompts, nearly all rollouts contain early catastrophic errors, and the resulting teacher signals are dominated by noise. Methods that uniformly sample prompts waste most of their gradient budget on uninformative hard prompts or trivially correct easy prompts, with the productive learning concentrated on prompts at the boundary of the student’s competence.

Failure-mode analysis. Fu et al. (2026) present a systematic empirical analysis of why sampled-token OPD is fragile in long-horizon settings. They identify three failure modes: (1) an imbalanced one-token signal that reduces distribution matching to a single sample, (2) unreliable teacher guidance on student-generated prefixes that deviate from the teacher’s training distribution, and (3) distortions caused by tokenizer or special-token mismatch. Their proposed fix, teacher top- K local support matching via truncated Reverse KL with top- p rollout sampling and special-token masking, produces more stable optimization across math reasoning and agentic training.

Step-level reweighting for tool-integrated reasoning. The weighting mechanisms above all operate at token or trajectory granularity. SOD (Zhong et al., 2026a) identifies a failure mode unique to tool-integrated reasoning (TIR) agents that demands an intermediate *step-level* granularity. In TIR settings, erroneous tool calls inject corrupted observations into the reasoning context, causing student-teacher divergence to accelerate sharply at subsequent steps rather than drift gradually as in text-only reasoning. Standard token-level OPD cannot detect this discontinuous state transition because it lacks visibility into inter-step boundaries. SOD adaptively reweights distillation strength at each reasoning step based on the accumulated step-level divergence, attenuating potentially misleading teacher signals in high-divergence regions while preserving dense guidance where teacher-student alignment remains intact. Using Qwen3-4B as teacher and Qwen3-0.6B/1.7B as students, SOD outperforms the second-best baseline by up to 20.86% across math, science, and code benchmarks, with the 0.6B student reaching 26.13% on AIME 2025 (average@32). The step-level granularity fills a gap between TIP’s per-token weighting (too fine for detecting

tool-call cascades) and SCOPE’s trajectory-level routing (too coarse for preserving useful supervision at well-aligned steps within a partially corrupted trajectory).

The token and sample weighting methods above operate *reactively*, attenuating the loss signal *after* the student has already generated a potentially uninformative rollout. A more proactive approach asks whether we can select prompts *before* generation to maximize the expected information gain per rollout, motivating the curriculum methods below.

6.2 Curriculum and Difficulty Adaptation

The sample weighting problem motivates a more principled solution. *Curriculum design* actively selects prompts matching the student’s current competence level. Where weighting operates on the loss *after* generation, curriculum design intervenes *before* generation by choosing which prompts to present. The design space spans multiple axes, including prompt difficulty, temporal depth within multi-turn interactions, semantic granularity of supervision, and initialization strategy, with each axis targeting a distinct failure mode of naive uniform sampling.

Competence-boundary sampling. PACED (Xu et al., 2026a) models prompt difficulty dynamically using a Beta-kernel distribution centered on the student’s current pass rate. The curriculum shifts toward harder prompts as the student improves, while avoiding prompts where the pass rate is near zero (where gradient SNR vanishes) or near one (where the student has already mastered the content). This “frontier sampling” strategy concentrates the gradient budget on the narrow band of prompts where learning is most efficient.

The mathematical justification connects to the gradient SNR analysis. For a prompt with pass rate p , the expected gradient magnitude under binary reward vanishes at both extremes ($p \rightarrow 0$ and $p \rightarrow 1$) and peaks at intermediate difficulty. PACED formalizes this through a general Beta-kernel weight family $w(p) = p^\alpha(1-p)^\beta$, where the symmetric case $\alpha = \beta = 1$ yields the familiar $p(1-p)$ with peak at $p^* = 0.5$, and asymmetric choices ($\alpha \neq \beta$) shift the peak to $p^* = \alpha/(\alpha + \beta)$, allowing the curriculum to be tilted toward harder problems ($\alpha < \beta$, peak left of 0.5) or easier problems ($\alpha > \beta$, peak right of 0.5) when prior knowledge of the pass-rate distribution motivates such tilting. PACED proves this Beta family is the leading-order minimax-optimal weighting under bounded misspecification of the SNR model, and shows empirically that the symmetric default $w(p) = p(1-p)$ already captures most of the achievable gain. Concretely, PACED estimates per-prompt difficulty $\hat{p}_i$ via a one-shot evaluation phase (typically $K=8$ rollouts per prompt) and weights each prompt by $\hat{p}_i(1-\hat{p}_i)$, concentrating training on the productive learning range while avoiding already-mastered or intractable problems. The single-pass estimation suffices because the Beta kernel is minimax-robust to stale pass rates, though iterative recomputation can yield modest additional gains.

This SNR analysis has a deeper implication that extends beyond curriculum design, as it explains *why* on-policy training naturally outperforms off-policy on reasoning tasks without explicit curriculum engineering. In off-policy settings, the problem difficulty distribution is fixed by the dataset. As the student improves, an increasing fraction of training data falls into the “too easy” zone where gradients are negligible, leading to diminishing returns. On-policy training, by contrast, automatically adjusts the effective difficulty because the student’s rollouts reflect its current capabilities, naturally generating sequences at the boundary of its competence. PACED makes this automatic adjustment explicit and tunable, but even naive on-policy training captures much of the benefit simply by sampling from p_θ rather than p_{data} .

Self-adjusting curricula. Several methods from earlier sections supply implicit curriculum mechanisms without explicit difficulty estimation. π -Play’s multi-agent framework (Section 5.3.2) automatically generates harder tasks as the student strengthens, PAINT (Tan & Hong, 2026) (Section 5.3.2) adapts its partial-solution disclosure based on student-teacher overlap, and PRISM (Wang et al., 2026d) (Section 5.2) shapes the policy landscape for subsequent RL through a structured multimodal warm-up. These complement PACED’s explicit difficulty estimation with emergent difficulty adaptation.

Dual-perspective curriculum. Where PACED treats the student as a black box whose pass rate fully characterizes prompt difficulty, Uni-OPD (Hou et al., 2026) argues that pass rate alone conflates two distinct failure modes: problems the student cannot solve (low pass rate due to capability gaps) and problems the student solves incorrectly in subtle ways (moderate pass rate but low-quality solutions). Separating these requires examining not just *whether* the student succeeds but *how* it fails. From the *student* perspective, Uni-OPD implements difficulty-aware and correctness-aware data balancing that promotes exploration of informative states during rollout sampling, extending PACED’s frontier-sampling insight from a single pass-rate metric to a dual-criterion that jointly considers problem difficulty and rollout correctness. From the *teacher* perspective, it tackles a subtler failure mode. Even when the student generates informative rollouts, aggregated token-level supervision can become inconsistent with outcome-level reward signals. A trajectory that receives higher token-level teacher scores may actually produce a worse final answer. Uni-OPD’s outcome-guided margin calibration restores order consistency between correct and incorrect trajectories, so that the token-level gradient direction aligns with the sequence-level quality ordering. This dual-perspective framework generalizes across LLMs and MLLMs (5 domains, 16 benchmarks) and accommodates both single-teacher and multi-teacher settings, implying that the curriculum and supervision reliability problems are largely orthogonal and can be solved independently.

Temporal curriculum for multi-turn distillation. PACED, Uni-OPD, and the self-adjusting methods all operate along the *difficulty* axis, selecting *which prompts* to train on. TCO (Wang et al., 2026c) opens an orthogonal axis by asking not which prompts but *how deep* within a single prompt’s trajectory to apply supervision, tackling a failure mode unique to multi-turn agent distillation. In standard OPD applied to multi-turn tasks, the student generates entire interaction trajectories and receives teacher supervision at every turn. TCO identifies “Trajectory-Level KL Instability,” where as conversation depth increases, KL divergence between student and teacher distributions escalates due to compounding errors in early turns that shift the state distribution, making subsequent teacher supervision increasingly misaligned with the student’s actual trajectory context. The proposed temporal curriculum progressively expands the horizon of teacher supervision rather than applying it uniformly across all turns from the start. Two scheduling variants target complementary failure modes. Forward-to-Backward (F2B) begins supervision at early turns and progressively extends to later turns, allowing the student to first establish reliable early-turn behavior before tackling the compounding-error regime. Backward-to-Forward (B2F) uses the teacher to generate early turns (supplying a clean state distribution) and applies supervision only at later turns, progressively requiring the student to handle earlier turns independently as training progresses. On ALFWorld, WebShop, and ScienceWorld, TCO delivers gains of up to +18 points over vanilla multi-turn OPD, with the F2B variant showing particular strength on tasks where early-turn errors are most consequential. The temporal curriculum insight connects directly to the gradient SNR framework developed above. Just as PACED identifies that prompts with near-zero pass rates yield negligible gradients, TCO identifies that later turns in long trajectories suffer analogous gradient degradation because the student’s state has drifted too far from where the teacher’s supervision is meaningful.

Off-policy cold start. Li et al. (2026d) identify that OPD can fail catastrophically when the student’s initial policy is too far from the teacher’s distribution (the “thinking-pattern mismatch” problem). Their proposed solution is an off-policy cold-start phase, a brief period of standard SFT on teacher-generated data that reduces the pattern gap before switching to on-policy training. This two-phase approach, off-policy warmup followed by on-policy refinement, has become a common industrial pattern (Section 8).

Retaining knowledge during on-policy training. A complementary concern is that on-policy training can cause *catastrophic forgetting* of capabilities the student acquired during pre-training. Chen et al. (2025) (“Retaining by Doing”) systematically compare forgetting in SFT versus RL and demonstrate that the mode-seeking nature of RL, which stems from its use of on-policy data, enables keeping prior knowledge intact when learning new tasks. Minimizing reverse KL concentrates mass on newly targeted modes without collapsing existing modes, whereas mode-covering SFT (forward KL) can overwrite prior distributions. The key finding is that on-policy data itself (rather than KL regularization or advantage

estimation) underlies RL’s robustness to forgetting, and that even *approximately* on-policy data (e.g., regenerated at the start of each epoch) suffices for substantially reducing forgetting in SFT settings. This observation offers a practical guideline: on-policy generation delivers a dual benefit, both exposing the student to its own error states (allowing correction) and preserving its existing distribution (preventing forgetting).

Semantic bootstrapping. Semantic Soft Bootstrapping (Mitra & Ulukus, 2025) offers a different perspective on curriculum design by leveraging in-context learning for self-distillation. Given a problem, the method generates multiple rollouts, partitions them into correct and incorrect solutions, then prompts the same model with both as in-context exemplars to produce a robust teacher response with verified reasoning. The student is subsequently trained via KL divergence to match the teacher’s token-level logits from the bare question alone (without the in-context hints). This posts +10.6% on MATH-500 over GRPO, showing that in-context self-supervision can complement gradient-based optimization.

The hybrid SFT+OPD pipeline. The off-policy cold start (Li et al. (2026d)) combined with on-policy refinement has emerged as a common industrial recipe. Empirically, a useful transition point from off-policy to on-policy training appears to be when the student’s generation distribution reaches sufficient overlap with the teacher’s. Transitioning too early can waste on-policy compute on uninformative rollouts (the student is too weak for meaningful self-generation). Transitioning too late tends to leave the student stuck near the off-policy ceiling. Qwen3 (Yang et al., 2025a) follows this two-phase approach, using off-policy distillation to establish basic reasoning and mode-switching capabilities before transitioning to on-policy distillation where the student generates its own sequences for logit-level alignment. DeepSeek-R1 (DeepSeek-AI et al., 2025) uses 100% off-policy (no transition), which works well at extreme teacher capability but may leave performance on the table for reasoning tasks where the student could exceed the teacher through on-policy exploration.

6.3 Compute Optimization

Token weighting (Section 6.1) reduces the effective gradient budget wasted on uninformative tokens, and curriculum design (Section 6.2) avoids generating uninformative rollouts altogether. Together they improve *sample efficiency*, extracting more learning per rollout. Yet neither eliminates the underlying *systems-level* bottleneck. OPD requires generating student rollouts *and* scoring them with the teacher at every training step, creating substantial compute overhead relative to off-policy SFT regardless of how efficiently each rollout is utilized. The methods below attack this orthogonal cost dimension.

The cost breakdown for a typical white-box OPD step consists of

1. **Student rollout:** autoregressive generation of $B \times K$ completions (batch size B , K rollouts per prompt), memory-bound due to KV cache growth. In practice this component typically dominates wall-clock time, since generation is sequential and cannot be trivially parallelized across tokens.
2. **Teacher scoring:** forward pass through the teacher to obtain full-vocabulary logits on student-generated sequences, compute-bound. When the teacher is considerably larger than the student, this pass can rival or exceed the student’s backward pass in cost.
3. **Student update:** standard backward pass computing gradients of the distillation loss, dominated by optimizer state memory. This component is usually the cheapest of the three but sets the lower bound on memory usage.

Each component admits independent optimization, and the methods below attack different bottlenecks.

Prefix truncation. FOPD (Zhang et al., 2026a) observes that the useful distillation signal (measured by per-token reverse-KL loss) is concentrated in the prefix of student-generated sequences, conjecturing that the student is weakest at high-level planning decisions captured in early tokens. Later tokens contribute diminishing marginal signal as the generation becomes increasingly constrained by the prefix context. Exploiting this observation, FOPD

truncates on-policy rollouts to a prefix of length k and reverts to off-policy KD for the remainder, reaching comparable quality at $2\text{--}47\times$ reduced training FLOPs depending on truncation ratio and task. The optimal truncation point k depends on the information density profile of the task. Using Qwen3-8B as teacher and Qwen3-8B-Base as student, a prefix scheduling strategy that progressively increases the training prefix from 1 token by a fixed increment ($\Delta L = 256$) per gradient step closely matches full OPD quality while substantially reducing compute.

Offline teacher caching. Lightning-OPD (Wu et al., 2026a) decouples student optimization from teacher inference entirely by precomputing teacher log-probabilities once over fixed SFT rollouts, eliminating the need for a live teacher server during training. The key enabling condition is *teacher consistency*, which requires that the same teacher model be used for both the SFT stage (generating training trajectories) and the OPD stage (providing the reference distribution). Under teacher consistency, Lightning-OPD provably shares the same optimum as standard online OPD, with bounded gradient discrepancy and an implicit regularization effect that prevents policy drift without explicit KL penalties. The framework achieves $4\times$ higher training efficiency than standard OPD, substantially lowering the barrier for academic research.

These two methods expose a deeper structural point. The “on-policy” requirement is not binary but admits degrees of approximation. FOPD shows that partial rollouts preserve most of the distributional alignment benefit. Lightning-OPD shows that fixed (stale) teacher signals remain effective when teacher consistency holds and the implicit regularization bounds policy drift. Together, they define a *fidelity-efficiency frontier* along which the degree of on-policy approximation trades off against compute cost, with full on-policy training at one extreme (maximum alignment, maximum cost) and fully offline SFT at the other (minimum cost, exposure bias ceiling).

Speculative knowledge distillation. SKD (Xu et al., 2025c) adapts speculative decoding to the distillation setting, where the student generates candidate tokens that the teacher verifies and selectively replaces based on its distribution, creating a data generation strategy that dynamically interpolates between supervised and on-policy KD. When student samples are consistently rejected, SKD degenerates to supervised KD. When they are consistently accepted, it degenerates to on-policy KD. DistillSpec (Zhou et al., 2024) inverts this approach. Instead of using speculation to speed up training, it uses on-policy KD to improve the draft model for speculative decoding at *inference* time. The expected wall-clock speedup of speculative decoding is

$$E[S] = \frac{\tau(x)}{c\gamma + 1} \quad \text{where } c = \frac{\text{cost(draft forward)}}{\text{cost(target forward)}} \quad (22)$$

with $\tau(x)$ denoting the block efficiency (expected tokens generated per verification step) and γ the block size. DistillSpec improves $\tau(x)$ by aligning the draft model’s distribution with the target model’s, yielding 10–45% speedup improvements over standard speculative decoding.

GPU memory management. Beyond FLOPs, the most severe constraint for white-box OPD is GPU memory. Holding a 70B teacher alongside a 7B student requires the teacher’s weights (~ 140 GB in BF16), the student’s weights and optimizer states (~ 84 GB), and the full-vocabulary logits tensor $[B, T, |V|]$ for distillation. Peak memory can easily exceed an 8×80 GB H100 node. Practitioners alleviate this via (1) *teacher quantization*, serving the teacher in FP8 or INT4 for a $2\text{--}4\times$ memory reduction at negligible precision loss, (2) *logit offloading*, immediately streaming logits to CPU RAM or retaining only the top- k entries (since 99% of the probability mass is typically concentrated in fewer than 100 tokens), and (3) *aggressive gradient checkpointing* to minimize student activation memory during the backward pass. Combining these techniques is typically required for viable white-box OPD on standard clusters.

Concrete cost example. To ground these formulas, consider distilling a 70B teacher into a 7B student on $8 \times \text{H100}$ GPUs (representative estimates based on scaling from reported benchmarks). Off-policy over 1B tokens involves three stages. The teacher generates the dataset offline (~ 200 GPU-hours), the student trains for ~ 100 GPU-hours, totaling

~300 GPU-hours. On-policy over 1B tokens, each step requires (1) student generation ($\sim 3\times$ forward-pass cost due to autoregressive decoding), (2) teacher scoring (one 70B forward pass), and (3) student backward pass, yielding $\sim 1,200\text{--}1,500$ GPU-hours, a $4\text{--}5\times$ overhead consistent with empirical measurements at smaller scales (Wu et al., 2026a). For reasoning tasks, on-policy methods consistently outperform off-policy SFT, and the off-policy performance ceiling is lower because static datasets cannot cover the combinatorial space of valid reasoning paths. This “ceiling gap” widens with task complexity, providing the strongest cost-benefit justification for on-policy methods. Section 7.4 formalizes this cost-quality tradeoff into a decision framework.

Asynchronous generation-training decoupling. NPD (Rang et al., 2026) attacks a major compute bottleneck of OPD, the synchronous coupling between student generation and gradient updates that forces the GPU to alternate between memory-bound decoding and compute-bound backpropagation. NPD reformulates on-policy distillation as a three-stage asynchronous pipeline. Student rollouts are batch-generated via vLLM, teacher top- k logits are computed through parallel prefill with sequence packing, and the student trains with a composite CE+KD loss on packed sequences. The key challenge in async OPD is policy lag, where the generation policy drifts from the training policy across async iterations, potentially degrading into off-policy learning. NPD counters this with Δ -IFD filtering, which estimates teacher-student discrepancy via an Instruction-Following Difficulty metric and rejects extreme out-of-distribution samples that would push updates outside a “safe proximal learning zone.” The resulting framework records $8.1\times$ throughput speedup over synchronous on-policy baselines while outperforming SFT by $+8.09\%$ averaged across 11 evaluation benchmarks spanning general, math, reasoning, and code tasks. NPD also shows that on-policy distillation produces superior initializations for subsequent RL, narrowing the exploration space for GRPO: the NPD $\rightarrow$ GRPO pipeline enables openPangu-Embedded-1B to reach 68.73% average accuracy, surpassing the larger Qwen3-1.7B (63.69%). The relationship between NPD and the other compute methods is one of orthogonal scope. FOPD reduces the *length* of each rollout, Lightning-OPD eliminates the *freshness* requirement for teacher signals, and NPD eliminates the *synchronization* requirement between generation and training. All three can be composed within a single system.

Drift-aware dynamic truncation. Prune-OPD (Yang et al., 2026d) attacks a subtler source of compute waste that arises specifically in long-horizon reasoning. As the student generates extended sequences, its prefix inevitably diverges from the teacher’s reasoning path, making subsequent teacher rewards increasingly unreliable. Rather than imposing a fixed truncation length (as in FOPD), Prune-OPD continuously monitors local student-teacher compatibility via per-position top- k overlap ratios. When this metric falls below a threshold, a *prefix-drift event* is recorded and subsequent OPD rewards are monotonically down-weighted. Once cumulative drift exceeds a budget, the rollout is truncated entirely, reallocating compute strictly to regions of reliable teacher supervision. Across diverse teacher-student pairs (DeepSeek-R1-Distill-Qwen-7B $\rightarrow$ R1-Distill-Qwen-1.5B, JustRL-DeepSeek-1.5B $\rightarrow$ R1-Distill-Qwen-1.5B, Qwen3-4B $\rightarrow$ Qwen3-1.7B/4B-Base), Prune-OPD reduces training time by 37.6–68.0% while preserving or improving performance on AMC, AIME, and HMMT. The adaptive behavior is key: when student-teacher compatibility remains high, Prune-OPD automatically preserves long-context supervision by expanding the training window, avoiding the information loss of fixed truncation. FOPD and Prune-OPD thus represent complementary truncation philosophies. FOPD uses a *scheduled* prefix that grows uniformly regardless of compatibility, while Prune-OPD uses a *reactive* budget that contracts or expands based on real-time drift detection, better suited for long-horizon tasks where drift timing varies across instances.

The training efficiency stack. The three subsections above form a layered efficiency stack. Token weighting (Section 6.1) reduces noise at the gradient level, so that each rollout contributes more effectively to learning. Curriculum design (Section 6.2) reduces waste at the sampling level, so that each rollout targets the student’s competence frontier. Compute optimization (this section) reduces overhead at the systems level, lowering the cost of producing and scoring each rollout. These layers are complementary rather than substitutive. PACED’s frontier sampling generates maximally informative prompts, TIP’s token weighting extracts maximal signal from the resulting rollouts, and FOPD’s prefix truncation minimizes the

generation cost of those rollouts. Deploying all three layers simultaneously can compound their individual $2\text{--}4\times$ efficiency gains, potentially reducing the effective cost gap between on-policy and off-policy training from $4\text{--}5\times$ to near parity while retaining the quality ceiling advantage.

7 Understanding OPD: Theory, Failure Modes, and Cost

The preceding sections addressed the mechanics of on-policy distillation, covering what to optimize (Section 4), where the signal comes from (Section 5), and how to stabilize training (Section 6). This section turns to the equally important questions of *why* and *when*: under what conditions does OPD succeed, how does it fail, what theoretical frameworks unify the methods surveyed above, and when is OPD preferable to simpler off-policy pipelines?

These four dimensions of understanding are not independent but form a cumulative argument. Identifying *when* OPD works (Section 7.1) establishes the preconditions that the theoretical framework must explain. Cataloging *how* it fails (Section 7.2) motivates the specific design choices, adaptive divergences, curriculum weighting, and hybrid RL objectives, whose theoretical grounding Section 7.3 supplies. Both the success conditions and the failure modes then inform the practical cost-benefit analysis in Section 7.4, which distills the theoretical insights into engineering guidelines. Resolving these questions moves OPD from an empirical art toward a principled engineering discipline with predictable outcomes (Zhang et al., 2026c; Li et al., 2026d).

7.1 Success Conditions

Li et al. (2026d) present the most systematic investigation of OPD training dynamics, identifying two necessary conditions for effective distillation. First, the student and teacher must share compatible reasoning patterns (operationalized as high overlap in their top- k token distributions). Even a superior teacher fails to transfer knowledge if its thinking style is structurally mismatched with the student’s (e.g., a non-thinking teacher distilling into a thinking student), because low initial distributional overlap cannot be recovered through training alone. Second, the teacher must supply genuinely new capabilities beyond what the student has already acquired. When both models are trained on the same data and recipe, they converge to similar distributions at their respective scales, leaving the teacher with little transferable signal, as reverse distillation experiments show that same-family 1.5B and 7B teachers become distributionally indistinguishable from the student’s perspective.

At the token level, successful OPD is characterized by progressive alignment on *high-probability overlap tokens*, a small shared token set that concentrates the bulk of the probability mass. When these conditions fail, recovery strategies include off-policy cold start (SFT initialization that reduces the pattern gap) and teacher-aligned prompt selection. Critically, the quality of OPD’s dense supervision degrades systematically with trajectory depth, with instability originating at later tokens and propagating backward, raising open questions about OPD’s applicability to long-horizon reasoning.

These findings point toward a general principle. OPD’s benefit scales with the *exploitable gap* between teacher and student. When this gap is too small (same-recipe models), the teacher offers no new information. When too large (thinking-pattern mismatch), the student cannot absorb the supervision. The productive regime lies in between, and identifying it remains a key diagnostic challenge.

Diagnostic checklist. Based on the surveyed literature, four pre-training diagnostics emerge as predictors of OPD benefit:

1. *Token overlap ratio*: Compute top- k overlap between teacher and student on a held-out set. If too low, apply off-policy cold start first.
2. *Pass rate on target prompts*: If the student cannot solve any problems in the target domain, OPD gradients will vanish. Use curriculum pacing (Xu et al., 2026a) or easier warm-up tasks.

3. *Teacher calibration*: Verify that the teacher’s confidence correlates with correctness on student-generated prefixes (Zhang et al., 2026b). Poorly calibrated teachers inject noise.
4. *Length stability*: Monitor output length during early training. Monotonic growth indicates the KL gradient asymmetry identified by Luo et al. (2026). Apply Stable-OPD corrections preemptively.

OPSD as compression, not correction. Kim & Lee (2026) sharpen the characterization of *what* on-policy self-distillation actually accomplishes in thinking-enabled mathematical reasoning. By applying OPSD separately to correct-only versus incorrect-only rollout groups, they isolate two potential mechanisms. Correct-only OPSD preserves accuracy while substantially shortening reasoning traces (compression), whereas incorrect-only OPSD damages accuracy (correction fails). Three alternative explanations (richer teacher context, mid-trace feedback reinjection, and extended training) are tested and ruled out, supporting the interpretation that the hindsight-conditioned self-teacher can identify redundancy in long thinking traces but does not reliably supply better alternative reasoning steps. This finding refines the success conditions above. OPSD’s productive regime appears to be “making the model express known solutions more efficiently” rather than “making the model solve harder problems”. The practical implication is a refined pipeline ordering. SFT establishes format compliance, RLVR expands the set of reachable correct trajectories, and OPSD compacts them for cheaper inference. Reversing the last two stages (compressing before exploration) tends to remove the raw material that exploration would produce, suggesting why post-RL placement is preferred. Across Qwen3-8B and AceReason-Nemotron-7B (with DeepSeek-R1-Distill-Qwen-7B as a supplementary pre-RL reference), the post-RLVR + correct-only OPSD configuration reaches the most favorable position in the accuracy-length plane in their experiments, consistent with OPSD’s mechanistic role rather than implementation details.

7.2 Failure Modes

Complementing the success conditions, a growing body of work catalogs the specific ways OPD can fail even when structural preconditions hold, as artifacts of the on-policy training dynamics themselves. Understanding these failure modes helps debug OPD pipelines in practice, and each mode has a corresponding mitigation strategy developed in the methods sections above.

We organize failure modes by their *root cause* instead of their symptom, as multiple failure modes can produce similar surface-level degradation (e.g., accuracy collapse) through different mechanisms.

The flawed prefix trap. As detailed in Section 6.1, when the student generates an erroneous prefix, the teacher’s conditional distribution over subsequent tokens is poorly calibrated because the teacher was never trained on such out-of-distribution inputs. Fu et al. (2026) formally characterize three specific failure modes, including (1) an imbalanced one-token signal that reduces distribution matching to a single sample, (2) unreliable teacher guidance on out-of-distribution student prefixes, and (3) tokenizer mismatch distortions. The theoretical analysis identifies a bias-variance tradeoff, where token-level OPD is biased relative to sequence-level Reverse KL but enjoys tighter worst-case variance bounds.

Two subtler failure modes compound the prefix trap. *Gradient SNR collapse* occurs when the student’s pass rate on a prompt approaches zero. In this state, all rollouts contain early errors, teacher signals are dominated by noise, and the gradient SNR vanishes. This explains why hard prompts yield no learning signal in standard OPD and motivates PACED’s competence-boundary curriculum (Section 6.2). *Epistemic suppression* (Kim et al., 2026a) manifests in self-distillation. When the process shortens reasoning traces, it disproportionately removes hedging phrases and uncertainty markers that serve as implicit attention flags for difficult sub-problems. This produces shorter but less calibrated reasoning where the model confidently commits to flawed steps. This calibration degradation connects directly to the teacher uncertainty challenge, where an overconfident student produces unreliable self-play targets, creating a self-reinforcing cycle that compounds with each iteration.

Self-play saturation (the Ouroboros problem). In self-distillation settings, the student optimizes against a target sharing its own inductive biases and architectural limitations, causing the hypothesis space to gradually collapse. If the model discovers a syntactic “hack” or a highly confident but flawed reasoning heuristic, no external signal penalizes it. The model self-reinforces this flawed trajectory, driving $p_\theta(y_{\text{flawed}}|x) \rightarrow 1$. Once entirely certain of its own hallucinations, gradients vanish, exploration ceases, and the policy is trapped. This is distinct from exposure bias (which arises from train-test distribution mismatch in off-policy settings). Saturation arises from the absence of distributional diversity in fully on-policy self-play, bearing structural similarities to GAN mode collapse, where a generator locks onto a narrow output manifold once the discriminator signal becomes non-informative.

The precision-recall tradeoff. Cha & Cho (2025) formalize a core tension in distillation. As the student concentrates mass on high-quality outputs (improving precision), it necessarily reduces coverage of the teacher’s full output distribution (reducing recall). In the OPD setting, this manifests as the “diversity collapse” problem. Aggressive Reverse KL distillation produces a student that scores highly on Pass@1 (precision) but catastrophically on Pass@ k (recall), because the student has collapsed onto a single reasoning strategy per prompt. On-policy generation can shift this tradeoff by allowing the student to explore its own distribution instead of being constrained to teacher-demonstrated paths, potentially improving recall without sacrificing precision. Temperature-scaled sampling during on-policy generation, as used in GKD and G-OPD, offers a direct control knob: higher sampling temperatures encourage distributional exploration (increasing recall) at the cost of noisier gradients (potentially reducing precision).

The calibration-capability gap. CaOPD (Zhang et al., 2026b) identifies a “Scaling Law of Miscalibration”. While OPD reliably improves task accuracy, the authors report that it tends to leave models in a state of severe overconfidence. They attribute the root cause to an information mismatch. Teacher supervision is formed under privileged context available during training (full reasoning traces, multiple samples), whereas the deployed model must report confidence using only its deployment-time context. By decoupling capability improvement from calibration degradation, CaOPD argues that the distilled model becomes more capable but less aware of its own uncertainty boundaries. This finding has potential implications for deployment. A distilled model that scores higher on benchmarks but does not reliably indicate when it might be wrong is arguably less safe than a weaker but better-calibrated baseline. The “illusion of certainty” connects to the saturation analysis above. Both indicate that OPD can produce students that lose metacognitive capabilities (uncertainty awareness, epistemic verbalization) even as task-level performance improves.

Agentic collapse in multi-turn OPD. The preceding failure modes all assume a single-turn response. Extending OPD to multi-turn tool-using agents surfaces a qualitatively different failure family, in which the interaction between teacher dynamics and trajectory structure produces pathologies unseen in single-turn settings. TT-OPD (Jeong, 2026) diagnoses three such pathologies on Healthcare AI Gym, and each maps onto a distinct structural mismatch that other agentic OPD methods also confront from complementary angles. *First, teacher-dynamics collapse.* Periodic hard-copy teacher resets cause a KL collapse (divergence drops abruptly from its accumulated value to near zero, e.g., $2.637 \rightarrow 0.343$ at a single copy event with period $T=30$), erasing the distillation gradient and driving accuracy monotonically downward. This is the multi-turn analogue of TCOD’s “Trajectory-Level KL Instability” (Section 6.2), but surfaced at the teacher-update granularity rather than the turn-depth granularity, and shows that stable teacher dynamics (EMA rather than hard copy) are a prerequisite for any other structural fix. *Second, trajectory-structure erosion.* Even with EMA teachers that eliminate KL collapse, the multi-turn structure itself erodes, as the student learns that verbose single-turn monologues maximize the sparse terminal reward more easily than coordinated tool-use sequences (turns collapse from 7.65 to 5.52 per episode under periodic resets, and still erode from 7.82 to 6.23 even with EMA teachers). This mirrors the reward-granularity mismatch that motivates Skill-SD’s skill-level decomposition and MAD-OPD’s step-level sampling (Section 8), indicating that sparse terminal rewards alone cannot sustain multi-turn structure regardless of which granularity is chosen. *Third, reward-hint runaway.* Adding outcome-conditioned teacher hints without length control

triggers a response explosion, where positive hints reinforce ever-longer reasoning until responses saturate the context limit and accuracy collapses from 54.5% to 49.0%. This is the multi-turn extension of the epistemic-suppression phenomenon documented in [Kim et al. \(2026a\)](#), as both arise when informative supervision signals that are helpful *on average* are applied without regularizers on the shape of the student’s response.

Taken together, these pathologies imply a general principle for agentic OPD. Stable teacher dynamics, explicit regularizers on trajectory structure, and granularity-matched credit assignment appear jointly necessary based on current evidence, and each existing method (TCOD, Skill-SD, MAD-OPD, TT-OPD) targets a different subset of the three.

Off-metric capability erosion. Beyond task-specific failure modes, [Wang \(2026\)](#) argue that distillation evaluation itself is structurally incomplete. Current practice reports what students retain (task scores) without accounting for what they lose (off-metric teacher capabilities that make those scores reliable). Reframing distillation as a lossy projection, they synthesize existing evidence into a taxonomy of off-metric losses, including calibration degradation (the student scores well but does not know when it is wrong), reasoning trace quality erosion (shorter traces that lose intermediate verification steps), and distributional narrowing (mode collapse invisible to pass@1 metrics). This taxonomy cuts across the specific failure modes cataloged above and offers a unifying evaluation lens. The proposed Distillation Loss Statement (primary metric, critical off-metric capabilities, preservation targets, stress distributions, observed erosion) provides a structured template for reporting distillation outcomes that goes beyond headline accuracy, complementing the diagnostic checklist in Section 7.1 with a deployment-oriented evaluation protocol.

7.3 Unified Theoretical Perspectives

The success conditions and failure modes above share deeper mathematical roots than their diverse surface manifestations suggest.

The unified f -divergence framework of [Wen et al. \(2023\)](#) subsumes Forward KL, Reverse KL, JSD, and total variation distance as special cases parameterized by a single convex generator function f . Since all these divergences can be expressed as expectations under the student’s on-policy distribution, the choice of divergence is ultimately a regularization decision. The generator function f determines the implicit weighting of likelihood ratios $p_T(y)/p_\theta(y)$, with Forward KL up-weighting regions where the student underestimates the teacher (coverage) and Reverse KL up-weighting regions of overestimation (precision). [Wu et al. \(2025\)](#) refine this picture for discrete LLM distributions, showing that under practical epoch budgets the mode-seeking/covering dichotomy is better understood through head-tail focus. Forward KL concentrates learning on the head of the teacher distribution while Reverse KL emphasizes the tail, explaining why adaptive methods that switch between divergences based on local geometry consistently outperform any fixed choice. At a more foundational level, [Cha & Cho \(2025\)](#) show that KD in generative models induces a precision-recall trade-off modulated by a single parameter, namely teacher entropy. As the teacher distribution becomes lower in entropy (more peaked), the distilled student concentrates probability mass on high-likelihood regions, improving sample precision at the cost of distributional coverage (recall). This tradeoff is validated in large-scale LLM experiments (SmolLM2 family), where lower-entropy teachers produce students that generate sharper, more fluent text but cover fewer modes of the output space. The connection to the “exploitable gap” identified empirically by [Li et al. \(2026d\)](#) is direct. The precision gain from distillation is largest when the teacher’s entropy is substantially lower than what the student would achieve from data alone, which is precisely when the teacher offers an informative compression of the output space that data-only training cannot replicate efficiently. A complementary unification comes from [Li et al. \(2026b\)](#), who show that GRPO and self-distillation tackle complementary failure modes and can be composed within a single framework via sample routing. Correct samples follow GRPO’s reward-aligned reinforcement path while failed samples receive dense logit-level correction from a feedback-conditioned self-teacher, suggesting that the two paradigms target distinct learning signals instead of being redundant approaches to the same objective.

Length inflation (OPD students’ outputs growing progressively longer during training) is a persistent practical concern. Luo et al. (2026) attribute the root cause to a feedback loop in the reverse-KL advantage. Once the student drifts toward long repetitive trajectories, the teacher’s log-probability increases faster than the student’s on those trajectories, producing large positive advantages that further encourage degenerate continuations. Their Stable-OPD framework adds a reference divergence term anchoring the student’s distribution to a pre-distillation baseline, combined with rollout mixing that blends on-policy student generations with off-policy gold demonstrations in a decaying ratio. These stabilization mechanisms post +7.2% on average over vanilla OPD by breaking the self-amplification cycle.

G-OPD’s equivalence derivation (Section 4.3) supports the view that standard OPD can be cast as a special case of dense KL-constrained RL, bridging the distillation and RL communities and making it more likely that advances in either field transfer to the other. This connection extends to DPO-based methods. Substituting the teacher’s log-probability as DPO’s implicit reward (Section 4.3) produces the same geometric mixture targeted by GKD’s λ -mixing, relating token-level KL distillation, preference optimization, and reward-regularized RL as three optimization trajectories toward the same family of target distributions parameterized by divergence choice and supervision density.

The unified KD+RL framework of Li et al. (2025) makes this complementarity precise by decomposing the hybrid gradient into a dense KD component (stable, analytically computable) and a Monte Carlo RL component (high variance, reward-driven). KD dampens variance in policy gradient estimation while RL prevents collapse onto suboptimal teacher modes, with an effective KD-to-RL weighting ratio decreasing over training as the student’s policy stabilizes. Song et al. (2026) further show that textual feedback from a teacher can expand the effective capabilities of RL beyond what outcome-only rewards attain, offering theoretical grounding for the verbal feedback methods (Lion, OVD, ThinkTuning) surveyed in Section 5.2. Text feedback carries a denser reward signal than binary outcomes, allowing the student to learn from partial successes that binary verification would discard.

Together, these perspectives converge on a central insight. On-policy distillation, reinforcement learning, and preference optimization are not separate training approaches but rather points on a continuous spectrum parameterized by the density of supervision, the choice of divergence, and the source of the training signal.

7.4 On-Policy vs. Off-Policy: A Decision Framework

The theoretical frameworks above establish that on-policy training offers principled advantages over off-policy alternatives. But theory alone cannot settle the engineering question of when this advantage justifies the additional infrastructure cost. This section distills the accumulated evidence into a practical decision framework.

The DeepSeek-R1 anomaly. DeepSeek-R1 (DeepSeek-AI et al., 2025) represents a prominent off-policy distillation success, training students from 1.5B to 70B parameters on $\sim 800\text{K}$ chain-of-thought traces via purely off-policy SFT. On AIME 2024 (pass@1), R1-Distill-Qwen-7B reaches 55.5% and R1-Distill-Qwen-32B reaches 72.6%, outperforming direct RL at the same scale (GRPO on Qwen2.5-32B-Base scores only 47.0%). Three factors plausibly explain this anomaly. An exceptionally strong 671B MoE teacher covers the problem space densely, diverse reasoning traces include backtracking and verification strategies, and relatively small student models maintain manageable effective conditional complexity under off-policy learning. The implication is that off-policy distillation often suffices when a massively capable teacher generates sufficiently diverse traces for a relatively small student. On-policy becomes more attractive as student capacity grows and diverges from the static training distribution.

The off-policy ceiling. DDT (Zhang et al., 2026c) supplies formal evidence for the generalization limits of off-policy approaches. The key factor driving on-policy generalization advantage is that generating training data from the current policy keeps learning aligned with the model’s evolving distribution. Off-policy performance is bounded by distributional mismatch that accumulates as the student’s policy drifts from the static training data,

because fixed datasets cannot cover the combinatorial space of valid reasoning paths. This gap widens with task complexity (more reasoning steps create more paths to cover), output diversity requirements (creative tasks vs. mathematics), and model capability (stronger students explore further from training data).

Distillation vs. direct RL. The choice between knowledge distillation and direct RL depends on three factors that interact non-trivially. Teacher quality relative to the reward signal determines whether dense per-token supervision outweighs sparse outcome rewards. Student scale modulates this tradeoff: smaller students ($\leq 7B$) benefit more from dense teacher supervision because their limited capacity cannot efficiently explore reward landscapes, while larger students ($\geq 32B$) have sufficient capacity for RL exploration. Target ceiling determines whether the goal is to match the teacher (distillation-bounded) or surpass it (RL-enabled). The hybrid approaches surveyed in Section 4.3 (G-OPD, KDRL, RLAD) resolve this tension by combining teacher logits as a dense variance-reducing baseline with sparse outcome rewards that drive exploration beyond the teacher’s ceiling. Empirically, this combination tends to outperform either approach alone in the settings tested so far (Yang et al., 2026c; Xu et al., 2025a; Zhang et al., 2026h).

The hybrid pipeline as standard recipe. Qwen3 (Yang et al., 2025a), DeepSeek-V4 (DeepSeek-AI, 2026), and MiMo-V2 (Xiaomi LLM-Core Team et al., 2026) converge on a two-phase pattern in which off-policy SFT establishes a strong distributional foundation and on-policy OPD or RL closes the remaining gap between the student’s generation distribution and the optimal policy. The transition point depends on the student’s initial capability. Weaker students benefit from extended off-policy warmup to reduce the teacher-student pattern gap (Li et al., 2026d), while stronger students can transition earlier because their on-policy rollouts already occupy high-density regions of the teacher’s distribution.

Compute-quality tradeoff. The expected compute cost for N tokens under each regime:

$$C_{\text{off}} \approx N \times (F_{\text{teacher}} + F_{\text{student}} + B_{\text{student}}) \quad (23)$$

$$C_{\text{on}} \approx N \times (G_{\text{student}} + \rho F_{\text{teacher}} + F_{\text{student}} + B_{\text{student}}) \quad (24)$$

where F, B denote forward and backward FLOPs, G is the autoregressive generation cost, and $\rho \in (0, 1]$ is the teacher supervision refresh rate (the fraction of steps requiring a fresh teacher forward pass). Because $G_{\text{student}} \gg F_{\text{student}}$, on-policy training incurs a substantial multiplier. For generic knowledge transfer where the off-policy ceiling is high, this overhead is rarely justified. For reasoning, code generation, and complex mathematics where the off-policy ceiling is low and marginal on-policy returns are high, the investment typically pays off. In practice, the most effective industrial pipelines tend to front-load cheap off-policy learning and reserve expensive on-policy compute for the final quality push (Section 8.5).

Summary decision rule. Use off-policy SFT alone when the teacher-student capacity ratio exceeds $10\times$ and the task admits bounded reasoning depth (factual QA, translation, summarization). Switch to on-policy OPD when any of three conditions holds: (1) the student exceeds $\sim 7B$ parameters and begins to explore regions uncovered by static teacher traces, (2) the target task involves multi-step reasoning where intermediate errors compound, or (3) the off-policy loss plateaus while held-out reward continues to improve under on-policy rollouts. In all other cases, the hybrid pipeline (off-policy warmup followed by on-policy refinement) yields the most robust cost-quality tradeoff.

8 Applications, Systems, and Emerging Domains

The preceding sections developed OPD from theory through methods to unified understanding. This section examines how these algorithmic advances translate into practice. The shift from research methods to industrial deployment suggests that the choice of OPD method is often shaped less by theoretical elegance than by practical constraints, including available compute, teacher access level, the need for multi-teacher coordination, system-level throughput bottlenecks, and domain-specific challenges such as cross-modal transfer and privacy. We organize industrial applications by their *deployment pattern* and emerging domains by the new challenges they introduce.

8.1 Industrial Deployment

Recent foundation models have increasingly incorporated OPD into their post-training pipelines, with four distinct deployment patterns that align with the theoretical progression surveyed above.

Two-phase distillation pipelines. The most common industrial pattern combines off-policy cold-start with on-policy refinement. Qwen3 (Yang et al., 2025a) employs strong-to-weak distillation where larger models (32B, 235B-A22B) supply logit-level supervision on student-generated on-policy sequences, with the student generating responses in either thinking or non-thinking mode. A key architectural choice is “Thinking Mode Fusion,” where the student learns to produce both short, direct answers (non-thinking mode) and extended chain-of-thought reasoning (thinking mode) by distilling from teachers operating in each mode independently. On-policy generation plays an important role here because the student’s mode-selection behavior (when to think vs. when to answer directly) is itself a learned policy that is difficult to capture with static datasets. Qwen3 reports that on-policy distillation outperforms direct reinforcement learning at roughly one-tenth of the GPU hours, and further improves pass@64 on AIME benchmarks where reinforcement learning from the same off-policy checkpoint does not. Gemma 2 (Gemma Team et al., 2024) embeds off-policy knowledge distillation directly into pre-training to bridge performance gaps across its model family. The 2B and 9B models are trained via logit-matching KD from a larger teacher (the 27B model itself is trained from scratch), with both small models distilled in parallel rather than via a 27B→9B→2B cascade. While this is not on-policy distillation per se (the student trains on fixed pre-training data rather than self-generated sequences), it represents an aggressive integration where distillation is a core training ingredient from the start, foreshadowing the post-training OPD patterns that later models adopt. MiMo-V2 (Xiaomi LLM-Core Team et al., 2026) combines distillation from multiple domain-specialized teachers with RL-based training, reaching strong performance on mathematical reasoning and agentic tasks with a 309B MoE model (15B active parameters).

OPD as model consolidation. A second pattern uses on-policy distillation for *merging* multiple specialized experts into a single deployable model. DeepSeek-V4 (DeepSeek-AI, 2026) entirely replaces the mixed RL stage of its predecessor (DeepSeek-V3.2) with pure multi-teacher OPD, consolidating more than ten domain-specific experts (mathematics, coding, agent, instruction-following) into a unified 1.6T-parameter model through full-vocabulary Reverse KL distillation. Computing the full $|V|$ -dimensional KL at each position (rather than the sampled-token approximation) ensures more stable gradients when merging experts with diverse specializations, while hidden-state caching with on-the-fly logit reconstruction makes this feasible across trillion-parameter teachers (Section 4.1). KAT-Coder-V2 (Li et al., 2026a) decomposes agentic coding into five expert domains, each undergoing independent SFT and RL training, and consolidates them via on-policy distillation, scoring 79.6% on SWEbench Verified. Nemotron-Cascade 2 (Yang et al., 2026e) applies the same principle to a 30B MoE model with only 3B activated parameters, reaching Gold Medal-level performance on IMO, IOI, and ICPC World Finals with $20\times$ fewer parameters than DeepSeek-V3.2-Speciale. CoPD (Gu et al., 2026) identifies a key limitation of sequential expert-then-distill pipelines, namely that large behavioral pattern gaps between fully-trained teachers and untrained students prevent effective capability absorption. The root cause is that experts trained in isolation develop idiosyncratic reasoning strategies that are mutually incompatible when merged. CoPD’s solution is *co-evolution*, running parallel RLVR expert training with interleaved bidirectional OPD so that experts serve as mutual teachers during their own training rather than after completion. This maintains consistent behavioral patterns among experts while preserving complementary knowledge. The result is a single model integrating text, image, and video reasoning capabilities that surpasses both mixed RLVR and sequential multi-teacher OPD (MOPD). The deeper insight is architectural. Co-evolution imposes a soft constraint on teacher diversity, preventing the specialization-consolidation gap that plagues post-hoc distillation from independently trained experts.

OPD for multi-budget reasoning. ORBIT (Liang et al., 2026) targets the deployment challenge of variable reasoning depth, where users need different compute budgets depending on query difficulty. It produces stage-wise expert policies through RLVR under succes-

sively tighter context budgets ($L_{k+1} = L_k/2$), then fuses them into a single model through mode-aware Reverse KL:

$$\mathcal{L}_{\text{fusion}}(\phi) = \mathbb{E}_{k \sim U(1,K)} \mathbb{E}_{\substack{q \sim \mathcal{D} \\ o \sim \pi_{\phi}(\cdot|q, p_k)}} \left[\sum_t \log \frac{\pi_{\phi}(o_t|q, p_k, o_{<t})}{\pi_{\theta_k}(o_t|q, p_k, o_{<t})} \right] \quad (25)$$

where p_k is the mode-specific prompt and $\{\pi_{\theta_k}\}$ are the frontier experts. Model merging initializes the student to mitigate the cold-start problem identified by Jang et al. (2026). Without initialization near the expert manifold, the student’s on-policy rollouts rarely reach the teachers’ high-density regions, starving the distillation objective of useful signal. Across DeepSeek-Distill-Qwen-1.5B, Qwen3-4B, and Openmath-Nemotron-7B, ORBIT produces models that smoothly interpolate between four reasoning modes (with budget tiers from 2K to 32K tokens depending on student size) with competitive performance at each budget level, matching the initial model’s accuracy at the highest budget while delivering significant token savings at lower budgets. The connection to the consolidation pattern above is instructive. Where KAT-Coder-V2 and Nemotron-Cascade 2 merge experts specialized by *domain*, ORBIT merges experts specialized by *compute budget*, suggesting that on-policy distillation is emerging as a general-purpose multi-teacher fusion tool for any axis of model specialization.

Agentic distillation. Extending OPD from single-turn generation to multi-turn agentic tasks introduces a distinct challenge: *error compounding*. In single-turn settings, a flawed token affects only subsequent tokens within the same sequence. In multi-turn agent settings, a flawed action at step t shifts the entire environmental state, causing all subsequent teacher supervision to operate on a distribution far from what the teacher encountered during its own training. This error compounding problem motivates three complementary solutions that attack it at different granularities.

At the *trajectory level*, TCOd (Wang et al., 2026c) prevents compounding by controlling the temporal depth of teacher supervision. Its temporal curriculum progressively expands the horizon, initially supervising only short interaction prefixes (where compounding is minimal) and gradually extending to full trajectories as the student’s early-turn reliability improves. This directly counters the “Trajectory-Level KL Instability” that TCOd identifies empirically, gaining up to +18 points over vanilla multi-turn OPD on ALFWorld, WebShop, and ScienceWorld.

At the *step level*, MAD-OPD’s OPAD variant (Wang et al., 2026b) offers an orthogonal decomposition. Rather than limiting *which* steps receive supervision, it treats each agentic step as an independent distillation unit, breaking the error propagation chain by preventing gradients from flowing across step boundaries. This step-level isolation is complemented by the multi-teacher debate ensemble, which targets a *second* agentic failure mode, namely single-teacher unreliability at later trajectory steps. When individual teachers err on complex sequential decisions, the collective intelligence of debating teachers yields more robust supervision than any single teacher can. Combined with task-adaptive divergence selection (JSD for agentic stability, Reverse KL for code), OPAD lifts the agentic average by +2.4% over single-teacher OPD.

At the *skill level*, OpenClaw-RL (Wang et al., 2026e) uses hindsight-guided OPD with a Process Reward Model that delivers step-level credit assignment, training a Qwen3-4B agent across personal, terminal, GUI, and SWE environments. Skill-SD (Wang et al., 2026a) applies importance-weighted Reverse KL with dynamic teacher synchronization for skill-conditioned learning, posting +14.0% over GRPO on AppWorld and +10.9% on Sokoban.

The progression from trajectory-level control (TCOD) through step-level isolation (MAD-OPD/OPAD) to skill-level decomposition (Skill-SD) exposes a general design principle. Effective agentic distillation requires matching the granularity of the distillation unit to the granularity at which errors compound. Coarse-grained approaches (full-trajectory OPD) fail because compounding is local. Fine-grained approaches (token-level) waste compute on positions where the agent’s behavior is already correct. The emerging consensus suggests that an effective granularity is the *decision boundary*, the point at which the agent commits to an action whose consequences are difficult to reverse. This parallels findings in hierarchical RL (Ross et al., 2011), where option boundaries similarly define the natural unit of credit

assignment, and suggests that future agentic OPD methods may benefit from automated decision-boundary detection instead of hand-specified decompositions.

These four patterns reflect a broader shift in how industry treats OPD, namely as a quality refinement tool (improving an existing pipeline’s outputs), an architectural simplification tool (collapsing a multi-expert system into a single deployable model), an inference-time control mechanism (allowing a single model to serve multiple latency-accuracy operating points), and an enabler for capability domains where traditional training falls short (agentic, multi-turn, multi-modal). All four patterns rely on on-policy generation rather than static dataset distillation, consistent with the theoretical advantages of distributional alignment surveyed above.

8.2 Emerging Domains

Multimodal OPD. Vision-language models face unique exposure bias challenges because multimodal interleaved generation must maintain coherence across modalities, and high-quality visual reasoning data is scarce compared to text. A natural question arises. Can on-policy distillation transfer knowledge *across* modality boundaries, and if so, what structural conditions make this feasible?

VOLD (Bousselham et al., 2025) answers affirmatively by showing that a text-only teacher can guide a VLM student through on-policy distillation without ever observing the visual input. The teacher supplies reasoning supervision on student-generated visual reasoning traces, leveraging text-based reasoning structure as a modality-agnostic scaffold. The authors report that an SFT cold-start phase is important for cross-modal distributional alignment. Without it, the student’s on-policy rollouts in the visual reasoning space are too far from the teacher’s text-based distribution for distillation to converge. On MMMU-Pro, VOLD improves Qwen2.5-VL-3B from 27.1% to 32.0%, consistent with the hypothesis that the reasoning structure itself (rather than perceptual grounding) carries much of the transferable knowledge.

This cross-modal principle extends naturally along two axes, temporal depth and modality self-alignment. Video-OPD (Li et al., 2026c) tackles temporal video grounding, where the student must localize events within video sequences, adding a temporal dimension that static image reasoning lacks. X-OPD (Cao et al., 2026) distills a separate text-only teacher (Qwen3-A3B-Instruct) into the speech LLM Qwen3-Omni-A3B by generating on-policy rollouts across both speech and text modalities and aligning them with the teacher’s text-conditioned reasoning via policy gradient optimization with a dual in-modal and cross-modal advantage function. CORD (Hu et al., 2026) pushes further into true intra-model self-distillation, where the stronger text mode of the *same* model serves as teacher for the weaker audio mode. CORD formalizes this with weighted Reverse KL between the audio and text modes of a single model:

$$\mathcal{L}_{\text{CORD}} = \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{\hat{y} \sim \pi_{\text{audio}}(\cdot|x)} \left[\sum_t w_t \cdot D_{\text{KL}}(\pi_{\text{audio}}(\cdot|x, \hat{y}_{<t}) \parallel \pi_{\text{text}}(\cdot|x, \hat{y}_{<t})) \right] \quad (26)$$

where the audio mode generates on-policy rollouts and the text mode supplies token-level targets, combined with GRPO for outcome reward maximization. CORD reports strong speech reasoning performance, suggesting that cross-modal self-distillation can serve as a general mechanism for aligning weaker modality modes to stronger ones within a single model.

KEPO (Yang et al., 2026b) identifies a prerequisite for all these multimodal OPD methods. When initial solve rates approach zero (common in specialized visual domains such as medical VQA), sparse RLVR fails due to exploration collapse before any meaningful on-policy data can be generated. By conditioning preference optimization on teacher knowledge, KEPO supplies the dense bootstrapping signal that allows subsequent on-policy training to converge. VISO (Lin et al., 2026) extends the paradigm to video reasoning through structured self-distillation, where a video-aware judge decomposes reasoning quality into multiple diagnostic dimensions and a direction-magnitude decoupling mechanism combines RL stability with fine-grained OPD credit assignment, recording approximately $2 \times$

faster convergence than RLVR baselines (see Section 5.3.1 for details). The progression from VOLD through CORD to KEPO and VISD thus maps a complete pipeline, from dense bootstrapping (KEPO) → cross-modal transfer from a text-only teacher (VOLD, X-OPD) → intra-model self-alignment (CORD) → structured self-distillation for temporal reasoning (VISD).

Text-to-image generation via flow matching. Flow-OPD (Fang et al., 2026b) extends on-policy distillation beyond language models entirely, targeting Flow Matching text-to-image models where multi-task alignment suffers from two bottlenecks absent in LLM settings. Scalar-valued rewards induce reward sparsity that GRPO alone cannot overcome, and jointly optimizing heterogeneous objectives (image quality, text rendering, prompt adherence) triggers gradient interference that produces a “seesaw effect” where improving one capability degrades another. Flow-OPD addresses both through a two-stage pipeline that first trains domain-specialized teacher models via single-reward GRPO fine-tuning, then consolidates them into a single student through on-policy distillation. Built on Stable Diffusion 3.5 Medium, Flow-OPD raises the GenEval score from 63 to 92 and OCR accuracy from 59 to 94, yielding an overall improvement of roughly 10 points over vanilla GRPO. The success of OPD in the continuous-output flow matching setting suggests that the distributional alignment benefits documented for discrete token spaces extend to continuous generation, broadening the scope of on-policy distillation beyond its LLM origins.

Step-distilled diffusion model fine-tuning. D-OPSD (Jiang et al., 2026) addresses a complementary challenge in the diffusion domain. Step-distilled models (e.g., Z-Image-Turbo, FLUX.2-klein) achieve few-step inference through careful trajectory distillation, but standard supervised fine-tuning breaks their few-step capability because it trains on external target-image states rather than states visited by the model’s own sampler. D-OPSD resolves this by exploiting the in-context capabilities of LLM/VLM-based diffusion encoders to construct on-policy self-distillation. The same model serves as both teacher (conditioned on multimodal features of both text prompt and target image) and student (conditioned on text features alone), with training minimizing the distributional divergence over the student’s own rollouts. This preserves few-step generation while enabling concept and style adaptation. Where Flow-OPD consolidates multiple teacher experts into one student, D-OPSD maintains a single model’s generative capability during task-specific fine-tuning, representing a self-distillation variant of the same continuous-output OPD paradigm.

GUI agent distillation. While GUI-SD (Section 5.3.1) tackles self-distillation for GUI grounding via visual PI, distilling compact agents from a larger teacher for full GUI task completion poses distinct challenges. Standard GKD-style OPD struggles in this domain because student outputs are often invalid JSON structures that the teacher cannot meaningfully evaluate, and the multi-solution nature of GUI tasks (multiple valid action sequences per instruction) makes single-reference supervision lossy. LiteGUI (Wu et al., 2026b) addresses both through Guided On-policy Distillation, which provides the teacher with oracle reference trajectories from a dynamically retrieved valid action set, reducing hallucinations in teacher supervision while allowing the teacher to evaluate the student’s output in context. Building on this foundation, a Multi-solution Dual-level GRPO framework jointly aligns macro-level subtask planning with micro-level execution matching, improving exploration in long-horizon GUI agent scenarios. Using Qwen3-VL-32B as teacher and 2B–3B scale students, LiteGUI achieves state-of-the-art performance among lightweight GUI agents on ScreenSpot-Pro, OS-World, and the newly introduced Lite-Bench, while remaining competitive with substantially larger models. LiteGUI represents the first systematic integration of on-policy distillation into the GUI agent domain, complementing GUI-SD’s self-distillation approach with a full teacher-student transfer paradigm.

Embodied intelligence and physical reasoning. Moving beyond language-only domains, OPD is increasingly applied to settings where the output space is continuous and errors are physically irreversible. These applications form a spectrum from high-level spatial reasoning through planning to low-level motor control, each demanding increasingly tight distributional alignment.

At the reasoning end of this spectrum, HY-Embodied-0.5 (Tencent Robotics X et al., 2026) employs on-policy distillation as the final stage of a multi-stage training pipeline (RL →

RFT $\rightarrow$ OPD) to compress a 32B expert VLM into a 2B embodied MoT model for spatial reasoning, outperforming similarly sized models on 16 out of 22 benchmarks. The task here is cognitive (understanding spatial relations), not physical, and OPD’s role is primarily model compression with minimal distributional mismatch.

Moving toward physical planning, OPD-AV (Afsharrad et al., 2026) applies GKD with $5\times$ compression for autonomous driving on nuScenes, distilling a Qwen3-8B SFT model into a 1.7B student. The challenge intensifies because planning errors compound through time. A suboptimal lane-change decision at t constrains all future maneuvers, making on-policy training particularly important for exposing the student to the consequences of its own imperfect decisions. GUI-SD (Zhang et al., 2026f) tackles an analogous compounding problem in graphical interfaces, where the student generates click trajectories while a privileged version of itself (with access to bounding boxes and Gaussian soft masks) supplies token-level KL supervision. Its entropy-guided distillation selectively weights tokens based on digit positional significance and teacher confidence, concentrating optimization on the most impactful and reliable coordinate positions while filtering unreliable supervision, and reaches state-of-the-art on 6 GUI grounding benchmarks.

At the motor control end, VLA-OPD (Zhong et al., 2026b) applies Reverse KL on self-generated robot trajectories, using a robust teacher model to provide dense token-level supervision to a Vision-Language-Action student. Because the action space is continuous rather than discrete, distributional mismatch manifests as physically unsafe trajectories (collisions, drops) instead of mere quality degradation. The student must learn precise motor control from the teacher’s dense token-level supervision on its own trajectories, attaining better sample efficiency than pure RL and better robustness than offline SFT.

The consistent finding across this spectrum is that on-policy generation becomes *more* critical as we move from cognitive to physical tasks. The diversity of valid outputs (multiple correct grasping trajectories, multiple valid temporal localizations, multiple acceptable spatial descriptions) makes off-policy data coverage impractical, and the consequence of distributional mismatch escalates from lower benchmark scores to physical failure. This gradient suggests that embodied OPD will increasingly require tighter integration with environment simulators that can supply the dense rollout data necessary for safe policy transfer.

Medical agents and tool-augmented clinical reasoning. Clinical decision making poses a different set of challenges as an emerging domain. The task is multi-turn by nature (history taking, test ordering, result interpretation, treatment selection), tool libraries are large and heterogeneous (medical knowledge bases, imaging interpreters, drug interaction checkers), and the terminal reward is sparse relative to the episode length. Healthcare AI Gym (Jeong, 2026) offers the first gymnasium-compatible environment at this scale, covering 10 clinical domains, 3.6K+ tasks, 135 domain-specific tools, and 828K medical passages, and uses it to systematically study how OPD behaves in tool-augmented clinical reasoning. Beyond benchmarking, the study exposes an agentic-textual transfer gap. Multi-turn agentic evaluation introduces systematic overhead on knowledge-recall benchmarks (MMLU-Med drops from 83.8% under logprob evaluation to 60–66% under agentic evaluation), indicating that multi-turn tool use trades parametric precision for retrieval-augmented reasoning. The proposed TT-OPD method (discussed in Sections 5.3.1 and 7.2) reaches the best score on 10 of 18 benchmarks, with an average +3.9 percentage-point improvement over a non-RL baseline, while sustaining stable multi-turn structure throughout training. Taken together, Healthcare AI Gym suggests that high-stakes specialist domains are not merely *scaled-up* agentic environments but require distillation objectives that explicitly decouple procedural competence from parametric recall, a distinction absent in current single-domain agentic benchmarks.

Privacy-preserving OPD. DP-OPD (Khadem et al., 2026) injects differential privacy noise (DP-SGD) into the student’s gradients during on-policy training while keeping the teacher frozen. The study shows that on-policy querying can extract useful utility while providing formal differential-privacy guarantees on the teacher’s training data, tackling a growing concern as distillation from proprietary models becomes standard practice. Without such

guarantees, on-policy querying carries a risk of memorizing and reproducing private training examples from the teacher.

8.3 System-Level Integration

The compute optimization methods surveyed in Section 6.3 target algorithmic efficiency, but industrial-scale OPD also requires systems-level solutions that span the full hardware-software stack.

Production OPD pipelines must orchestrate three concurrent workloads, namely student rollout generation (autoregressive, memory-bound), teacher scoring (compute-bound forward pass), and student gradient update (compute-bound backward pass). Frameworks like OpenRLHF (Hu et al., 2024) and veRL (Sheng et al., 2024) decompose these into separate process groups with distinct parallelism strategies. Teacher inference uses tensor parallelism across 4–8 GPUs for low latency, student generation uses pipeline parallelism for throughput, and the optimizer uses ZeRO-3 for memory efficiency. The main engineering challenge is synchronization. Since the student’s policy changes after each gradient step, all pending rollouts become stale. Asynchronous architectures that tolerate slight policy staleness gain higher throughput at the cost of noisier gradients. High-throughput serving engines (vLLM (Kwon et al., 2023), TensorRT-LLM (NVIDIA, 2024)) serve the generation phase, with continuous batching and PagedAttention allowing efficient rollout generation across large prompt pools. White-box OPD additionally requires exposing the full vocabulary logits rather than just the sampled token. The communication cost is substantial. A 70B teacher with 7B student on $8 \times \text{H100}$ must transmit logit data of approximately 16 GB per batch ($([B, T, |V|] \times 2 \text{ bytes in BF16, e.g., } B=16, T=4096, |V|=128\text{K})$), feasible within a single node via NVLink (900 GB/s) but requiring careful tensor slicing or top- k sparsification for multi-node setups.

DeepSeek-V4 (DeepSeek-AI, 2026) offers the most detailed public account of full-vocabulary OPD systems engineering at trillion-parameter scale, tackling the specific challenges of distilling from 10+ teachers each potentially comprising trillions of parameters. Rather than materializing the full $|V| > 100\text{K}$ logit tensor for each teacher (which is prohibitive when distilling from 10+ teachers simultaneously), the framework caches only the last-layer hidden states in a centralized buffer. It reconstructs logits on-the-fly through the corresponding prediction head at training time, virtually eliminating memory pressure while incurring negligible recomputation cost. Training samples are ordered by teacher identity within each mini-batch, so that at most one teacher’s prediction head resides in GPU memory at any time. All weight loading and offloading proceed asynchronously without blocking the critical path. The framework further integrates FP4 quantization-aware training during the OPD phase itself, quantizing MoE expert weights and attention QK activations so that the distilled model is natively adapted to deployment-time precision from the start rather than requiring post-hoc quantization.

8.4 When to Use On-Policy vs. Off-Policy

The theoretical framework above produces concrete engineering implications. On-policy training is not universally superior, and its benefit depends on whether the specific setting satisfies conditions that justify the additional compute overhead.

The choice between off-policy and on-policy distillation tends to reduce to a small set of conditions. Off-policy SFT is often sufficient when the target distribution is well-covered by the teacher’s beam-search outputs (instruction following, factual QA) and three conditions hold simultaneously, including an exceptionally strong teacher ($> 500\text{B}$), diverse reasoning traces, and a relatively small student ($\leq 70\text{B}$). DeepSeek-R1 (DeepSeek-AI et al., 2025) exemplifies this regime. The transition to on-policy becomes attractive when (1) the student exhibits clear exposure bias (performing well on teacher-style but not user-generated prompts), (2) the task involves multi-step reasoning where compounding errors degrade quality (Ross et al., 2011; Gudibande et al., 2023), or (3) the student needs to exceed the teacher via reward-guided exploration (Yang et al., 2026c). A rough formal criterion is whether the inference-time state visitation $d_{\pi_\theta}(s)$ diverges substantially from training-time $d_{\mathcal{D}}(s)$.

Within on-policy methods, the choice between logit-level and reward-level supervision is influenced by student scale. Smaller students ($<7B$) often benefit from dense logit supervision (GKD, DistiLLM) to bootstrap competence, while larger students ($\geq 7B$) with reliable verifiers tend to benefit from reward-guided OPD (KDRL (Xu et al., 2025a), G-OPD (Yang et al., 2026c)) that allows discovery beyond the teacher ceiling. The hybrid approach combining off-policy SFT warm-up with on-policy refinement consistently performs well across the surveyed literature (Qwen3 (Yang et al., 2025a), MiMo-V2 (Xiaomi LLM-Core Team et al., 2026)). At a deeper level, OPD and RLHF share much of their computational structure, including on-policy rollouts, per-trajectory signals, and policy gradient updates. The main distinction is the signal source (teacher distribution vs. reward model preference). Methods like G-OPD and REOPOLD (Ko et al., 2026) formalize this relationship, with practical consequence. For reward-guided OPD, much of the existing RLHF infrastructure applies directly (the rollout-score-update loop is essentially the same), which helps explain why frameworks like OpenRLHF (Hu et al., 2024) serve both communities. White-box logit-level OPD additionally requires teacher co-hosting and full-vocabulary logit transfer, an extension that existing frameworks support through their teacher-inference process groups.

8.5 The Distillation Tax: Compute Budget Allocation

Based on the compute analysis in Section 6.3 and the staged pipelines described in recent industrial reports (DeepSeek-AI et al., 2025; Yang et al., 2025a), a recurring pattern is observable in compute budget allocation. Several prominent industrial pipelines approximate a three-stage structure:

1. **Off-policy warm-up** (majority of budget): Standard SFT on teacher-generated data supplies cheap, high-bandwidth learning that establishes the student’s basic capabilities and aligns its distribution close enough to the teacher for effective on-policy training.
2. **On-policy logit distillation** (largest on-policy share): Token-level OPD with adaptive divergences (GKD, DistiLLM, ToDi) closes the distribution gap on the student’s own generation trajectories, mitigating exposure bias where it matters most.
3. **Reward-guided refinement** (final smaller share): RL-augmented objectives (G-OPD, KDRL, RLAD) push the student beyond the teacher’s capability ceiling on reasoning tasks, using sparse outcome rewards to discover new solution paths.

This staged approach front-loads cheap, high-bandwidth learning and reserves expensive on-policy compute for the final quality push. Qwen3 (Yang et al., 2025a) exemplifies a related pattern, with its four-stage post-training pipeline (cold-start SFT $\rightarrow$ reasoning RL $\rightarrow$ thinking-mode fusion SFT $\rightarrow$ general RL) interleaving supervised and on-policy stages, while distillation from larger models (32B, 235B-A22B) serves as the primary transfer mechanism for smaller variants, achieving comparable performance to RL at roughly 1/10 the GPU cost. On-policy training tends to be most useful after off-policy warm-up has brought the student close to the teacher’s distribution, since this reduces the wasted rollouts that occur when the student-teacher gap is too large (Section 7.1).

9 Open Problems and Future Directions

The methods surveyed above follow a recognizable trajectory, moving from fixed objectives to adaptive ones, from external teachers to self-generated signals, and from single-turn generation to multi-turn trajectories. Continuing along this trajectory, distillation, reinforcement learning, and alignment look less like separate training stages than facets of a single optimization process in which the model generates, evaluates, and refines its own behavior under evolving constraints. The open problems below are organized along this direction.

Distillation scaling laws. While parameter-scaling laws for pre-training are well-established (Kaplan et al., 2020; Hoffmann et al., 2022), the study of distillation-specific scaling has only recently begun. Busbridge et al. (2025) proposed the first distillation scaling laws under an off-policy setting, exposing a non-monotone relationship between compute

budget and optimal teacher size: as compute increases, the optimal teacher first grows until it slightly exceeds the student, then plateaus, and eventually *decreases* because inference cost with large teachers dominates the compute budget. Meanwhile, optimal student distillation tokens continue to scale as a power law. DeepSeek-R1 (DeepSeek-AI et al., 2025) offers complementary off-policy evidence on the student-size axis (Section 7.4): on AIME 2024, performance scales as 28.9% $\rightarrow$ 55.5% $\rightarrow$ 69.7% $\rightarrow$ 72.6% for student sizes 1.5B $\rightarrow$ 7B $\rightarrow$ 14B $\rightarrow$ 32B, with the steepest gain between 1.5B and 7B and sharply diminishing returns beyond 14B, indicating that student capacity follows a power-law saturation pattern even under off-policy distillation.

Yet no equivalent framework exists for on-policy distillation, where the generation cost of student rollouts introduces an additional compute axis absent from off-policy settings. By analogy with these off-policy findings, one plausible conjectured functional form for the on-policy distillation loss is a joint power-law:

$$L(N_S, N_T, D_{\text{on}}) = E + \frac{A}{N_S^\alpha} + \frac{B}{N_T^\beta} + \frac{C}{D_{\text{on}}^\gamma} + f(N_S, N_T) \quad (27)$$

where E is the irreducible task entropy, the first three terms capture independent scaling of student size, teacher size, and on-policy rollout budget, and $f(N_S, N_T)$ models the capacity-gap interference between student and teacher. This form is speculative and has not been validated empirically for on-policy settings. Whether a separable power-law with an additive interaction term captures the true scaling behavior remains an open question, and the actual relationship may be substantially more complex. Nevertheless, even as a heuristic scaffold, this conjectured form helps organize the constraints that existing off-policy evidence imposes. The non-monotone teacher-size finding implies that B/N_T^β alone is insufficient and the interaction term f must capture the regime where teacher inference cost overwhelms quality gains. Qwen3 (Yang et al., 2025a) further shows that distilling from multiple specialized teachers yields better scaling than from a single monolithic teacher of equivalent total parameters, indicating that f must additionally account for teacher diversity. A key open question is whether the on-policy rollout term C/D_{on}^γ interacts synergistically or antagonistically with the teacher-size term, that is, whether more on-policy data can compensate for a weaker teacher, or whether teacher quality sets a hard ceiling regardless of rollout volume. Future work could conduct controlled grid searches that independently vary N_S , N_T , and D_{on} under on-policy training to disentangle these effects, addressing allocation questions such as “Given a fixed GPU budget, is it more effective to invest in a larger teacher or in more student rollouts?” Establishing such scaling laws could move OPD compute allocation closer to the predictable footing that pre-training scaling enjoys.

Uncertainty-aware feedback. Even with perfect scaling laws, a second foundational gap remains. Teachers currently supply point-estimate probabilities. When the teacher itself is uncertain about a token, its logits carry little useful information, yet the student treats all teacher signals equally. A natural next step is for OPD frameworks to incorporate the teacher’s *epistemic uncertainty*, allowing the student to discount feedback on queries where the teacher is guessing. Process reward models (Lightman et al., 2024) offer a potential path by scoring each reasoning step independently, allowing identification of precisely where teacher supervision is reliable versus speculative.

The connection to the failure-mode analysis of Section 7.2 is direct. The flawed prefix trap is at its core a problem of unrecognized teacher uncertainty, and explicitly modeling this uncertainty would yield a principled solution. The core intuition is to suppress the distillation gradient at positions where the teacher’s output entropy approaches its maximum, effectively treating such positions as uninformative. TIP (Xu et al., 2026b) offers empirical support for this direction, showing that teacher-student divergence-based filtering outperforms entropy-only schemes, suggesting that *relative* uncertainty (how surprised the teacher is by the student’s continuation) matters more than the teacher’s absolute confidence.

The next frontier is *predictive* uncertainty. Rather than measuring the teacher’s entropy on the student’s current prefix, the goal would be to estimate what the teacher’s reliability would be several tokens ahead, allowing the student to avoid committing to reasoning paths that will lead to regions of teacher incompetence. Systematically avoiding these paths

before the student wastes compute exploring them represents a major unexploited efficiency gain.

Agent-level, continual, and lifelong distillation. The uncertainty problem intensifies when distillation extends beyond single-turn generation. Extending to multi-turn agentic settings requires new credit assignment mechanisms, as the student must learn long-horizon tool use, environmental state tracking, and error recovery across entire interaction trajectories. OpenClaw-RL (Wang et al., 2026e), Skill-SD (Wang et al., 2026a), TCOd (Wang et al., 2026c), and MAD-OPD’s OPAD variant (Wang et al., 2026b) represent early steps, but the systematic degradation of teacher supervision with trajectory depth (Li et al., 2026d) remains a hard open problem. MAD-OPD’s step-level sampling strategy offers a partial mitigation by treating each agentic step as an independent distillation unit instead of propagating errors across the full trajectory, while TCOd’s temporal curriculum controls how much of the trajectory receives teacher supervision at each training stage. Together, these methods suggest that the multi-turn distillation frontier benefits from *both* temporal decomposition (breaking trajectories into manageable units) and progressive exposure (gradually increasing horizon complexity).

Three unsolved sub-problems define this frontier. *First, environment non-stationarity.* Agentic environments change in response to actions, making direct trajectory comparison meaningless unless the teacher delivers counterfactual evaluations conditioned on the student’s actual state. *Second, tool-use combinatorics.* Modern agent frameworks expose dozens of tools with structured argument schemas, suggesting distillation may need to operate at the tool-call level instead of individual tokens, with credit assignment mechanisms attributing trajectory success to specific tool selections. *Third, safety-critical exploration.* A single incorrect file write in coding agents or misguided click in web-browsing agents can trigger irreversible actions, requiring safety constraints that prevent exploring catastrophically dangerous action sequences.

On a separate axis, production models undergo repeated fine-tuning cycles (safety patches, capability updates), raising the question of how distillation objectives should adapt when the teacher itself evolves. SDFT (Shenfeld et al., 2026) and OEL (Ye et al., 2026a) offer initial answers, but *lifelong distillation*, attaining re-alignment without catastrophic forgetting, without retraining from scratch, and without accumulating distributional drift, connects to continual learning with the added complication that both teacher and student distributions are non-stationary.

Efficiency frontiers. The agentic and multi-turn settings above amplify what is already a central systems bottleneck of OPD. The considerable compute overhead of on-policy training remains a barrier. While FOPD (Zhang et al., 2026a), Lightning-OPD (Wu et al., 2026a), and SKD (Xu et al., 2025c) have made progress, the theoretical minimum cost of on-policy distillation (the information-theoretic lower bound on how much on-policy data is needed to reach a given quality level) remains unknown. A promising underexplored direction is *selective teacher inference*. MiniPLM (Gu et al., 2025) shows that selecting training instances based on the log-probability discrepancy between teacher and a reference model (“Difference Sampling”) substantially improves data efficiency under off-policy KD. Transplanting this principle to on-policy rollouts, using teacher-reference divergence to prioritize which student-generated sequences deserve dense teacher feedback, could yield significant compute savings by skipping expensive teacher inference on uninformative rollouts where the student has already converged.

Latent-space and cross-modal distillation. Efficiency improvements tackle the *cost* of OPD, but a separate limitation constrains its *scope*. Current OPD operates exclusively in vocabulary space, where cross-architecture alignment requires dedicated mechanisms (DSKD’s dual-space projections, ULD’s probability-space transport). A more radical approach would distill in a *shared latent space*, bypassing the vocabulary bottleneck entirely and allowing distillation between architecturally dissimilar models that share no structural overlap. The extension to multimodal settings raises additional open questions. How does exposure bias manifest when multiple modalities must maintain coherence? What are appropriate divergence choices when teacher and student operate in different modality spaces? Can the unified f -divergence framework be extended to handle continuous signal spaces (images,

audio) alongside discrete tokens? VOLD (Bousselham et al., 2025), X-OPD (Cao et al., 2026), and CORD (Hu et al., 2026) offer empirical starting points, but the theoretical foundations for cross-modal OPD remain to be established.

Privacy and evaluation methodology. As OPD becomes more powerful and more widely deployed, questions of trust, both in the training process and in its evaluation, grow increasingly urgent. DP-OPD (Khadem et al., 2026) offers an early principled framework for differentially private distillation, with an asymmetric design applying DP-SGD only to the student while keeping the teacher frozen. The on-policy formulation is well-suited to the DP setting because training on self-generated prefixes reduces compounding errors from DP noise. Open questions include scaling DP-OPD to reasoning-capable LLMs and integrating DP with self-distillation. On the evaluation front, Fu et al. (2026) demonstrate that token-level OPD can appear successful on standard benchmarks while failing on distribution-shifted prompts, highlighting a gap between benchmark performance and true generalization. The precision-recall framework of Cha & Cho (2025) shows distillation induces a deep trade-off between sample precision and distributional coverage, suggesting that evaluation protocols need to probe whether the student has learned underlying causal mechanisms rather than superficial correlations.

Diagnostic tools for failure modes. A natural extension of uncertainty-aware training is developing tools to detect when distillation breaks down entirely. As the scale and complexity of OPD increase, the failure modes diagnosed in Section 7.2 demand more systematic diagnostic tools. While current literature identifies isolated traps, such as the flawed prefix trap and epistemic suppression (Kim et al., 2026a), the field lacks a unified framework to detect and quantify these pathologies during training. Real-time diagnostic probes that track gradient signal-to-noise ratios, representation collapse, and teacher-student divergence dynamics without incurring the full forward-pass cost would fill this gap. Drawing inspiration from gradient profiling in pre-training, such tools could automatically surface when a student enters the Ouroboros self-play saturation or when a multi-turn agent begins to suffer from trajectory-structure erosion (Jeong, 2026). Such probes would shift OPD debugging from post-hoc benchmark failure analysis to proactive mid-training intervention.

Cross-architecture scalability. A structural barrier intertwined with the efficiency frontier is the scalability of cross-architecture distillation. As explored in Section 5.1, mapping between heterogeneous model families introduces significant friction. While DSKD (Zhang et al., 2025b) and ULD (Boizard et al., 2025) offer initial solutions via dual-space projection and optimal transport, these methods have primarily been validated on relatively small scale differences. The viability of these mappings when the capacity gap spans orders of magnitude, such as distilling a 400B MoE teacher into a 1B dense student, remains largely unproven. At such extremes, the representational bottleneck may preclude simple linear projections or vocabulary-level alignments. Developing non-linear, hierarchical alignment mechanisms that bridge massive architectural divides without prohibitive compute overhead remains an open challenge with significant practical implications.

Unified scheduling of OPD and RLVR. The convergence of distillation and reinforcement learning highlighted throughout this survey surfaces an important scheduling challenge. While joint optimization objectives (G-OPD (Yang et al., 2026c), KDRL (Xu et al., 2025a), RLAD (Zhang et al., 2026h)) show the value of combining dense teacher guidance with sparse outcome rewards, the temporal scheduling of these signals remains largely heuristic. The co-evolutionary framework of CoPD (Gu et al., 2026) points toward continuous interleaving, but the question of how to allocate compute between imitating a teacher and exploring beyond it across the training lifecycle is still unanswered. A rigorous theoretical framework for this allocation would depend on the student’s relative competence, suggesting that the appropriate schedule may not be a static hyperparameter but a dynamic policy that transitions from pure distillation (when the capability gap is large) to pure RL (when the student matches the teacher). Formulating this schedule as an active learning problem appears to be a promising frontier.

Closing the distillation-RL loop. The open problems above (scaling, uncertainty, agents, efficiency, cross-modal transfer, privacy) are individually significant, but they share a

common resolution pathway. The field has historically treated distillation and RL as separate, linearly staged processes. Static distillation saturates when the student exhausts the teacher’s knowledge, while pure RL diverges when reward signals are sparse. [Chen et al. \(2025\)](#) show that RL retains more pre-training capabilities than SFT because on-policy rollouts maintain distributional proximity to the student’s prior. KDRL ([Xu et al., 2025a](#)) jointly optimizes KD and RL, reducing the catastrophic forgetting common to sequential pipelines. REOPOLD ([Ko et al., 2026](#)) offers a further instantiation, with its relaxed on-policy framework combining dense KD gradients for early-stage stability with RL-style reward signals for late-stage exploration. CoPD ([Gu et al., 2026](#)) shows that the sequential pipeline (train experts first, then distill) can be replaced entirely by *co-evolutionary* training where RLVR and bidirectional OPD are interleaved (Section 8.1), maintaining behavioral proximity that enables continuous knowledge absorption instead of post-hoc transfer. A natural extrapolation of this co-evolution principle is *simultaneous* optimization in which distillation and RL operate as complementary signals across training. Principled methods for interleaving these objectives, deciding when to switch, and preventing RL-discovered capabilities from being overwritten during subsequent distillation phases remain open.

Toward self-improving systems. A natural endpoint of the trends above, from external teachers to self-distillation, from static datasets to on-policy generation, from single-shot training to iterative refinement, is a system that continues to improve with limited human intervention. π -Play ([Zhang et al., 2026g](#)) and SSD ([Zhang et al., 2026d](#)) supply early existence proofs, showing models that improve through self-generated curricula and self-filtered training data. The theoretical ceiling of such systems, whether they can exceed the capability of any fixed teacher given sufficient compute and environmental grounding, connects to open questions about the limits of self-play in imperfect-information games. The saturation analysis of Section 7.2 identifies the boundary conditions: pure self-play converges to the model’s own prior distribution unless grounded by external verification or internally generated privileged information. Whether it is possible to raise this ceiling without access to a stronger oracle, attaining genuine capability *creation* rather than mere capability *redistribution*, remains an open question at the intersection of distillation, RL, and learning theory. Progress in this direction is likely to depend on verification mechanisms that scale without human labeling.

10 Conclusion

This survey has systematically examined on-policy distillation for large language models, synthesizing contributions across foundational formulations, algorithmic innovations, theoretical analyses, and industrial deployments. We organize this fast-growing literature along three core design axes rather than surface-level categories. *Objective function* design progresses from fixed divergences through adaptive routing to RL-augmented objectives that lift the teacher ceiling. *Signal source* architecture spans from dense white-box logits through black-box API access to self-distillation that removes the external teacher entirely. *Training dynamics* encompasses token weighting, curriculum adaptation, and compute optimization.

Key Findings. Three patterns recur across the methods surveyed. *First*, shifting from off-policy to on-policy training distributions, from $y \sim p_{\text{data}}$ to $y \sim p_{\theta}$, has been associated with consistent accuracy gains across mathematical reasoning, code generation, and instruction following (Tables 2 and 3). The improvement is commonly attributed to reduced exposure bias, as the student learns to recover from its own errors instead of memorizing idealized trajectories, with the performance gap reportedly widening for longer reasoning chains where error compounding is more severe ([Gu et al., 2024](#); [Agarwal et al., 2024](#)). *Second*, divergence selection appears to benefit from being adapted to both task and token position rather than fixed globally. Reverse KL has been observed to do well on reasoning tasks where mode-seeking can prevent hallucination, while Forward KL preserves the diversity often needed for open-ended generation. Adaptive methods such as ToDi ([Jung et al., 2025](#)) and EOPD ([Jin et al., 2026a](#)), which select divergences based on local distributional geometry, represent the current frontier. *Third*, the boundary between distillation and reinforcement learning continues to narrow. G-OPD ([Yang et al., 2026c](#)), RLAD ([Zhang et al., 2026h](#)), and

REOPOLD (Ko et al., 2026) suggest that combining dense teacher supervision with sparse outcome rewards can be effective, with KD providing gradient stability while RL permits exploration beyond the teacher’s capability. The convergence is not only conceptual but also yields a shared toolkit (trust regions, advantage estimation, policy gradient variance reduction) that both communities can draw on.

Practical Takeaways. With white-box teacher access, token-level methods (GKD (Agarwal et al., 2024), DistiLLM (Ko et al., 2024)) tend to offer a favorable quality-compute tradeoff for general instruction following. For reasoning-intensive tasks, sequence-level methods (MiniLLM (Gu et al., 2024)) or hybrid approaches (FOPD (Zhang et al., 2026a), PAGED (Xu et al., 2026a), Lightning-OPD (Wu et al., 2026a)) are often preferred despite higher gradient variance inherent to REINFORCE-style estimators (Gu et al., 2024). With only API access, GAD (Ye et al., 2025) and OVD (Xiong et al., 2026) can support on-policy learning through adversarial and verbal formulations, respectively. Self-distillation methods (SPIN (Chen et al., 2024), OPSD (Zhao et al., 2026b), SD-ZERO (He et al., 2026)) remove the need for a separate teacher entirely, which makes them attractive under resource constraints. The substantial compute overhead of on-policy generation can often be reduced via speculative decoding (Xu et al., 2025c), prefix truncation (Zhang et al., 2026a), or offline precomputation (Wu et al., 2026a).

Looking Ahead. The progression from off-policy imitation to on-policy self-correction echoes a broader pattern in machine learning. Classical supervised learning assumes a fixed target distribution, whereas on-policy distillation treats the learner’s own distribution as part of the training objective. The methods surveyed here progressively relax the assumptions under which distillation was originally formulated as a compression technique. Fixed objectives are replaced by adaptive ones, external teachers by self-generated signals, and single-turn optimization by multi-turn trajectory learning. A natural continuation of these trends is a training regime in which models generate, evaluate, and update their own behavior more continuously rather than only during a one-shot training run. The absence of distillation scaling laws (Section 9), the open problem of teacher uncertainty quantification, and the challenge of lifelong adaptation without catastrophic forgetting (Kirkpatrick et al., 2017) appear to be among the main barriers between the current state of the art and such a regime. As LLMs evolve from static text generators into more interactive reasoning agents, the principles of on-policy alignment surveyed here seem likely to become a more integral component of the LLM training stack rather than an optional enhancement.

Limitations. This survey focuses exclusively on methods that involve on-policy generation from the student during training. We deliberately exclude general knowledge distillation (e.g., feature-based, attention transfer), model compression techniques orthogonal to the training regime (pruning, quantization), and methods that use student-generated data only at inference time (self-consistency, majority voting). Our coverage is current through April 2026. Given the field’s rapid pace (more than ten new OPD papers per month through early 2026), some concurrent work may be absent. The method selection considerations in Section 3.3 reflect current empirical evidence and hardware constraints. As compute costs decrease and training frameworks mature, the cost-benefit calculus will shift toward more aggressive on-policy methods.

Reproducibility challenges. A significant concern across the OPD literature is the difficulty of fair comparison. Methods are evaluated on different base models, different training compute budgets, different benchmark versions (e.g., MATH-500 vs. full MATH), and with different numbers of rollouts per prompt. Few papers control for all these variables simultaneously, making cross-method comparison unreliable from published numbers alone. Our Tables 2 and 3 report original-paper results rather than controlled reproductions, and readers should interpret cross-row comparisons with appropriate caution. We advocate for standardized OPD benchmarking protocols that fix the base model, compute budget, and evaluation suite, analogous to what HELM (Liang et al., 2023) offers for general LLM evaluation.

Broader impact. OPD carries immediate implications for AI accessibility, safety, and environmental sustainability.

Regarding accessibility, effective distillation allows smaller, more efficient models that run on consumer hardware, democratizing advanced reasoning capabilities that would otherwise require expensive API access to frontier models. The self-distillation methods surveyed here (OPSD, SD-ZERO, SSD) stand out in this regard, as they permit capability improvement without any external teacher, allowing resource-constrained researchers and developers to iterate rapidly.

Regarding safety, OPD introduces both risks and mitigations. The risk is that effective distillation of frontier capabilities lowers barriers to misuse. The mitigation is that the teacher retains a supervisory role during training, offering a natural point for injecting safety constraints, alignment objectives, and capability boundaries. MSD (Qin et al., 2026) concretely instantiates safety-aware OPD, transferring safety alignment from high-resource to low-resource languages through on-policy self-distillation without requiring target-language safety data. By adaptively upweighting safety-critical tokens where teacher confidence and student disagreement are both high, MSD generalizes to unseen languages and more challenging jailbreak attacks, showing that OPD can serve as a scalable mechanism for multilingual safety alignment. As the field matures, we expect safety-aware OPD, where the teacher not only transfers knowledge but also enforces alignment properties during the interactive training loop, to be an important research direction.

Regarding environmental sustainability, OPD’s compute overhead (4–5 $\times$ over off-policy SFT, Section 6.3) raises concerns at industrial scale. Yet the resulting smaller, more capable students can substantially reduce downstream inference costs, and the net carbon footprint over a model’s lifetime can be lower when amortized across billions of inference calls. The broader convergence of distillation, reinforcement learning, and self-play into a more unified on-policy training picture suggests that future LLMs may be trained less as one-pass artifacts and more as systems that are refined periodically through their interactions with users and environments.

References

- Amirhossein Afsharrad, Amirhesam Abedsoltan, Ahmadreza Moradipari, and Sanjay Lall. On-Policy Distillation of Language Models for Autonomous Vehicle Motion Planning. *arXiv preprint arXiv:2604.07944*, 2026. URL <https://arxiv.org/abs/2604.07944>.
- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-Policy Distillation of Language Models: Learning from Self-Generated Mistakes. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2306.13649>.
- Nicolas Boizard, Kevin El Haddad, Céline Hudelot, and Pierre Colombo. Towards Cross-Tokenizer Distillation: the Universal Logit Distillation Loss for LLMs. *Transactions on Machine Learning Research*, 2025. URL <https://arxiv.org/abs/2402.12030>.
- Walid Bousselham, Hilde Kuehne, and Cordelia Schmid. VOLD: Reasoning Transfer from LLMs to Vision-Language Models via On-Policy Distillation. *arXiv preprint arXiv:2510.23497*, 2025. URL <https://arxiv.org/abs/2510.23497>.
- Dan Busbridge, Amitis Shidani, Floris Weers, Jason Ramapuram, Etai Littwin, and Russ Webb. Distillation Scaling Laws. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2025. URL <https://arxiv.org/abs/2502.08606>.
- Yuang Cai and Yuyu Yuan. $\mathcal{X}$ -KD: General Experiential Knowledge Distillation for Large Language Models. *arXiv preprint arXiv:2602.12674*, 2026. URL <https://arxiv.org/abs/2602.12674>.
- Di Cao, Dongjie Fu, Hai Yu, Siqi Zheng, Xu Tan, and Tao Jin. X-OPD: Cross-Modal On-Policy Distillation for Capability Alignment in Speech LLMs. *arXiv preprint arXiv:2603.24596*, 2026. URL <https://arxiv.org/abs/2603.24596>.

- Yihan Cao, Yanbin Kang, Zhengming Xing, and Ruijie Jiang. Delta Knowledge Distillation for Large Language Models. *arXiv preprint arXiv:2509.14526*, 2025. URL <https://arxiv.org/abs/2509.14526>.
- Sungmin Cha and Kyunghyun Cho. Why Knowledge Distillation Works in Generative Models: A Minimal Working Explanation. *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. URL <https://arxiv.org/abs/2505.13111>.
- Howard Chen, Noam Razin, Karthik Narasimhan, and Danqi Chen. Retaining by Doing: The Role of On-Policy Data in Mitigating Forgetting. *arXiv preprint arXiv:2510.18874*, 2025. URL <https://arxiv.org/abs/2510.18874>.
- Zixiang Chen, Yihe Deng, Huizhuo Yuan, Kaixuan Ji, and Quanquan Gu. Self-Play Fine-Tuning Converts Weak Language Models to Strong Language Models. *Proceedings of ICML*, 2024. URL <https://arxiv.org/abs/2401.01335>.
- DeepSeek-AI. DeepSeek-V4: Towards highly efficient million-token context intelligence. Technical report, DeepSeek, 2026. URL https://huggingface.co/deepseek-ai/DeepSeek-V4-Pro/blob/main/DeepSeek_V4.pdf.
- DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, Bing Xue, Bingxuan Wang, Bochao Wu, Bei Feng, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Qu, Hui Li, Jianzhong Guo, Jiashi Li, Jiawei Wang, Jingchang Chen, Jingyang Yuan, Junjie Qiu, Junlong Li, J. L. Cai, Jiaqi Ni, Jian Liang, Jin Chen, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Liang Zhao, Litong Wang, Liyue Zhang, Lei Xu, Leyi Xia, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Meng Li, Miaojun Wang, Mingming Li, Ning Tian, Panpan Huang, Peng Zhang, Qiancheng Wang, Qinyu Chen, Qiushi Du, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, R. J. Chen, R. L. Jin, Ruyi Chen, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shengfeng Ye, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shuting Pan, S. S. Li, Shuang Zhou, Shaoqing Wu, Shengfeng Ye, Tao Yun, Tian Pei, Tianyu Sun, T. Wang, Wangding Zeng, Wanbiao Zhao, Wen Liu, Wenfeng Liang, Wenjun Gao, Wenqin Yu, Wentao Zhang, W. L. Xiao, Wei An, Xiaodong Liu, Xiaohan Wang, Xiaokang Chen, Xiaotao Nie, Xin Cheng, Xin Liu, Xin Xie, Xingchao Liu, Xinyu Yang, Xinyuan Li, Xuecheng Su, Xuheng Lin, X. Q. Li, Xiangyue Jin, Xiaojin Shen, Xiaosha Chen, Xiaowen Sun, Xiaoxiang Wang, Xinnan Song, Xinyi Zhou, Xianzu Wang, Xinxia Shan, Y. K. Li, Y. Q. Wang, Y. X. Wei, Yang Zhang, Yanhong Xu, Yao Li, Yao Zhao, Yaofeng Sun, Yaohui Wang, Yi Yu, Yichao Zhang, Yifan Shi, Yiliang Xiong, Ying He, Yishi Piao, Yisong Wang, Yixuan Tan, Yiyang Ma, Yiyuan Liu, Yongqiang Guo, Yuan Ou, Yuduan Wang, Yue Gong, Yuheng Zou, Yujia He, Yunfan Xiong, Yuxiang Luo, Yuxiang You, Yuxuan Liu, Yuyang Zhou, Y. X. Zhu, Yanhong Xu, Yanping Huang, Yaohui Li, Yi Zheng, Yuchen Zhu, Yunxian Ma, Ying Tang, Yukun Zha, Yuting Yan, Z. Z. Ren, Zehui Ren, Zhangli Sha, Zhe Fu, Zhean Xu, Zhenda Xie, Zhengyan Zhang, Zhewen Hao, Zhicheng Ma, Zhigang Yan, Zhiyu Wu, Zihui Gu, Zijia Zhu, Zijun Liu, Zilin Li, Ziwei Xie, Ziyang Song, Zizheng Pan, Zhen Huang, Zhipeng Xu, Zhongyu Zhang, and Zhen Zhang. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. *Nature*, 645:633–638, 2025. URL <https://arxiv.org/abs/2501.12948>.
- Ken Ding. HDPO: Hybrid Distillation Policy Optimization via Privileged Self-Distillation. *arXiv preprint arXiv:2603.23871*, 2026. URL <https://arxiv.org/abs/2603.23871>.
- Jiabin Fan, Guoqing Luo, Michael Bowling, and Lili Mou. KETCHUP: K-Step Return Estimation for Sequential Knowledge Distillation. *Proceedings of EACL*, 2026. URL <https://arxiv.org/abs/2504.19024>.
- Junfeng Fang, Zhepei Hong, Mao Zheng, Mingyang Song, Gengsheng Li, Houcheng Jiang, Dan Zhang, Haiyun Guo, Xiang Wang, and Tat-Seng Chua. Rubric-based On-policy Distillation. *arXiv preprint arXiv:2605.07396*, 2026a.

Zhen Fang, Wenxuan Huang, Yu Zeng, Yiming Zhao, Shuang Chen, Kaituo Feng, Yunlong Lin, Lin Chen, Zehui Chen, Shaosheng Cao, and Feng Zhao. Flow-OPD: On-Policy Distillation for Flow Matching Models. *arXiv preprint arXiv:2605.08063*, 2026b.

Yuqian Fu, Haohuan Huang, Kaiwen Jiang, Jiakai Liu, Zhuo Jiang, Yuanheng Zhu, and Dongbin Zhao. Revisiting On-Policy Distillation: Empirical Failure Modes and Simple Fixes. *arXiv preprint arXiv:2603.25562*, 2026. URL <https://arxiv.org/abs/2603.25562>.

Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, Léonard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ramé, Johan Ferret, Peter Liu, Pouya Tafti, Abe Friesen, Michelle Casbon, Sabela Ramos, Ravin Kumar, Charline Le Lan, Sammy Jerome, Anton Tsitsulin, Nino Vieillard, Piotr Stanczyk, Sertan Girgin, Nikola Momchev, Matt Hoffman, Shantanu Thakoor, Jean-Bastien Grill, Behnam Neyshabur, Olivier Bachem, Alanna Walton, Aliaksei Severyn, Alicia Parrish, Aliya Ahmad, Allen Hutchison, Alvin Abdagic, Amanda Carl, Amy Shen, Andy Brock, Andy Coenen, Anthony Laforge, Antonia Paterson, Ben Bastian, Bilal Piot, Bo Wu, Brandon Royal, Charlie Chen, Chintu Kumar, Chris Perry, Chris Welty, Christopher A. Choquette-Choo, Danila Sinopalnikov, David Weinberger, Dimple Vijaykumar, Dominika Rogozńska, Dustin Herbison, Elisa Bandy, Emma Wang, Eric Noland, Erica Moreira, Evan Senter, Evgenii Eltyshev, Francesco Visin, Gabriel Rasskin, Gary Wei, Glenn Cameron, Gus Martins, Hadi Hashemi, Hanna Klimczak-Plucińska, Harleen Batra, Harsh Dhand, Ivan Nardini, Jacinda Mein, Jack Zhou, James Svensson, Jeff Stanway, Jetha Chan, Jin Peng Zhou, Joana Carrasqueira, Joana Iljazi, Jocelyn Becker, Joe Fernandez, Joost van Amersfoort, Josh Gordon, Josh Lipschultz, Josh Newlan, Ju yeong Ji, Kareem Mohamed, Kartikeya Badola, Kat Black, Katie Millican, Keelin McDonell, Kelvin Nguyen, Kiranbir Sodhia, Kish Greene, Lars Lowe Sjoesund, Lauren Usui, Laurent Sifre, Lena Heuermann, Leticia Lago, Lilly McNealus, Livio Baldini Soares, Logan Kilpatrick, Lucas Dixon, Luciano Martins, Machel Reid, Manvinder Singh, Mark Iverson, Martin Görner, Mat Velloso, Mateo Wirth, Matt Davidow, Matt Miller, Matthew Rahtz, Matthew Watson, Meg Risdal, Mehran Kazemi, Michael Moynihan, Ming Zhang, Minsuk Kahng, Minwoo Park, Mofi Rahman, Mohit Khatwani, Natalie Dao, Nenshad Bardoliwalla, Nesh Devanathan, Neta Dumai, Nilay Chauhan, Oscar Wahltinez, Pankil Botarda, Parker Barnes, Paul Barham, Paul Michel, Pengchong Jin, Petko Georgiev, Phil Culliton, Pradeep Kuppala, Ramona Comanescu, Ramona Merhej, Reena Jana, Reza Ardeshtir Rokni, Rishabh Agarwal, Ryan Mullins, Samaneh Saadat, Sara Mc Carthy, Sarah Cogan, Sarah Perrin, Sébastien M. R. Arnold, Sebastian Krause, Shengyang Dai, Shruti Garg, Shruti Sheth, Sue Ronstrom, Susan Chan, Timothy Jordan, Ting Yu, Tom Eccles, Tom Hennigan, Tomas Kocisky, Tulsee Doshi, Vihan Jain, Vikas Yadav, Vilobh Meshram, Vishal Dharmadhikari, Warren Barkley, Wei Wei, Wenming Ye, Woohyun Han, Woosuk Kwon, Xiang Xu, Zhe Shen, Zhitao Gong, Zichuan Wei, Victor Cotruta, Phoebe Kirk, Anand Rao, Minh Giang, Ludovic Peran, Tris Warkentin, Eli Collins, Joelle Barral, Zoubin Ghahramani, Raia Hadsell, D. Sculley, Jeanine Banks, Anca Dragan, Slav Petrov, Oriol Vinyals, Jeff Dean, Demis Hassabis, Koray Kavukcuoglu, Clement Farabet, Elena Buchatskaya, Sebastian Borgeaud, Noah Fiedel, Armand Joulin, Kathleen Kenealy, Robert Dadashi, and Alek Andreev. Gemma 2: Improving Open Language Models at a Practical Size. *arXiv preprint arXiv:2408.00118*, 2024. URL <https://arxiv.org/abs/2408.00118>.

Naibin Gu, Chenxu Yang, Qingyi Si, Chuanyu Qin, Dingyu Yao, Peng Fu, Zheng Lin, Weiping Wang, Nan Duan, and Jiaqi Wang. Co-Evolving Policy Distillation. *arXiv preprint arXiv:2604.27083*, 2026. URL <https://arxiv.org/abs/2604.27083>.

Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge Distillation of Large Language Models. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2306.08543>.

Yuxian Gu, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. MiniPLM: Knowledge Distillation for Pre-Training Language Models. *Proceedings of ICLR*, 2025. URL <https://arxiv.org/abs/2410.17215>.

- Arnav Gudibande, Eric Wallace, Charlie Snell, Xinyang Geng, Hao Liu, Pieter Abbeel, Sergey Levine, and Dawn Song. The False Promise of Imitating Proprietary LLMs. *arXiv preprint arXiv:2305.15717*, 2023. URL <https://arxiv.org/abs/2305.15717>.
- Yinghui He, Simran Kaur, Adithya Bhaskar, Yongjin Yang, Jiarui Liu, Narutatsu Ri, Liam Fowl, Abhishek Panigrahi, Danqi Chen, and Sanjeev Arora. Self-Distillation Zero: Self-Revision Turns Binary Rewards into Dense Supervision. *arXiv preprint arXiv:2604.12002*, 2026. URL <https://arxiv.org/abs/2604.12002>.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Jordan Hoffmann, Sebastian Borgeaud, Arthur Mensch, Elena Buchatskaya, Trevor Cai, Eliza Rutherford, Diego de Las Casas, Lisa Anne Hendricks, Johannes Welbl, Aidan Clark, Tom Hennigan, Eric Noland, Katie Millican, George van den Driessche, Bogdan Damoc, Aurelia Guy, Simon Osindero, Karen Simonyan, Erich Elsen, Jack W. Rae, Oriol Vinyals, and Laurent Sifre. Training Compute-Optimal Large Language Models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. URL <https://arxiv.org/abs/2203.15556>.
- Wenjin Hou, Shangpin Peng, Weinong Wang, Zheng Ruan, Yue Zhang, Zhenglin Zhou, Mingqi Gao, Yifei Chen, Kaiqi Wang, Hongming Yang, Chengquan Zhang, Zhuotao Tian, Han Hu, Yi Yang, Fei Wu, and Hehe Fan. Uni-opd: Unifying on-policy distillation with a dual-perspective recipe. *arXiv preprint arXiv:2605.03677*, 2026.
- Cheng-Yu Hsieh, Chun-Liang Li, Chih-Kuan Yeh, Hootan Nakhost, Yasuhisa Fujii, Alexander Ratner, Ranjay Krishna, Chen-Yu Lee, and Tomas Pfister. Distilling Step-by-Step! Outperforming Larger Language Models with Less Training Data and Smaller Model Sizes. *Findings of ACL*, 2023. URL <https://arxiv.org/abs/2305.02301>.
- Jian Hu, Xibin Tao, Weixun Hou, Quanzhi Liu, Zhengyu Feng, Yang Xu, Qing Ye, Qi Li, Zenan Li, and Shuai Zhang. OpenRLHF: An easy-to-use, scalable and high-performance RLHF framework. *arXiv preprint arXiv:2405.11143*, 2024.
- Jing Hu, Danxiang Zhu, Xianlong Luo, Dan Zhang, Shuwei He, Yishu Lei, Haitao Zheng, Shikun Feng, Jingzhou He, Yu Sun, Hua Wu, and Haifeng Wang. CORD: Bridging the Audio-Text Reasoning Gap via Weighted On-policy Cross-modal Distillation. *arXiv preprint arXiv:2601.16547*, 2026. URL <https://arxiv.org/abs/2601.16547>.
- Haiduo Huang, Jiangcheng Song, Yadong Zhang, and Pengju Ren. SelecTKD: Selective Token-Weighted Knowledge Distillation for LLMs. *arXiv preprint arXiv:2510.24021*, 2025. URL <https://arxiv.org/abs/2510.24021>.
- Jonas Hübner, Frederike Lübeck, Lejs Behric, Anton Baumann, Marco Bagatella, Daniel Marta, Ido Hakimi, Idan Shenfeld, Thomas Kleine Buening, Carlos Guestrin, and Andreas Krause. Reinforcement Learning via Self-Distillation. *arXiv preprint arXiv:2601.20802*, 2026. URL <https://arxiv.org/abs/2601.20802>.
- Ijun Jang, Jewon Yeom, Juan Yeo, Hyunggu Lim, and Taesup Kim. Stable On-Policy Distillation through Adaptive Target Reformulation. In *Findings of ACL*, 2026. URL <https://arxiv.org/abs/2601.07155>.
- Minbyul Jeong. Healthcare AI GYM for medical agents. *arXiv preprint arXiv:2605.02943*, 2026.
- Nan Jia, Haojin Yang, Xing Ma, Jiesong Lian, Shuailiang Zhang, Weipeng Zhang, Ke Zeng, Xunliang Cai, and Zequn Sun. Asymmetric On-Policy Distillation: Bridging Exploitation and Imitation at the Token Level. *arXiv preprint arXiv:2605.06387*, 2026. URL <https://arxiv.org/abs/2605.06387>.
- Dengyang Jiang, Xin Jin, Dongyang Liu, Zanyi Wang, Mingzhe Zheng, Ruoyi Du, Xiangpeng Yang, Qilong Wu, Zhen Li, Peng Gao, Harry Yang, and Steven C.H. Hoi. D-OPSD: On-policy self-distillation for continuously tuning step-distilled diffusion models. *arXiv preprint arXiv:2605.05204*, 2026.

- Yuxin Jiang, Chunkit Chan, Mingyang Chen, and Wei Wang. Lion: Adversarial Distillation of Proprietary Large Language Models. *Proceedings of EMNLP*, 2023. URL <https://arxiv.org/abs/2305.12870>.
- Woogyel Jin, Taywon Min, Yongjin Yang, Swanand Ravindra Kadhe, Yi Zhou, Dennis Wei, Nathalie Baracaldo, and Kimin Lee. Entropy-Aware On-Policy Distillation of Language Models. *arXiv preprint arXiv:2603.07079*, 2026a. URL <https://arxiv.org/abs/2603.07079>.
- Yiqiao Jin, Yiyang Wang, Lucheng Fu, Yijia Xiao, Yinyi Luo, Haoxin Liu, B. Aditya Prakash, Josiah Hester, Jindong Wang, and Srijan Kumar. UniSD: Towards a Unified Self-Distillation Framework for Large Language Models. *arXiv preprint arXiv:2605.06597*, 2026b.
- Seongryong Jung, Suwan Yoon, DongGeon Kim, and Hwanhee Lee. ToDi: Token-wise Distillation via Fine-Grained Divergence Control. In *Proceedings of EMNLP*, 2025. URL <https://arxiv.org/abs/2505.16297>.
- Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling laws for neural language models. *arXiv preprint arXiv:2001.08361*, 2020.
- Fatemeh Khadem, Sajad Mousavi, Yi Fang, and Yuhong Liu. DP-OPD: Differentially Private On-Policy Distillation for Language Models. *arXiv preprint arXiv:2604.04461*, 2026. URL <https://arxiv.org/abs/2604.04461>.
- Gyeongman Kim, Doohyuk Jang, and Eunho Yang. PromptKD: Distilling Student-Friendly Knowledge for Generative Language Models via Prompt Tuning. In *Findings of EMNLP*, 2024. URL <https://arxiv.org/abs/2402.12842>.
- Jaehoon Kim and Dongha Lee. OPSD Compresses What RLVR Teaches: A Post-RL Compaction Stage for Reasoning Models. *arXiv preprint arXiv:2605.06188*, 2026. URL <https://arxiv.org/abs/2605.06188>.
- Jeonghye Kim, Xufang Luo, Minbeom Kim, Sangmook Lee, Dohyung Kim, Jiwon Jeon, Dongsheng Li, and Yuqing Yang. Why Does Self-Distillation (Sometimes) Degrade the Reasoning Capability of LLMs? *arXiv preprint arXiv:2603.24472*, 2026a. URL <https://arxiv.org/abs/2603.24472>.
- Minsang Kim and Seung Jun Baek. Explain in Your Own Words: Improving Reasoning via Token-Selective Dual Knowledge Distillation. *Proceedings of ICLR*, 2026. URL <https://arxiv.org/abs/2603.13260>.
- Yeongmin Kim, Donghyeok Shin, Mina Kang, Byeonghu Na, and Il-Chul Moon. Distillation of Large Language Models via Concrete Score Matching. In *International Conference on Learning Representations (ICLR)*, 2026b. URL <https://arxiv.org/abs/2509.25837>.
- Yoon Kim and Alexander M. Rush. Sequence-level knowledge distillation. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pp. 1317–1327, 2016.
- John Kirchenbauer, Abhimanyu Hans, Brian Bartoldson, Micah Goldblum, Ashwinee Panda, and Tom Goldstein. Multi-Token Prediction via Self-Distillation. *arXiv preprint arXiv:2602.06019*, 2026. URL <https://arxiv.org/abs/2602.06019>.
- James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, et al. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences*, 114(13):3521–3526, 2017.
- Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. DistiLLM: Towards Streamlined Distillation for Large Language Models. *Proceedings of ICML*, 2024. URL <https://arxiv.org/abs/2402.03898>.

- Jongwoo Ko, Tianyi Chen, Sungnyun Kim, Tianyu Ding, Luming Liang, Ilya Zharkov, and Se-Young Yun. DistiLLM-2: A Contrastive Approach Boosts the Distillation of LLMs. In *Proceedings of ICML*, 2025. URL <https://arxiv.org/abs/2503.07067>.
- Jongwoo Ko, Sara Abdali, Young Jin Kim, Tianyi Chen, and Pashmina Cameron. Scaling Reasoning Efficiently via Relaxed On-Policy Distillation. *arXiv preprint arXiv:2603.11137*, 2026. URL <https://arxiv.org/abs/2603.11137>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the ACM SIGOPS 29th Symposium on Operating Systems Principles (SOSP)*, 2023.
- Fengxiang Li, Han Zhang, Haoyang Huang, Jinghui Wang, Jinhua Hao, Kun Yuan, Mengtong Li, Minglei Zhang, Pengcheng Xu, Wenhao Zhuang, Yizhen Shao, Zongxian Feng, Can Tang, Chao Wang, Chengxiao Tong, Fan Yang, Gang Xiong, Haixuan Gao, Han Gao, Hao Wang, Haochen Liu, Hongliang Sun, Jiabao Li, Jingwen Chang, Jun Du, Junyi Peng, Leizhen Cui, Meimei Jing, Mingqi Wu, Shangpeng Yan, Shaotong Qi, Suzhe Xu, Wenxuan Zhao, Xianda Sun, Xuan Xie, Yanbo Wang, Yao Xia, Yinghan Cui, Yingpeng Chen, Yong Wang, Yuze Shi, Zhiwei Shen, Ziyu Wang, Ming Sun, Lin Ye, and Bin Chen. KAT-Coder-V2 Technical Report. *arXiv preprint arXiv:2603.27703*, 2026a. URL <https://arxiv.org/abs/2603.27703>.
- Gengsheng Li, Tianyu Yang, Junfeng Fang, Mingyang Song, Mao Zheng, Haiyun Guo, Dan Zhang, Jinqiao Wang, and Tat-Seng Chua. Unifying Group-Relative and Self-Distillation Policy Optimization via Sample Routing. *arXiv preprint arXiv:2604.02288*, 2026b. URL <https://arxiv.org/abs/2604.02288>.
- Jiaze Li, Hao Yin, Haoran Xu, Boshen Xu, Wenhui Tan, Zewen He, Jianzhong Ju, Zhenbo Luo, and Jian Luan. Video-OPD: Efficient Post-Training of Multimodal Large Language Models for Temporal Video Grounding via On-Policy Distillation. *arXiv preprint arXiv:2602.02994*, 2026c. URL <https://arxiv.org/abs/2602.02994>.
- Yaxuan Li, Yuxin Zuo, Bingxiang He, Jinqian Zhang, Chaojun Xiao, Cheng Qian, Tianyu Yu, Huan ang Gao, Wenkai Yang, Zhiyuan Liu, and Ning Ding. Rethinking On-Policy Distillation of Large Language Models: Phenomenology, Mechanism, and Recipe. *arXiv preprint arXiv:2604.13016*, 2026d. URL <https://arxiv.org/abs/2604.13016>.
- Yingru Li, Ziniu Li, and Jiakai Liu. A Note on Hybrid Online Reinforcement and Imitation Learning for LLMs: Formulations and Algorithms. *arXiv preprint arXiv:2512.23097*, 2025. URL <https://arxiv.org/abs/2512.23097>.
- Kun Liang, Clive Bai, Xin Xu, Chenming Tang, Sanwoo Lee, Weijie Liu, Saiyong Yang, and Yunfang Wu. ORBIT: On-policy Exploration-Exploitation for Controllable Multi-Budget Reasoning. *arXiv preprint arXiv:2601.08310*, 2026. URL <https://arxiv.org/abs/2601.08310>.
- Percy Liang, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, Benjamin Newman, Binhang Yuan, Bobby Yan, Ce Zhang, Christian Cosgrove, Christopher D. Manning, Christopher Ré, Diana Acosta-Navas, Drew A. Hudson, Eric Zelikman, Esin Durmus, Faisal Ladhak, Frieda Rong, Hongyu Ren, Huaxiu Yao, Jue Wang, Keshav Santhanam, Laurel Orr, Lucia Zheng, Mert Yuksekgonul, Mirac Suzgun, Nathan Kim, Neel Guha, Niladri Chatterji, Omar Khattab, Peter Henderson, Qian Huang, Ryan Chi, Sang Michael Xie, Shibani Santurkar, Surya Ganguli, Tatsunori Hashimoto, Thomas Icard, Tianyi Zhang, Vishrav Chaudhary, William Wang, Xuechen Li, Yifan Mai, Yuhui Zhang, and Yuta Koreeda. Holistic evaluation of language models. *Transactions on Machine Learning Research*, 2023.
- Wenjie Liao, Like Wu, Liangjie Zhao, Shihui Xu, and Shigeru Fujimura. IRIS: Interpolative Rényi Iterative Self-play for Large Language Model Fine-Tuning. *arXiv preprint arXiv:2604.20933*, 2026. URL <https://arxiv.org/abs/2604.20933>.

- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2305.20050>.
- Hao Lin, Kunyang Lv, Xu Jiang, Jingqi Tian, Zhongjing Du, Jiayu Ding, Qiaoman Zhang, and Hongbo Jin. VISD: Enhancing Video Reasoning via Structured Self-Distillation. *arXiv preprint arXiv:2605.06094*, 2026. URL <https://arxiv.org/abs/2605.06094>.
- Kevin Lu and Thinking Machines Lab. On-policy distillation. Thinking Machines Lab: Connectionism, 2025. URL <https://www.thinkingmachines.ai/blog/on-policy-distillation>. doi: 10.64434/tml.20251026.
- Feng Luo, Yu-Neng Chuang, Guanchu Wang, Zicheng Xu, Xiaotian Han, Tianyi Zhang, and Vladimir Braverman. Demystifying OPD: Length Inflation and Stabilization Strategies for Large Language Models. *arXiv preprint arXiv:2604.08527*, 2026. URL <https://arxiv.org/abs/2604.08527>.
- Yuanjie Lyu, Chengyu Wang, Jun Huang, and Tong Xu. From Correction to Mastery: Reinforced Distillation of Large Language Model Agents. *arXiv preprint arXiv:2509.14257*, 2025. URL <https://arxiv.org/abs/2509.14257>.
- Ethan Mendes, Jungsoo Park, and Alan Ritter. Didactic to Constructive: Turning Expert Solutions into Learnable Reasoning. *arXiv preprint arXiv:2602.02405*, 2026. URL <https://arxiv.org/abs/2602.02405>.
- Purbesh Mitra and Sennur Ulukus. Semantic Soft Bootstrapping: Long Context Reasoning in LLMs without Reinforcement Learning. *arXiv preprint arXiv:2512.05105*, 2025. URL <https://arxiv.org/abs/2512.05105>.
- NVIDIA. TensorRT-LLM: High-performance inference framework for large language models. <https://github.com/NVIDIA/TensorRT-LLM>, 2024.
- Minjae Oh, Sangjun Song, Gyubin Choi, Yunho Choi, and Yohan Jo. KL for a KL: On-Policy Distillation with Control Variate Baseline. *arXiv preprint arXiv:2605.07865*, 2026.
- Emiliano Penalosa, Dheeraj Vattikonda, Nicolas Gontier, Alexandre Lacoste, Laurent Charlin, and Massimo Caccia. Privileged Information Distillation for Language Models. *arXiv preprint arXiv:2602.04942*, 2026. URL <https://arxiv.org/abs/2602.04942>.
- Jingyu Peng, Maolin Wang, Hengyi Cai, Yuchen Li, Kai Zhang, Shuaiqiang Wang, Dawei Yin, and Xiangyu Zhao. AdaSwitch: Balancing Exploration and Guidance in Knowledge Distillation via Adaptive Switching. *arXiv preprint arXiv:2510.07842*, 2025. URL <https://arxiv.org/abs/2510.07842>.
- Ruiyang Qin, Qingzhuo Wang, Dongrui Liu, Qiang Li, Zhihua Wei, and Wen Shen. Multilingual Safety Alignment via Self-Distillation. *arXiv preprint arXiv:2605.02971*, 2026. URL <https://arxiv.org/abs/2605.02971>.
- Miao Rang, Zhenni Bi, Hang Zhou, Kai Han, Xuechun Wang, An Xiao, Xinghao Chen, Yunhe Wang, and Hanting Chen. Near-Policy Distillation: Accelerating On-Policy Distillation via Asynchronous Generation and Selective Packing. *arXiv preprint arXiv:2605.05940*, 2026. URL <https://arxiv.org/abs/2605.05940>.
- Stéphane Ross, Geoffrey J. Gordon, and J. Andrew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics*, pp. 627–635, 2011.
- Aswin RRV, Jacob Dineen, Divij Handa, Md Nayem Uddin, Mihir Parmar, Chitta Baral, and Ben Zhou. ThinkTuning: Instilling Cognitive Reflections without Distillation. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025. URL <https://arxiv.org/abs/2508.07616>.

- Hejian Sang, Yuanda Xu, Zhengze Zhou, Ran He, Zhipeng Wang, and Jiachen Sun. CRISP: Compressed Reasoning via Iterative Self-Policy Distillation. *arXiv preprint arXiv:2603.05433*, 2026. URL <https://arxiv.org/abs/2603.05433>.
- Idan Shenfeld, Mehul Damani, Jonas Hübner, and Pulkit Agrawal. Self-Distillation Enables Continual Learning. *arXiv preprint arXiv:2601.19897*, 2026. URL <https://arxiv.org/abs/2601.19897>.
- Guangming Sheng, Chi Cao, Zilingfeng Gao, et al. veRL: Volcano engine reinforcement learning for LLM. *arXiv preprint arXiv:2409.19256*, 2024.
- Makoto Shing, Kou Misaki, Han Bao, Sho Yokoi, and Takuya Akiba. TAID: Temporally Adaptive Interpolated Distillation for Efficient Knowledge Transfer in Language Models. *Proceedings of ICLR*, 2025. URL <https://arxiv.org/abs/2501.16937>.
- Aasheesh Singh, Vishal Vaddina, and Dagnachew Birru. ORPO-Distill: Mixed-Policy Preference Optimization for Cross-Architecture LLM Distillation. *NeurIPS Workshop on Efficient Reasoning*, 2025. URL <https://arxiv.org/abs/2509.25100>.
- Yuda Song, Lili Chen, Fahim Tajwar, Remi Munos, Deepak Pathak, J. Andrew Bagnell, Aarti Singh, and Andrea Zanette. Expanding the Capabilities of Reinforcement Learning via Text Feedback. *arXiv preprint arXiv:2602.02482*, 2026. URL <https://arxiv.org/abs/2602.02482>.
- Alex Stein, Furong Huang, and Tom Goldstein. GATES: Self-Distillation under Privileged Context with Consensus Gating. *arXiv preprint arXiv:2602.20574*, 2026. URL <https://arxiv.org/abs/2602.20574>.
- Jie Sun, Mao Zheng, Mingyang Song, Qiyong Zhong, Yilin Cheng, Bichuan Feng, Pengfei Liu, Junfeng Fang, and Xiang Wang. SimCT: Recovering Lost Supervision for Cross-Tokenizer On-Policy Distillation. *arXiv preprint arXiv:2605.07711*, 2026.
- Zhiqian Tan and Yinrong Hong. PAINT: Partial-Solution Adaptive Interpolated Training for Self-Distilled Reasoners. *arXiv preprint arXiv:2604.26573*, 2026. URL <https://arxiv.org/abs/2604.26573>.
- Tencent Robotics X, HY Vision Team, Xumin Yu, Zuyan Liu, Ziyi Wang, He Zhang, Yongming Rao, Fangfu Liu, Yani Zhang, Ruowen Zhao, Oran Wang, Yves Liang, Haitao Lin, Minghui Wang, Yubo Dong, Kevin Cheng, Bolin Ni, Rui Huang, Han Hu, Zhengyou Zhang, Linus, and Shunyu Yao. HY-Embodied-0.5: Embodied Foundation Models for Real-World Agents. *arXiv preprint arXiv:2604.07430*, 2026. URL <https://arxiv.org/abs/2604.07430>.
- Hao Wang, Guozhi Wang, Han Xiao, Yufeng Zhou, Yue Pan, Jichao Wang, Ke Xu, Yafei Wen, Xiaohu Ruan, Xiaoxin Chen, and Honggang Qi. Skill-SD: Skill-Conditioned Self-Distillation for Multi-turn LLM Agents. *arXiv preprint arXiv:2604.10674*, 2026a. URL <https://arxiv.org/abs/2604.10674>.
- Jianze Wang, Ying Liu, Jinlong Chen, Xuchun Hu, Qilong Zhang, Yu Cao, Jun Wang, Hua Yang, Yong Xie, and Qianglong Chen. MAD-OPD: Breaking the Ceiling in On-Policy Distillation via Multi-Agent Debate. *arXiv preprint arXiv:2605.01347*, 2026b. URL <https://arxiv.org/abs/2605.01347>.
- Jiaqi Wang, Wenhao Zhang, Weijie Shi, Yaliang Li, and James Cheng. TCO: Exploring Temporal Curriculum in On-Policy Distillation for Multi-turn Autonomous Agents. *arXiv preprint arXiv:2604.24005*, 2026c. URL <https://arxiv.org/abs/2604.24005>.
- Sudong Wang, Weiquan Huang, Xiaomin Yu, Zuhao Yang, Hehai Lin, Keming Wu, Chaojun Xiao, Chen Chen, Wenxuan Wang, Beier Zhu, Yunjian Zhang, and Chengwei Qin. Beyond SFT-to-RL: Pre-alignment via Black-Box On-Policy Distillation for Multimodal RL. *arXiv preprint arXiv:2604.28123*, 2026d. URL <https://arxiv.org/abs/2604.28123>.
- Wenshuo Wang. Knowledge Distillation Must Account for What It Loses. *arXiv preprint arXiv:2604.25110*, 2026.

- Yinjie Wang, Xuyang Chen, Xiaolong Jin, Mengdi Wang, and Ling Yang. OpenClaw-RL: Train Any Agent Simply by Talking. *arXiv preprint arXiv:2603.10165*, 2026e. URL <https://arxiv.org/abs/2603.10165>.
- Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f-Divergence Minimization for Sequence-Level Knowledge Distillation. *Proceedings of ACL*, 2023. URL <https://arxiv.org/abs/2307.15190>.
- Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. Rethinking Kullback-Leibler Divergence in Knowledge Distillation for Large Language Models. *Proceedings of COLING*, 2025. URL <https://arxiv.org/abs/2404.02657>.
- Yecheng Wu, Song Han, and Han Cai. Lightning OPD: Efficient Post-Training for Large Reasoning Models with Offline On-Policy Distillation. *arXiv preprint arXiv:2604.13010*, 2026a. URL <https://arxiv.org/abs/2604.13010>.
- Yubin Wu, Zicheng Cai, Liping Ning, Hua Wang, Zhi Chen, Yaohua Tang, and Hao Chen. LiteGUI: Distilling Compact GUI Agents with Reinforcement Learning. *arXiv preprint arXiv:2605.07505*, 2026b.
- Xiaomi LLM-Core Team, Bangjun Xiao, Bingquan Xia, Bo Yang, Bofei Gao, Bowen Shen, Chen Zhang, Chenhong He, Chiheng Lou, Fuli Luo, Gang Wang, Gang Xie, Hailin Zhang, Hanglong Lv, Hanyu Li, Heyu Chen, Hongshen Xu, Houbin Zhang, Huaqiu Liu, Jiangshan Duo, Jianyu Wei, Jiebao Xiao, Jinhao Dong, Jun Shi, Junhao Hu, Kainan Bao, Kang Zhou, Lei Li, Liang Zhao, Linghao Zhang, Peidian Li, Qianli Chen, Shaohui Liu, Shihua Yu, Shijie Cao, Shimao Chen, Shouqiu Yu, Shuo Liu, Tianling Zhou, Weijiang Su, Weikun Wang, Wenhan Ma, Xiangwei Deng, Bohan Mao, Bowen Ye, Can Cai, Chenghua Wang, Chengxuan Zhu, Chong Ma, Chun Chen, Chunan Li, Dawei Zhu, Deshan Xiao, Dong Zhang, Duo Zhang, Fangyue Liu, Feiyu Yang, Fengyuan Shi, Guoan Wang, Hao Tian, Hao Wu, Heng Qu, Hongfei Yi, Hongxu An, Hongyi Guan, Xing Zhang, Yifan Song, Yihan Yan, Yihao Zhao, Yingchun Lai, Yizhao Gao, Yu Cheng, Yuanyuan Tian, Yudong Wang, Zhen Tang, Zhengju Tang, Zhengtao Wen, Zhichao Song, Zhixian Zheng, Zihan Jiang, Jian Wen, Jiarui Sun, Jiawei Li, Jinlong Xue, Jun Xia, Kai Fang, Menghang Zhu, Nuo Chen, Qian Tu, Qihao Zhang, Qiying Wang, Rang Li, Rui Ma, Shaolei Zhang, Shengfan Wang, Shicheng Li, Shuhao Gu, Shuhuai Ren, Sirui Deng, Tao Guo, Tianyang Lu, Wei Ji Zhuang, Weikang Zhang, Weimin Xiong, Wenshan Huang, Wenyu Yang, Xin Zhang, Xing Yong, Xu Wang, Xueyang Xie, Yilin Jiang, Yixin Yang, Yongzhe He, Yu Tu, Yuanliang Dong, Yuchen Liu, Yue Ma, Yue Yu, Yuxing Xiang, Zhaojun Huang, Zhenru Lin, Zhipeng Xu, Zhiyang Chen, Zhonghua Deng, Zihan Zhang, and Zihao Yue. MiMo-V2-Flash Technical Report. *arXiv preprint arXiv:2601.02780*, 2026. URL <https://arxiv.org/abs/2601.02780>.
- Xurong Xie, Zhucun Xue, Jiafu Wu, Jian Li, Yabiao Wang, Xiaobin Hu, Yong Liu, and Jiangning Zhang. LLM-Oriented Token-Adaptive Knowledge Distillation. *Proceedings of AAAI*, 2026. URL <https://arxiv.org/abs/2510.11615>.
- Jing Xiong, Hui Shen, Shansan Gong, Yuxin Cheng, Jianghan Shen, Chaofan Tao, Haochen Tan, Haoli Bai, Lifeng Shang, and Ngai Wong. OVD: On-policy Verbal Distillation. *arXiv preprint arXiv:2601.21968*, 2026. URL <https://arxiv.org/abs/2601.21968>.
- Hongling Xu, Qi Zhu, Heyuan Deng, Jinpeng Li, Lu Hou, Yasheng Wang, Lifeng Shang, Ruifeng Xu, and Fei Mi. KDRL: Post-Training Reasoning LLMs via Unified Knowledge Distillation and Reinforcement Learning. *arXiv preprint arXiv:2506.02208*, 2025a. URL <https://arxiv.org/abs/2506.02208>.
- Shicheng Xu, Liang Pang, Yunchang Zhu, Jia Gu, Zihao Wei, Jingcheng Deng, Feiyang Pan, Huawei Shen, and Xueqi Cheng. RLKD: Distilling LLMs’ Reasoning via Reinforcement Learning. *arXiv preprint arXiv:2505.16142*, 2025b. URL <https://arxiv.org/abs/2505.16142>.
- Wenda Xu, Rujun Han, Zifeng Wang, Long T. Le, Dhruv Madeka, Lei Li, William Yang Wang, Rishabh Agarwal, Chen-Yu Lee, and Tomas Pfister. Speculative Knowledge Distillation: Bridging the Teacher-Student Gap Through Interleaved Sampling. *Proceedings of ICLR*, 2025c. URL <https://arxiv.org/abs/2410.11325>.

- Xiaohan Xu, Ming Li, Chongyang Tao, Tao Shen, Reynold Cheng, Jinyang Li, Can Xu, Dacheng Tao, and Tianyi Zhou. A Survey on Knowledge Distillation of Large Language Models. *arXiv preprint arXiv:2402.13116*, 2024. URL <https://arxiv.org/abs/2402.13116>.
- Yuanda Xu, Hejian Sang, Zhengze Zhou, Ran He, and Zhipeng Wang. PACED: Distillation and On-Policy Self-Distillation at the Frontier of Student Competence. *arXiv preprint arXiv:2603.11178*, 2026a. URL <https://arxiv.org/abs/2603.11178>.
- Yuanda Xu, Hejian Sang, Zhengze Zhou, Ran He, Zhipeng Wang, and Alborz Geramifard. TIP: Token Importance in On-Policy Distillation. *arXiv preprint arXiv:2604.14084*, 2026b. URL <https://arxiv.org/abs/2604.14084>.
- Jianhao Yan, Yafu Li, Zican Hu, Zhi Wang, Ganqu Cui, Xiaoye Qu, Yu Cheng, and Yue Zhang. Learning to Reason under Off-Policy Guidance. *arXiv preprint arXiv:2504.14945*, 2025. URL <https://arxiv.org/abs/2504.14945>.
- Shaotian Yan, Kaiyuan Liu, Chen Shen, Bing Wang, Sinan Fan, Jun Zhang, Yue Wu, Zheng Wang, and Jieping Ye. Distribution-Aligned Sequence Distillation for Superior Long-CoT Reasoning. *arXiv preprint arXiv:2601.09088*, 2026. URL <https://arxiv.org/abs/2601.09088>.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Yang, Jiayi Yang, Jing Zhou, Jingren Zhou, Junyang Lin, Kai Dang, Keqin Bao, Kexin Yang, Le Yu, Lianghao Deng, Mei Li, Mingfeng Xue, Mingze Li, Pei Zhang, Peng Wang, Qin Zhu, Rui Men, Ruize Gao, Shixuan Liu, Shuang Luo, Tianhao Li, Tianyi Tang, Wenbiao Yin, Xingzhang Ren, Xinyu Wang, Xinyu Zhang, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yinger Zhang, Yu Wan, Yuqiong Liu, Zekun Wang, Zeyu Cui, Zhenru Zhang, Zhipeng Zhou, and Zihan Qiu. Qwen3 Technical Report. *arXiv preprint arXiv:2505.09388*, 2025a. URL <https://arxiv.org/abs/2505.09388>.
- Chenxu Yang, Chuanyu Qin, Qingyi Si, Minghui Chen, Naibin Gu, Dingyu Yao, Zheng Lin, Weiping Wang, Jiaqi Wang, and Nan Duan. Self-Distilled RLVR. *arXiv preprint arXiv:2604.03128*, 2026a. URL <https://arxiv.org/abs/2604.03128>.
- Fan Yang, Rui Meng, Trudi Di Qi, Ali Ezzati, and Yuxin Wen. KEPO: Knowledge-Enhanced Preference Optimization for Reinforcement Learning with Reasoning. *arXiv preprint arXiv:2602.00400*, 2026b. URL <https://arxiv.org/abs/2602.00400>.
- Ling Yang, Zhaochen Yu, Tianjun Zhang, Minkai Xu, Joseph E. Gonzalez, Bin Cui, and Shuicheng Yan. SuperCorrect: Advancing Small LLM Reasoning with Thought Template Distillation and Self-Correction. *Proceedings of ICLR*, 2025b. URL <https://arxiv.org/abs/2410.09008>.
- Wenkai Yang, Weijie Liu, Ruobing Xie, Kai Yang, Saiyong Yang, and Yankai Lin. Learning beyond Teacher: Generalized On-Policy Distillation with Reward Extrapolation. *arXiv preprint arXiv:2602.12125*, 2026c. URL <https://arxiv.org/abs/2602.12125>.
- Zhicheng Yang, Zhijiang Guo, Yifan Song, Minrui Xu, Yongxin Wang, Yiwei Wang, Xiaodan Liang, and Jing Tang. Prune-OPD: Efficient and Reliable On-Policy Distillation for Long-Horizon Reasoning. *arXiv preprint arXiv:2605.07804*, 2026d.
- Zhuolin Yang, Zihan Liu, Yang Chen, Wenliang Dai, Boxin Wang, Sheng-Chieh Lin, Chankyu Lee, Yangyi Chen, Dongfu Jiang, Jiafan He, Renjie Pi, Grace Lam, Nayeon Lee, Alexander Bukharin, Mohammad Shoeybi, Bryan Catanzaro, and Wei Ping. Nemotron-Cascade 2: Post-Training LLMs with Cascade RL and Multi-Domain On-Policy Distillation. *arXiv preprint arXiv:2603.19220*, 2026e. URL <https://arxiv.org/abs/2603.19220>.
- Tianzhu Ye, Li Dong, Zewen Chi, Xun Wu, Shaohan Huang, and Furu Wei. Black-Box On-Policy Distillation of Large Language Models. *arXiv preprint arXiv:2511.10643*, 2025. URL <https://arxiv.org/abs/2511.10643>.

- Tianzhu Ye, Li Dong, Qingxiu Dong, Xun Wu, Shaohan Huang, and Furu Wei. Online Experiential Learning for Language Models. *arXiv preprint arXiv:2603.16856*, 2026a. URL <https://arxiv.org/abs/2603.16856>.
- Tianzhu Ye, Li Dong, Xun Wu, Shaohan Huang, and Furu Wei. On-Policy Context Distillation for Language Models. *arXiv preprint arXiv:2602.12275*, 2026b. URL <https://arxiv.org/abs/2602.12275>.
- Xin Yu, Liuchen Liao, Yiwen Zhang, Yingchen Yu, Lingzhou Xue, and Qinzhen Guo. Preference-based self-distillation: Beyond KL matching via reward regularization. *arXiv preprint arXiv:2605.05040*, 2026.
- Dongxu Zhang, Zhichao Yang, Sepehr Janghorbani, Jun Han, Andrew Ressler, Qian Qian, Gregory D. Lyng, Sanjit Singh Batra, and Robert E. Tillman. Fast and Effective On-policy Distillation from Reasoning Prefixes. *arXiv preprint arXiv:2602.15260*, 2026a. URL <https://arxiv.org/abs/2602.15260>.
- Jiaxin Zhang, Xiangyu Peng, Qinglin Chen, Qinyuan Ye, Caiming Xiong, and Chien-Sheng Wu. The Illusion of Certainty: Decoupling Capability and Calibration in On-Policy Distillation. *arXiv preprint arXiv:2604.16830*, 2026b. URL <https://arxiv.org/abs/2604.16830>.
- Miaosen Zhang, Yishan Liu, Shuxia Lin, Xu Yang, Qi Dai, Chong Luo, Weihao Jiang, Peng Hou, Anxiang Zeng, Xin Geng, and Baining Guo. Towards On-Policy SFT: Distribution Discriminant Theory and its Applications in LLM Training. *arXiv preprint arXiv:2602.12222*, 2026c. URL <https://arxiv.org/abs/2602.12222>.
- Ruixiang Zhang, Richard He Bai, Huangjie Zheng, Navdeep Jaitly, Ronan Collobert, and Yizhe Zhang. Embarrassingly Simple Self-Distillation Improves Code Generation. *arXiv preprint arXiv:2604.01193*, 2026d. URL <https://arxiv.org/abs/2604.01193>.
- Songming Zhang, Xue Zhang, Tong Zhang, Bojie Hu, Yufeng Chen, and Jinan Xu. AlignDis-til: Token-Level Language Model Alignment as Adaptive Policy Distillation. In *Proceedings of ACL*, 2025a. URL <https://arxiv.org/abs/2503.02832>.
- Xinsen Zhang, Zhenkai Ding, Tianjun Pan, Run Yang, Chun Kang, Xue Xiong, and Jingnan Gu. OPSDL: On-Policy Self-Distillation for Long-Context Language Models. *arXiv preprint arXiv:2604.17535*, 2026e. URL <https://arxiv.org/abs/2604.17535>.
- Xue Zhang, Songming Zhang, Yunlong Liang, Fandong Meng, Yufeng Chen, Jinan Xu, and Jie Zhou. A Dual-Space Framework for General Knowledge Distillation of Large Language Models. *arXiv preprint arXiv:2504.11426*, 2025b. URL <https://arxiv.org/abs/2504.11426>.
- Yan Zhang, Daiqing Wu, Huawen Shen, Yu Zhou, and Can Ma. Learn where to click from yourself: On-policy self-distillation for gui grounding. *arXiv preprint arXiv:2605.00642*, 2026f.
- Yaocheng Zhang, Yuanheng Zhu, Wenyue Chong, Songjun Tu, Qichao Zhang, Jiajun Chai, Xiaohan Wang, Wei Lin, Guojun Yin, and Dongbin Zhao. π -Play: Multi-Agent Self-Play via Privileged Self-Distillation without External Data. *arXiv preprint arXiv:2604.14054*, 2026g. URL <https://arxiv.org/abs/2604.14054>.
- Zhaoyang Zhang, Shuli Jiang, Yantao Shen, Yuting Zhang, Dhananjay Ram, Shuo Yang, Zhuowen Tu, Wei Xia, and Stefano Soatto. Reinforcement-aware Knowledge Distillation for LLM Reasoning. *arXiv preprint arXiv:2602.22495*, 2026h. URL <https://arxiv.org/abs/2602.22495>.
- Anhao Zhao, Ziyang Chen, Junlong Tong, Yingqi Fan, Fanghua Ye, Shuhao Li, Yunpu Ma, Wenjie Li, and Xiaoyu Shen. On-Policy Supervised Fine-Tuning for Efficient Reasoning. *arXiv preprint arXiv:2602.13407*, 2026a. URL <https://arxiv.org/abs/2602.13407>.
- Siyan Zhao, Zhihui Xie, Mengchen Liu, Jing Huang, Guan Pang, Feiyu Chen, and Aditya Grover. Self-Distilled Reasoner: On-Policy Self-Distillation for Large Language Models. *arXiv preprint arXiv:2601.18734*, 2026b. URL <https://arxiv.org/abs/2601.18734>.

- Binbin Zheng, Xing Ma, Yiheng Liang, Jingqing Ruan, Xiaoliang Fu, Kepeng Lin, Benchang Zhu, Ke Zeng, and Xunliang Cai. SCOPE: Signal-Calibrated On-Policy Distillation Enhancement with Dual-Path Adaptive Weighting. *arXiv preprint arXiv:2604.10688*, 2026. URL <https://arxiv.org/abs/2604.10688>.
- Qihuang Zhong, Liang Ding, Li Shen, Juhua Liu, Bo Du, and Dacheng Tao. Revisiting Knowledge Distillation for Autoregressive Language Models. *Proceedings of ACL*, 2024. URL <https://arxiv.org/abs/2402.11890>.
- Qiyong Zhong, Mao Zheng, Mingyang Song, Xin Lin, Jie Sun, Houcheng Jiang, Xiang Wang, and Junfeng Fang. SOD: Step-wise On-policy Distillation for Small Language Model Agents. *arXiv preprint arXiv:2605.07725*, 2026a.
- Zhide Zhong, Haodong Yan, Junfeng Li, Junjie He, Tianran Zhang, and Haoang Li. VLA-OPD: Bridging Offline SFT and Online RL for Vision-Language-Action Models via On-Policy Distillation. *arXiv preprint arXiv:2603.26666*, 2026b. URL <https://arxiv.org/abs/2603.26666>.
- Yongchao Zhou, Kaifeng Lyu, Ankit Singh Rawat, Aditya Krishna Menon, Afshin Ros-tamizadeh, Sanjiv Kumar, Jean-François Kagy, and Rishabh Agarwal. DistillSpec: Improving Speculative Decoding via Knowledge Distillation. *Proceedings of ICLR*, 2024. URL <https://arxiv.org/abs/2310.08461>.
- Matthieu Zimmer, Xiaotong Ji, Tu Nguyen, and Haitham Bou Ammar. Rethinking Large Language Model Distillation: A Constrained Markov Decision Process Perspective. *arXiv preprint arXiv:2509.22921*, 2025. URL <https://arxiv.org/abs/2509.22921>.